
Amazon Simple Storage Service

Console User Guide



Amazon Simple Storage Service: Console User Guide

Copyright © 2019 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

Welcome to the Amazon S3 Console User Guide	1
Changing the Console Language	2
Creating and Configuring a Bucket	3
Creating a Bucket	3
More Info	8
Deleting a Bucket	8
More Info	10
Emptying a Bucket	10
Viewing Bucket Properties	11
Enabling or Disabling Versioning	12
Enabling Default Encryption	13
More Info	16
Enabling Server Access Logging	16
Enabling Object-Level Logging	19
More Info	21
Configuring Static Website Hosting	21
Redirecting Website Requests	24
Advanced Settings	25
Setting Up a Destination for Event Notifications	25
Enabling and Configuring Event Notifications	27
Enabling Transfer Acceleration	31
Uploading, Downloading, and Managing Objects	33
Uploading S3 Objects	33
Uploading Files and Folders by Using Drag and Drop	34
Uploading Files by Pointing and Clicking	39
More Info	41
Downloading S3 Objects	42
Related Topics	44
Deleting Objects	44
More Info	46
Undeleting Objects	46
More Info	47
Restoring Archived S3 Objects	47
Archive Retrieval Options	48
Restoring an Archived S3 Object	48
Upgrade an In-Progress Restore	50
Checking Archive Restore Status and Expiration Date	51
Locking S3 Objects	52
More Info	54
Viewing an Overview of an Object	55
More Info	56
Viewing Object Versions	57
More Info	58
Viewing Object Properties	58
Adding Encryption to an Object	60
More Info	63
Adding Metadata to an Object	63
Adding System-Defined Metadata	64
Adding User-Defined Metadata	66
Adding Tags to an Object	68
More Info	71
Using Folders	71
Creating a Folder	72
Deleting Folders	73

Making Folders Public	75
Batch Operations	76
Creating a Batch Operations Job	76
More Info	77
Managing Batch Operations Jobs	77
Storage Management	78
Creating a Lifecycle Policy	78
Creating Cross-Region Replication Rules	81
Adding a CRR Rule When the Destination Bucket Is in the Same AWS Account	82
Adding a CRR Rule When the Destination Bucket Is in a Different AWS Account	89
More Info	95
Managing Cross-Region Replication Rules	95
More Info	97
Configuring Storage Class Analysis	97
Configuring Amazon S3 Inventory	101
Destination Bucket Policy	103
Grant Amazon S3 Permission to Encrypt Using Your AWS KMS Key	104
Configuring Request Metrics	105
Configuring a Request Metrics Filter	107
Setting Permissions	110
Blocking Public Access	110
Access Status	111
More Info	111
Editing Bucket Public Access Settings	112
Editing Public Access Settings for an S3 Bucket	112
Editing Public Access Settings for Multiple S3 Buckets	113
More Info	114
Editing Account Public Access Settings	115
More Info	116
Setting Object Permissions	116
More Info	119
Setting ACL Bucket Permissions	119
More Info	122
Adding a Bucket Policy	123
More Info	124
Allowing Cross-Domain Resource Sharing with CORS	124
More Info	125
Document History	126
Earlier Updates	126
AWS Glossary	128

Welcome to the Amazon S3 Console User Guide

Welcome to the *Amazon Simple Storage Service Console User Guide* for the Amazon Simple Storage Service (Amazon S3) console.

Amazon S3 provides virtually limitless storage on the internet. This guide explains how you can manage buckets, objects, and folders in Amazon S3 by using the AWS Management Console, a browser-based graphical user interface for interacting with AWS services.

For detailed conceptual information about how Amazon S3 works, see [What Is Amazon S3?](#) in the *Amazon Simple Storage Service Developer Guide*. The developer guide also has detailed information about Amazon S3 features and code examples to support those features.

Topics

- [Creating and Configuring an S3 Bucket \(p. 3\)](#)
- [Uploading, Downloading, and Managing Objects \(p. 33\)](#)
- [Storage Management \(p. 78\)](#)
- [Setting Bucket and Object Access Permissions \(p. 110\)](#)

How Do I Change the Language of the Amazon S3 Console?

You can change the display language of the Amazon S3 console. Several languages are supported.

To change the console language

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Scroll down until you see the bar at the bottom of the window, and then choose the language on the left side of the bar.



3. Choose the language that you want from the menu.



Creating and Configuring an S3 Bucket

Amazon S3 is cloud storage for the Internet. To upload your data (photos, videos, documents etc.), you first create a bucket in one of the AWS Regions. You can then upload your data objects to the bucket.

Every object you store in Amazon S3 resides in a bucket. You can use buckets to group related objects in the same way that you use a directory to group files in a file system.

Amazon S3 creates buckets in the AWS Region that you specify. You can choose any AWS Region that is geographically close to you to optimize latency, minimize costs, or address regulatory requirements. For example, if you reside in Europe, you might find it advantageous to create buckets in the EU (Ireland) or EU (Frankfurt) regions. For a list of Amazon S3 AWS Regions, see [Regions and Endpoints](#) in the *Amazon Web Services General Reference*.

You are not charged for creating a bucket. You are only charged for storing objects in the bucket and for transferring objects out of the bucket. For more information about pricing, see [Amazon Simple Storage Service \(S3\) FAQs](#).

Amazon S3 bucket names are globally unique, regardless of the AWS Region in which you create the bucket. You specify the name at the time you create the bucket. For bucket naming guidelines, see [Bucket Restrictions and Limitations](#) in the *Amazon Simple Storage Service Developer Guide*.

The following topics explain how to use the Amazon S3 console to create, delete, and manage buckets.

Topics

- [How Do I Create an S3 Bucket? \(p. 3\)](#)
- [How Do I Delete an S3 Bucket? \(p. 8\)](#)
- [How Do I Empty an S3 Bucket? \(p. 10\)](#)
- [How Do I View the Properties for an S3 Bucket? \(p. 11\)](#)
- [How Do I Enable or Suspend Versioning for an S3 Bucket? \(p. 12\)](#)
- [How Do I Enable Default Encryption for an S3 Bucket? \(p. 13\)](#)
- [How Do I Enable Server Access Logging for an S3 Bucket? \(p. 16\)](#)
- [How Do I Enable Object-Level Logging for an S3 Bucket with AWS CloudTrail Data Events? \(p. 19\)](#)
- [How Do I Configure an S3 Bucket for Static Website Hosting? \(p. 21\)](#)
- [How Do I Redirect Requests to an S3 Bucket Hosted Website to Another Host? \(p. 24\)](#)
- [Advanced Settings for S3 Bucket Properties \(p. 25\)](#)


How Do I Create an S3 Bucket?

Before you can upload data to Amazon S3, you must create a bucket in one of the AWS Regions to store your data in. After you create a bucket, you can upload an unlimited number of data objects to the bucket.

Buckets have configuration properties, including their geographical region, who has access to the objects in the bucket, and other metadata.

To create an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Choose **Create bucket**.



3. On the **Name and region** page, type a name for your bucket and choose the AWS Region where you want the bucket to reside. Complete the fields on this page as follows:
 - a. For **Bucket name**, type a unique DNS-compliant name for your new bucket. Follow these naming guidelines:
 - The name must be unique across all existing bucket names in Amazon S3.
 - The name must not contain uppercase characters.
 - The name must start with a lowercase letter or number.
 - The name must be between 3 and 63 characters long.
 - After you create the bucket you cannot change the name, so choose wisely.
 - Choose a bucket name that reflects the objects in the bucket because the bucket name is visible in the URL that points to the objects that you're going to put in your bucket.

For information about naming buckets, see [Rules for Bucket Naming](#) in the *Amazon Simple Storage Service Developer Guide*.

- b. For **Region**, choose the AWS Region where you want the bucket to reside. Choose a Region close to you to minimize latency and costs, or to address regulatory requirements. Objects stored in a Region never leave that Region unless you explicitly transfer them to another Region. For a list of Amazon S3 AWS Regions, see [Regions and Endpoints](#) in the *Amazon Web Services General Reference*.
- c. (Optional) If you have already set up a bucket that has the same settings that you want to use for the new bucket that you want to create, you can set it up quickly by choosing **Copy settings from an existing bucket**, and then choosing the bucket whose settings you want to copy.

The settings for the following bucket properties are copied: versioning, tags, and logging.

- d. Do one of the following:
 - If you copied settings from another bucket, choose **Create**. You're done, so skip the following steps.
 - If not, choose **Next**.

The screenshot shows the 'Create bucket' wizard in the Amazon S3 console. The wizard has a blue header with the title 'Create bucket' and a close button (X). Below the header is a progress bar with four steps: 1. Name and region (active), 2. Configure options, 3. Set permissions, and 4. Review. The main content area is dark blue and contains the following fields:

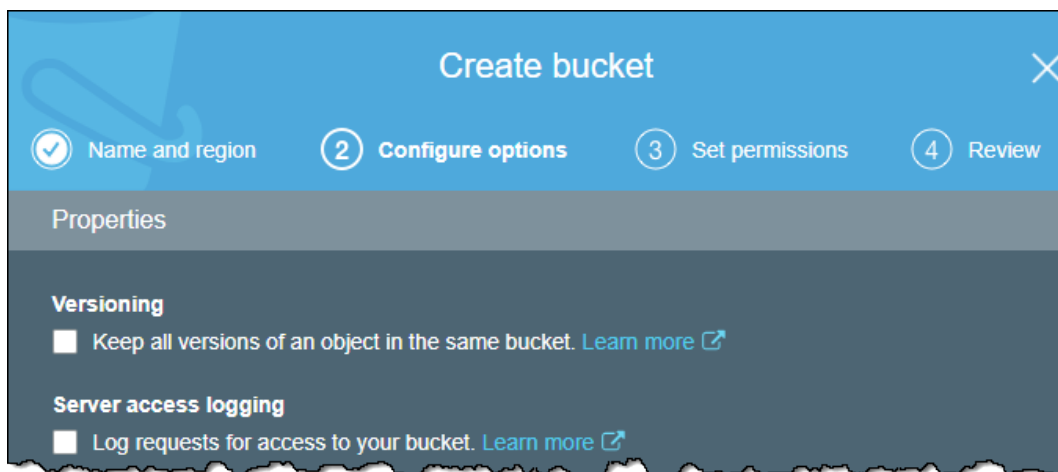
- Name and region**
 - Bucket name**: A text input field containing 'admin-created'.
 - Region**: A dropdown menu showing 'US West (Oregon)'.
 - Copy settings from an existing bucket**: A section with a dropdown menu showing 'Select bucket (optional)' and a count of '47 Buckets'.

At the bottom of the form are three buttons: 'Create' (disabled), 'Cancel', and 'Next' (active).

4. On the **Configure options** page, you can configure the following properties and Amazon CloudWatch metrics for the bucket. Or, you can configure these properties and CloudWatch metrics later, after you create the bucket.
 - a. **Versioning**

Select **Keep all versions of an object in the same bucket.** to enable object versioning for the bucket. For more information on enabling versioning, see [How Do I Enable or Suspend Versioning for an S3 Bucket?](#) (p. 12).
 - b. **Server access logging**

Select **Log requests for access to your bucket.** to enable server access logging on the bucket. Server access logging provides detailed records for the requests that are made to your bucket. For more information about enabling server access logging, see [How Do I Enable Server Access Logging for an S3 Bucket?](#) (p. 16).



Create bucket

1 Name and region 2 Configure options 3 Set permissions 4 Review

Properties

Versioning

☐ Keep all versions of an object in the same bucket. [Learn more](#)

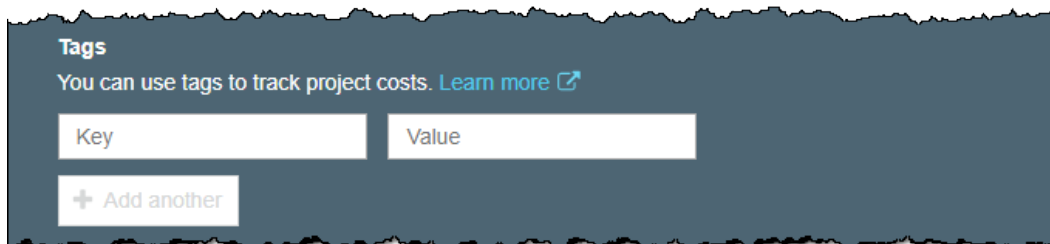
Server access logging

☐ Log requests for access to your bucket. [Learn more](#)

c. **Tags**

You can use cost allocation bucket tags to annotate billing for your use of a bucket. Each tag is a key-value pair that represents a label that you assign to a bucket.

To add a tag, type a *Key* and a *Value*. Choose **Add another** to add another tag. For more information about cost allocation tags, see [Using Cost Allocation S3 Bucket Tags](#) in the *Amazon Simple Storage Service Developer Guide*.



Tags

You can use tags to track project costs. [Learn more](#)

Key Value

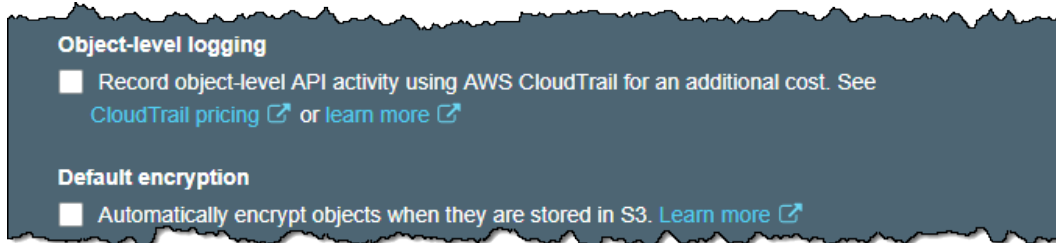
+ Add another

d. **Object-level logging**

Select **Record object-level API activity by using CloudTrail for an additional cost.** to enable object-level logging with CloudTrail. For more information about enabling object-level logging, see [How Do I Enable Object-Level Logging for an S3 Bucket with AWS CloudTrail Data Events?](#) (p. 19).

e. **Default encryption**

Select **Automatically encrypt objects when they are stored in S3.** to enable default encryption for the bucket. You can enable default encryption for a bucket so that all objects are encrypted when they are stored in the bucket. For more information about enabling default encryption, see [How Do I Enable Default Encryption for an S3 Bucket?](#) (p. 13).



Object-level logging

☐ Record object-level API activity using AWS CloudTrail for an additional cost. See [CloudTrail pricing](#) or [learn more](#)

Default encryption

☐ Automatically encrypt objects when they are stored in S3. [Learn more](#)

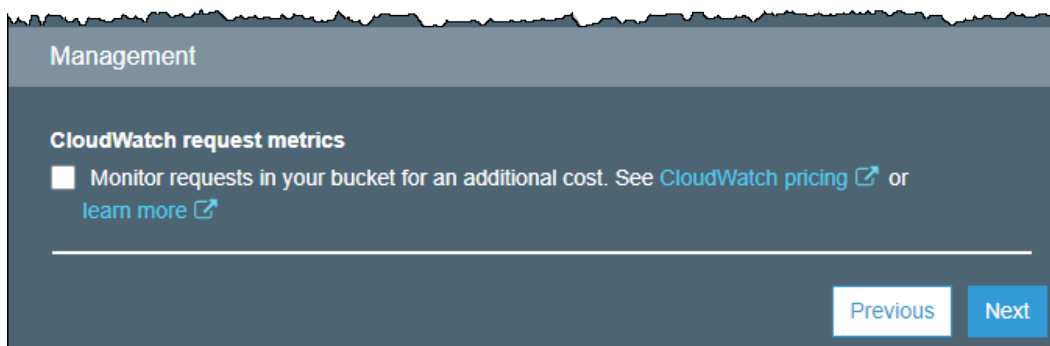
f. **Object lock**

Select **Permanently allow objects in this bucket to be locked** if you want to be able to lock objects in the bucket. Object lock requires that versioning be enabled on the bucket. For more information about object locking, see [Introduction to Amazon S3 Object Lock](#) in the *Amazon Simple Storage Service Developer Guide*



g. **CloudWatch request metrics**

Select **Monitor requests in your bucket for an additional cost** to configure CloudWatch request metrics for the bucket. For more information about CloudWatch request metrics, see [How Do I Configure Request Metrics for an S3 Bucket?](#) (p. 105).



5. Choose **Next**.

6. On the **Set permissions** page, you manage the permissions that are set on the bucket that you are creating.

Under **Public access settings for this bucket**, we recommend that you do not change the default settings that are listed under **Manage public access control lists (ACLs) for this bucket** or **Manage public bucket policies for this bucket**. You can change the permissions after you create the bucket. For more information about setting bucket permissions, see [How Do I Set ACL Bucket Permissions?](#) (p. 119).

Warning

We highly recommend that you keep the default access settings for blocking public access to the bucket that you are creating. Public access means that anyone in the world can access the objects in the bucket.

If you intend to use the bucket to store Amazon S3 server access logs, in the drop-down list under **Manage system permissions**, choose **Grant Amazon S3 Log Delivery group write access to this bucket**. For more information about server access logs, see [How Do I Enable Server Access Logging for an S3 Bucket?](#) (p. 16).

The screenshot shows the 'Create bucket' wizard in the AWS Management Console, specifically the 'Set permissions' step (indicated by a circled 3). The wizard has four steps: 1. Name and region, 2. Configure options, 3. Set permissions, and 4. Review. A note at the top states: 'Note: You can grant access to specific users after you create the bucket.' The main section is titled 'Public access settings for this bucket' and includes a link to 'Learn more' about Amazon S3 block public access settings. Below this, there are two sections: 'Manage public access control lists (ACLs) for this bucket' and 'Manage public bucket policies for this bucket'. Both sections have two checkboxes each, all of which are checked and marked as '(Recommended)'. The ACL section includes 'Block new public ACLs and uploading public objects' and 'Remove public access granted through public ACLs'. The policy section includes 'Block new public bucket policies' and 'Block public and cross-account access if bucket has public policies'. At the bottom, there is a 'Manage system permissions' section with a dropdown menu currently set to 'Do not grant Amazon S3 Log Delivery group write access to this bucket'. At the bottom right, there are 'Previous' and 'Next' buttons.

When you're done configuring permissions on the bucket, choose **Next**.

7. On the **Review** page, verify the settings. If you want to change something, choose **Edit**. If your current settings are correct, choose **Create bucket**.

More Info

- [How Do I Delete an S3 Bucket? \(p. 8\)](#)
- [How Do I Set ACL Bucket Permissions? \(p. 119\)](#)

How Do I Delete an S3 Bucket?

You can delete an empty bucket, and when you're using the AWS Management Console, you can delete a bucket that contains objects. If you delete a bucket that contains objects, all the objects in the bucket are permanently deleted.

When you delete a bucket with versioning enabled, all versions of all the objects in the bucket are permanently deleted. For more information about versioning, see [Managing Objects in a Versioning-Enabled Bucket](#) in the *Amazon Simple Storage Service Developer Guide*.

Before deleting a bucket, consider the following:

- Bucket names are unique. If you delete a bucket, another AWS user can use the name.
- When you delete a bucket that contains objects, all the objects in the bucket are permanently deleted, including objects that transitioned to the Amazon S3 Glacier storage class.
- If the bucket hosts a static website, and you created and configured an Amazon Route 53 hosted zone as described in [Create and Configure Amazon Route 53 Hosted Zone](#): You must clean up the Route 53 hosted zone settings that are related to the bucket as described in [Delete the Route 53 Hosted Zone](#).
- If the bucket receives log data from Elastic Load Balancing (ELB): We recommend that you stop the delivery of ELB logs to the bucket before deleting it. After you delete the bucket, if another user creates a bucket using the same name, your log data could potentially be delivered to that bucket. For

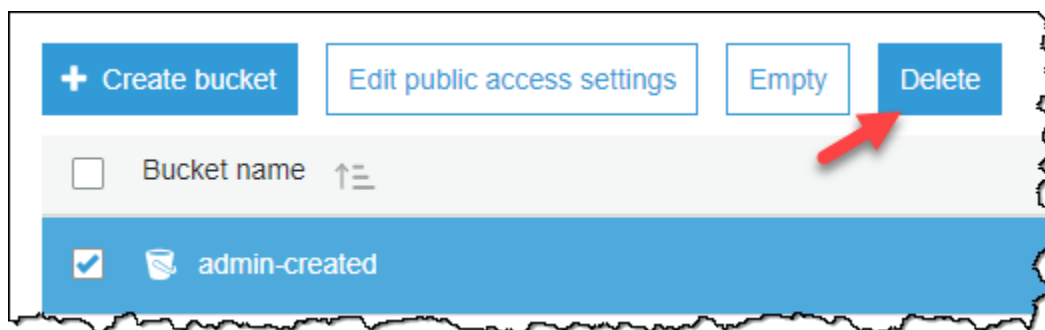
information about ELB access logs, see [Access Logs](#) in the *User Guide for Classic Load Balancers* and [Access Logs](#) in the *User Guide for Application Load Balancers*.

Important

If you want to continue to use the same bucket name, don't delete the bucket. We recommend that you empty the bucket and keep it. After a bucket is deleted, the name becomes available to reuse, but the name might not be available for you to reuse for various reasons. For example, it might take some time before the name can be reused, and some other account could create a bucket with that name before you do.

To delete an S3 bucket

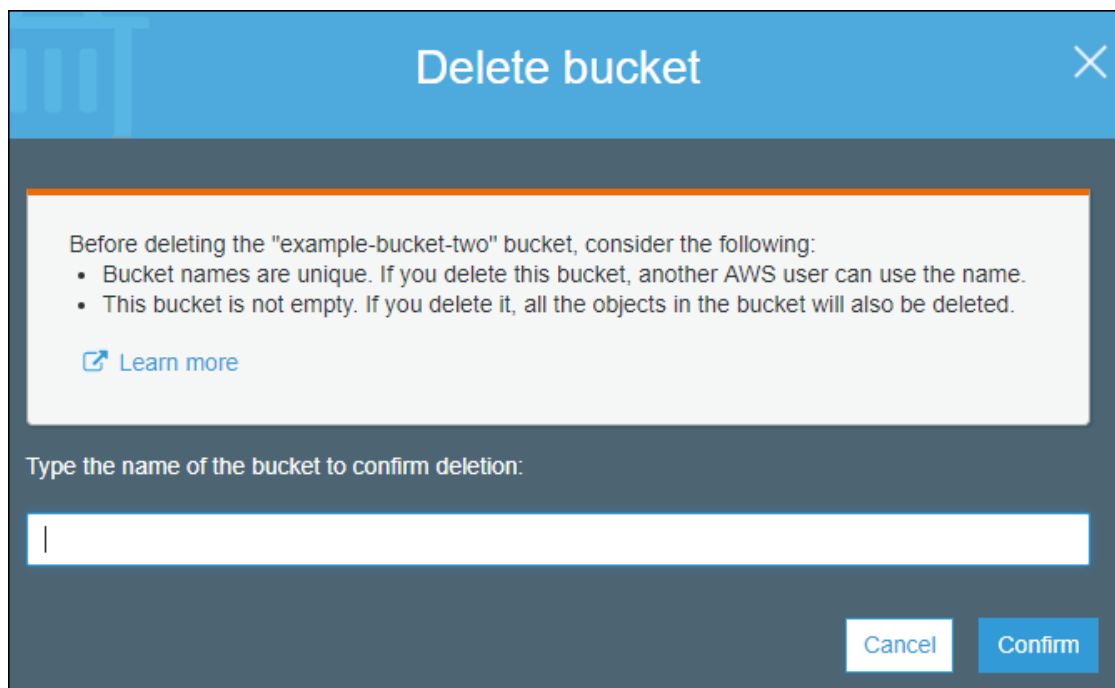
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the bucket icon next to the name of the bucket that you want to delete and then choose **Delete bucket**.



3. In the **Delete bucket** dialog box, type the name of the bucket that you want to delete for confirmation, and then choose **Confirm**.

Note

The text in the dialog box changes depending on whether the bucket is empty, is used for a static website, or is used for ELB access logs.



More Info


- [How Do I Empty an S3 Bucket?](#) (p. 10)
- [How Do I Delete Objects from an S3 Bucket?](#) (p. 44)

How Do I Empty an S3 Bucket?

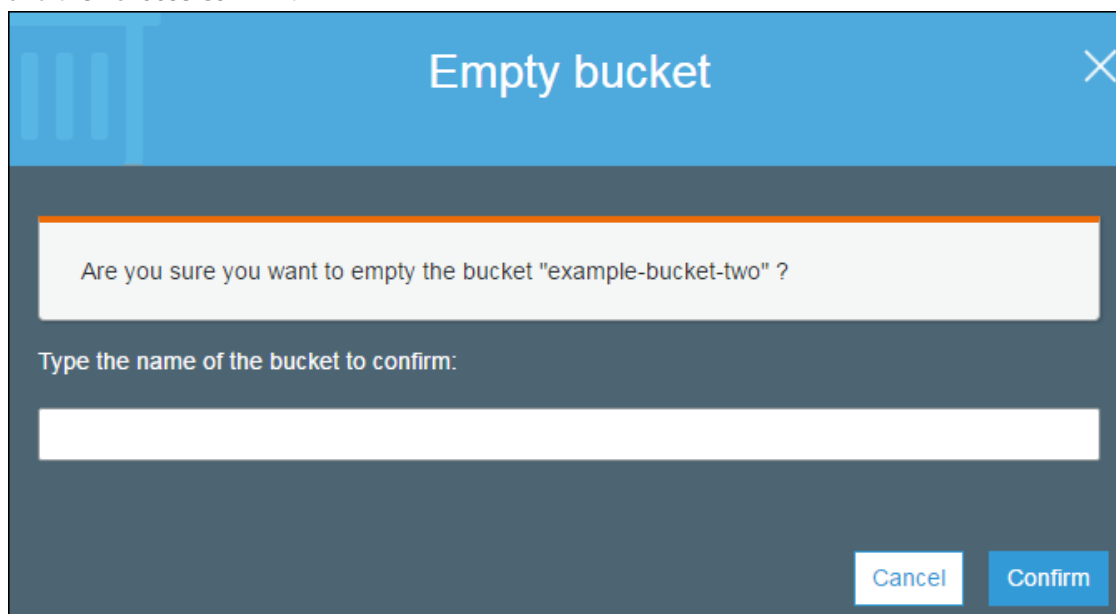
You can empty a bucket, which deletes all of the objects in the bucket without deleting the bucket. When you empty a bucket with versioning enabled, all versions of all the objects in the bucket are deleted. For more information, see [Managing Objects in a Versioning-Enabled Bucket](#) and [Deleting/Emptying a Bucket](#) in the *Amazon Simple Storage Service Developer Guide*.

To empty an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to empty and then choose **Empty**.



3. In the **Empty bucket** dialog box, type the name of the bucket you want to empty for confirmation and then choose **Confirm**.




How Do I View the Properties for an S3 Bucket?

This topic explains how to view the properties for an S3 bucket.

To view the properties for an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to view the properties for.



3. Choose **Properties**.




4. On the **Properties** page, you can configure the following properties for the bucket.
 - a. **Versioning** – Versioning enables you to keep multiple versions of an object in one bucket. By default, versioning is disabled for a new bucket. For information about enabling versioning, see [How Do I Enable or Suspend Versioning for an S3 Bucket? \(p. 12\)](#).
 - b. **Server access logging** – Server access logging provides detailed records for the requests that are made to your bucket. By default, Amazon S3 does not collect server access logs. For information about enabling server access logging, see [How Do I Enable Server Access Logging for an S3 Bucket? \(p. 16\)](#).
 - c. **Static website hosting** – You can host a static website on Amazon S3. To enable static website hosting, choose **Static website hosting** and then specify the settings you want to use. For more information, see [How Do I Configure an S3 Bucket for Static Website Hosting? \(p. 21\)](#).
 - d. **Object-level logging** – Object-level logging records object-level API activity by using CloudTrail data events. For information about enabling object-level logging, see [How Do I Enable Object-Level Logging for an S3 Bucket with AWS CloudTrail Data Events? \(p. 19\)](#).
 - e. **Tags** – With AWS cost allocation, you can use bucket tags to annotate billing for your use of a bucket. A tag is a key-value pair that represents a label that you assign to a bucket. To add tags, choose **Tags**, and then choose **Add tag**. For more information, see [Using Cost Allocation Tags for S3 Buckets](#) in the *Amazon Simple Storage Service Developer Guide*.
 - f. **Transfer acceleration** – Amazon S3 Transfer Acceleration enables fast, easy, and secure transfers of files over long distances between your client and an S3 bucket. For information about enabling transfer acceleration, see [How Do I Enable Transfer Acceleration for an S3 Bucket? \(p. 31\)](#).
 - g. **Events** – You can enable certain Amazon S3 bucket events to send a notification message to a destination whenever the events occur. To enable events, choose **Events** and then specify the settings you want to use. For more information, see [How Do I Enable and Configure Event Notifications for an S3 Bucket? \(p. 27\)](#).
 - h. **Requester Pays** – You can enable Requester Pays so that the requester (instead of the bucket owner) pays for requests and data transfers. For more information, see [Requester Pays Buckets](#) in the *Amazon Simple Storage Service Developer Guide*.

How Do I Enable or Suspend Versioning for an S3 Bucket?

Versioning enables you to keep multiple versions of an object in one bucket. This section describes how to enable object versioning on a bucket. For more information about versioning support in Amazon S3, see [Object Versioning](#) and [Using Versioning](#) in the *Amazon Simple Storage Service Developer Guide*.

To enable or disable versioning on an S3 bucket

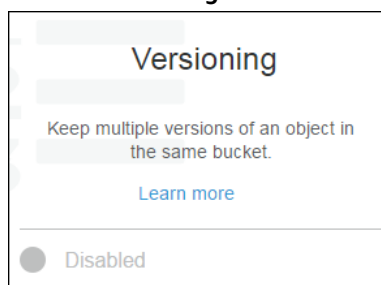
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to enable versioning for.




3. Choose **Properties**.



4. Choose **Versioning**.



5. Choose **Enable versioning** or **Suspend versioning**, and then choose **Save**.



How Do I Enable Default Encryption for an S3 Bucket?

Amazon S3 default encryption provides a way to set the default encryption behavior for an S3 bucket. You can set default encryption on a bucket so that all objects are encrypted when they are stored in the bucket. The objects are encrypted using server-side encryption with either Amazon S3-managed keys (SSE-S3) or AWS KMS-managed keys (SSE-KMS).

When you use server-side encryption, Amazon S3 encrypts an object before saving it to disk in its data centers and decrypts it when you download the objects. For more information about protecting data using server-side encryption and encryption key management, see [Protecting Data Using Server-Side Encryption](#) in the *Amazon Simple Storage Service Developer Guide*.


Default encryption works with all existing and new S3 buckets. Without default encryption, to encrypt all objects stored in a bucket, you must include encryption information with every object storage

request. You must also set up an S3 bucket policy to reject storage requests that don't include encryption information.

There are no new charges for using default encryption for S3 buckets. Requests to configure the default encryption feature incur standard Amazon S3 request charges. For information about pricing, see [Amazon S3 Pricing](#). For SSE-KMS encryption key storage, AWS Key Management Service (AWS KMS) charges apply and are listed at [AWS KMS Pricing](#).

To enable default encryption on an S3 bucket

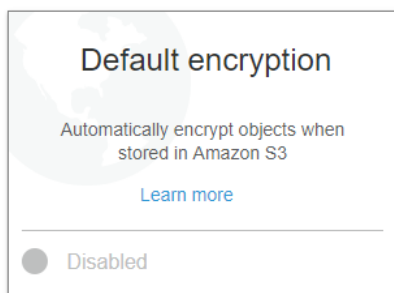
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want.



3. Choose **Properties**.



4. Choose **Default encryption**.



5. Choose **AES-256** or **AWS-KMS**.
 - a. To use keys that are managed by Amazon S3 for default encryption, choose **AES-256**. For more information about using Amazon S3 server-side encryption to encrypt your data, see [Protecting Data with Amazon S3-Managed Encryption Keys](#) in the *Amazon Simple Storage Service Developer Guide*.

Default encryption

This property does not affect existing objects in your bucket.

☐ None

☒ AES-256
Use Server-Side Encryption with Amazon S3-Managed Keys (SSE-S3)

☐ AWS-KMS
Use Server-Side Encryption with AWS KMS-Managed Keys (SSE-KMS)

Amazon S3 evaluates and applies bucket policies before applying bucket encryption settings. Even if you enable bucket encryption settings, your PUT requests without encryption information will be rejected if you have bucket policies to reject such PUT requests. Check your bucket policy and modify it if required.

[View bucket policy](#)

[Cancel](#) [Save](#)

Important

You might need to update your bucket policy when enabling default encryption. For more information, see [Moving to Default Encryption from Using Bucket Policies for Encryption Enforcement](#) in the *Amazon Simple Storage Service Developer Guide*.

- b. To use keys that are managed by AWS KMS for default encryption, choose **AWS-KMS**, and then choose a master key from the list of the AWS KMS master keys that you have created. Type the Amazon Resource Name (ARN) of the AWS KMS key to use. You can find the ARN for your AWS KMS key in the IAM console, under **Encryption keys**. Or, you can choose a key name from the drop-down list.

Default encryption

This property does not affect existing objects in your bucket.

☐ None

☐ AES-256
Use Server-Side Encryption with Amazon S3-Managed Keys (SSE-S3)

☒ AWS-KMS
Use Server-Side Encryption with AWS KMS-Managed Keys (SSE-KMS)

Custom KMS ARN

Type to search

aws/s3

ca-key

Custom KMS ARN

[View bucket policy](#)

Cancel Save

Important

If you use the AWS KMS option for your default encryption configuration, you are subject to the RPS (requests per second) limits of AWS KMS. For more information about AWS KMS limits and how to request a limit increase, see [AWS KMS limits](#).

For more information about creating an AWS KMS key, see [Creating Keys](#) in the *AWS Key Management Service Developer Guide*. For more information about using AWS KMS with Amazon S3, see [Protecting Data with AWS KMS–Managed Keys](#) in the *Amazon Simple Storage Service Developer Guide*.

6. Choose **Save**.

More Info

- [Amazon S3 Default Encryption for S3 Buckets](#) in the *Amazon Simple Storage Service Developer Guide*
- [How Do I Add Encryption to an S3 Object?](#) (p. 60)

How Do I Enable Server Access Logging for an S3 Bucket?

This topic describes how to enable server access logging for an Amazon S3 bucket using the AWS Management Console. For information about how to enable logging programmatically and details about how logs are delivered, see [Server Access Logging](#) in the *Amazon Simple Storage Service Developer Guide*.

By default, Amazon Simple Storage Service (Amazon S3) doesn't collect server access logs. When you enable logging, Amazon S3 delivers access logs for a source bucket to a target bucket that you choose. The target bucket must be in the same AWS Region as the source bucket.

Server access logging provides detailed records for the requests that are made to an S3 bucket. Server access logs are useful for many applications. For example, access log information can be useful in security and access audits. It can also help you learn about your customer base and understand your Amazon S3 bill.

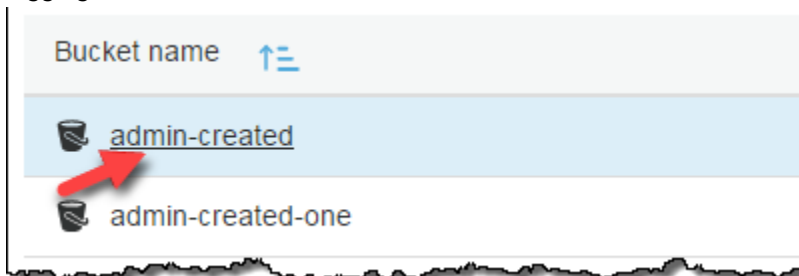
An access log record contains details about the requests that are made to a bucket. This information can include the request type, the resources that are specified in the request, and the time and date that the request was processed. For more information, see [Server Access Log Format](#) in the *Amazon Simple Storage Service Developer Guide*.

Important

There is no extra charge for enabling server access logging on an Amazon S3 bucket. However, any log files that the system delivers to you will accrue the usual charges for storage. (You can delete the log files at any time.) We do not assess data transfer charges for log file delivery, but we do charge the normal data transfer rate for accessing the log files.

To enable server access logging for an S3 bucket


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to enable server access logging for.



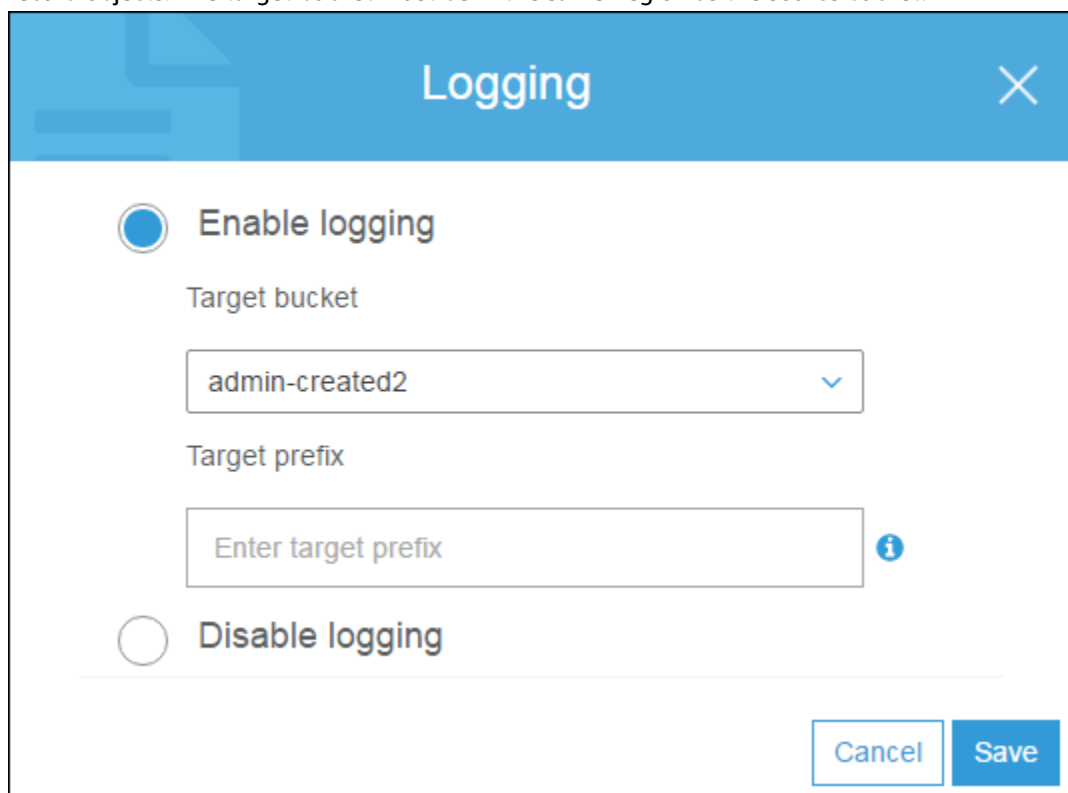
3. Choose **Properties**.



4. Choose **Server access logging**.



5. Choose **Enable Logging**. For **Target**, choose the name of the bucket that you want to receive the log record objects. The target bucket must be in the same Region as the source bucket.



6. (Optional) For **Target prefix**, type a key name prefix for log objects, so that all of the log object names begin with the same string.
7. Choose **Save**.

You can view the logs in the target bucket. If you specified a prefix, the prefix shows as a folder in the target bucket in the console. After you enable server access logging, it might take a few hours before the logs are delivered to the target bucket. For more information about how and when logs are delivered, see [Server Access Logging](#) in the *Amazon Simple Storage Service Developer Guide*.

More Info

[How Do I View the Properties for an S3 Bucket? \(p. 11\)](#)

How Do I Enable Object-Level Logging for an S3 Bucket with AWS CloudTrail Data Events?

This section describes how to enable an AWS CloudTrail trail to log data events for objects in an S3 bucket by using the Amazon S3 console. CloudTrail supports logging Amazon S3 object-level API operations such as `GetObject`, `DeleteObject`, and `PutObject`. These events are called data events. By default, CloudTrail trails don't log data events, but you can configure trails to log data events for S3 buckets that you specify, or to log data events for all the Amazon S3 buckets in your AWS account.

Important


Additional charges apply for data events. For more information, see [AWS CloudTrail Pricing](#).

To configure a trail to log data events for an S3 bucket, you can use either the AWS CloudTrail console or the Amazon S3 console. If you are configuring a trail to log data events for all the Amazon S3 buckets in your AWS account, it's easier to use the CloudTrail console. For information about using the CloudTrail console to configure a trail to log S3 data events, see [Data Events](#) in the *AWS CloudTrail User Guide*.

The following procedure shows how to use the Amazon S3 console to enable a CloudTrail trail to log data events for an S3 bucket.

To enable CloudTrail data events logging for objects in an S3 bucket

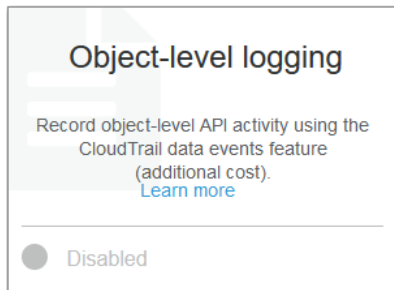
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want.



3. Choose **Properties**.



4. Choose **Object-level logging**.



5. Choose an existing CloudTrail trail in the drop-down menu. The trail you select must be in the same AWS Region as your bucket, so the drop-down list contains only trails that are in the same Region as the bucket or trails that were created for all Regions.

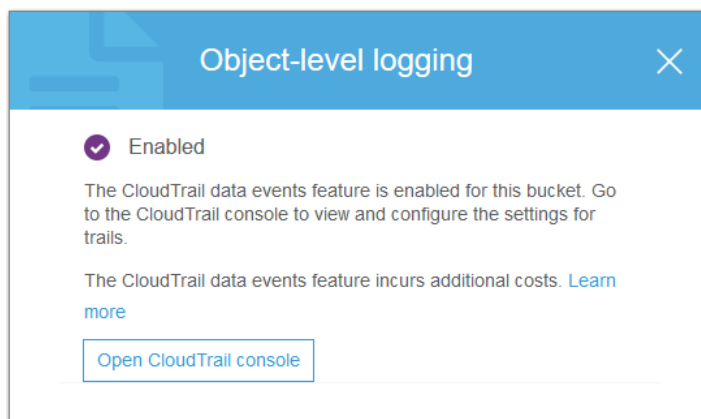
If you need to create a trail, choose the **CloudTrail console** link to go to the CloudTrail console. For information about how to create trails in the CloudTrail console, see [Creating a Trail with the Console](#) in the *AWS CloudTrail User Guide*.

The screenshot shows the 'Object-level logging' console window. At the top, there's a blue header with the title 'Object-level logging' and a close button. Below the header, a toggle switch is in the 'Disabled' position. Underneath, it says 'CloudTrail trail for recording activity' with an information icon and a refresh icon. A dropdown menu shows 'example-trail'. A note states: 'You can only choose a trail in the same region as the bucket. To create a trail, go to the [CloudTrail console](#)'. Below this, the 'Events' section has two checkboxes: 'Read' and 'Write', both of which are currently unchecked. At the bottom right, there are 'Cancel' and 'Create' buttons.

6. Under **Events**, select **Read** to specify that you want CloudTrail to log Amazon S3 read APIs such as `GetObject`. Select **Write** to log Amazon S3 write APIs such as `PutObject`. Select both **Read** and **Write** to log both read and write object APIs. For a list of supported data events that CloudTrail logs for Amazon S3 objects, see [Amazon S3 Object-Level Actions Tracked by CloudTrail Logging](#) in the *Amazon Simple Storage Service Developer Guide*.

This screenshot is similar to the previous one, but the 'Events' section is highlighted with a red rectangle. In this state, both the 'Read' and 'Write' checkboxes are checked. The 'Create' button at the bottom right is now highlighted in blue, indicating it is the active or recommended action.

7. Choose **Create** to enable object-level logging for the bucket.



To disable object-level logging for the bucket, you must go to the CloudTrail console and remove the bucket name from the trail's **Data events**.

Note

If you use the CloudTrail console or the Amazon S3 console to configure a trail to log data events for an S3 bucket, the Amazon S3 console shows that object-level logging is enabled for the bucket.

For information about enabling object-level logging when you create an S3 bucket, see [How Do I Create an S3 Bucket? \(p. 3\)](#).

More Info

- [How Do I View the Properties for an S3 Bucket? \(p. 11\)](#)
- [Logging Amazon S3 API Calls By Using AWS CloudTrail](#) in the *Amazon Simple Storage Service Developer Guide*
- [Working with CloudTrail Log Files](#) in the *AWS CloudTrail User Guide*


How Do I Configure an S3 Bucket for Static Website Hosting?

You can host a static website on Amazon S3. On a static website, individual web pages include static content and they might also contain client-side scripts. By contrast, a dynamic website relies on server-side processing, including server-side scripts such as PHP, JSP, or ASP.NET. Amazon S3 does not support server-side scripting.

The following is a quick procedure to configure an Amazon S3 bucket for static website hosting in the S3 console. If you're looking for more in-depth information, as well as walkthroughs on using a custom domain name for your static website or speeding up your website, see [Hosting a Static Website on Amazon S3](#) in the *Amazon Simple Storage Service Developer Guide*.

To configure an S3 bucket for static website hosting

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to enable static website hosting for.



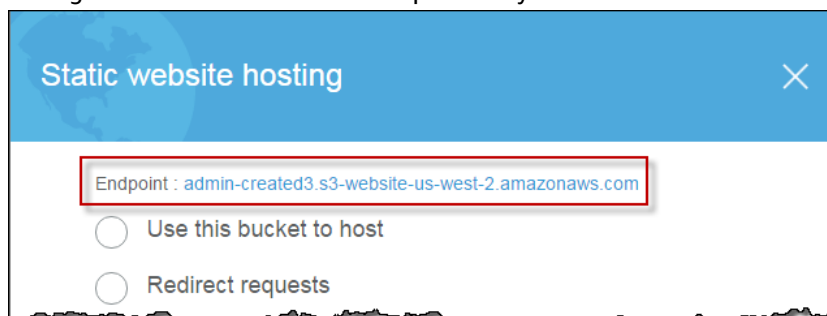
3. Choose **Properties**.



4. Choose **Static website hosting**.

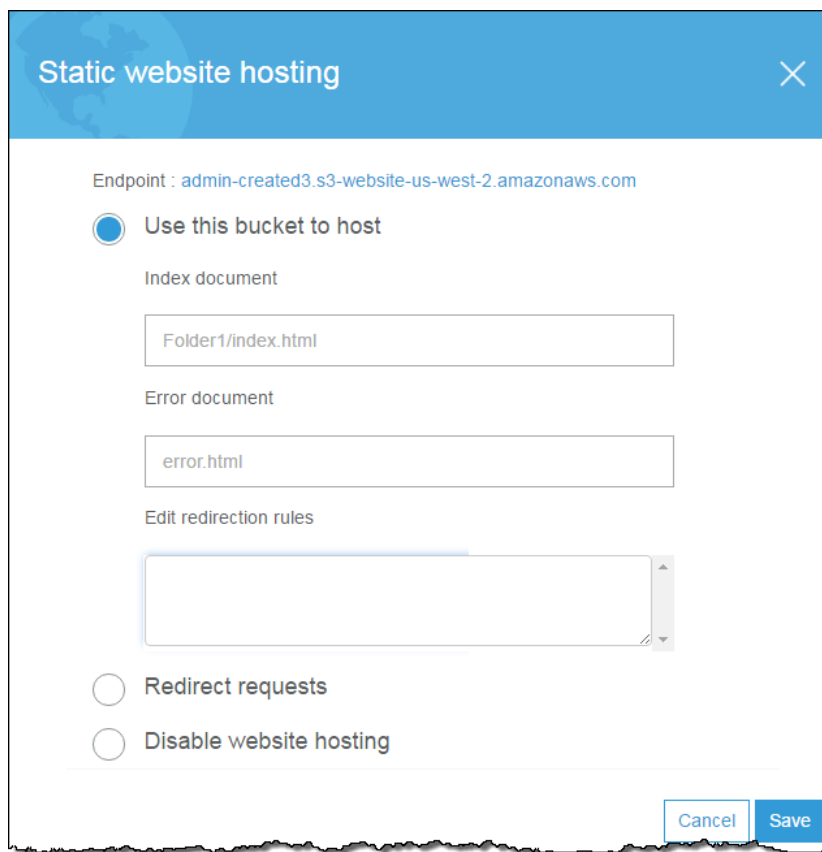


After you enable your bucket for static website hosting, web browsers can access all of your content through the Amazon S3 website endpoint for your bucket.



5. Choose **Use this bucket to host**.

- For **Index Document**, type the name of the index document, which is typically named `index.html`. When you configure a bucket for website hosting, you must specify an index document. Amazon S3 returns this index document when requests are made to the root domain or any of the subfolders. For more information, see [Configure a Bucket for Website Hosting](#) in the *Amazon Simple Storage Service Developer Guide*.
- (Optional) For 4XX class errors, you can optionally provide your own custom error document that provides additional guidance for your users. For **Error Document**, type the name of the file that contains the custom error document. If an error occurs, Amazon S3 returns an HTML error document. For more information, see [Custom Error Document Support](#) in the *Amazon Simple Storage Service Developer Guide*.
- (Optional) If you want to specify advanced redirection rules, in the **Edit redirection rules** text area, use XML to describe the rules. For example, you can conditionally route requests according to specific object key names or prefixes in the request. For more information, see [Configure a Bucket for Website Hosting](#) in the *Amazon Simple Storage Service Developer Guide*.



Static website hosting

Endpoint : admin-created3.s3-website-us-west-2.amazonaws.com

☒ Use this bucket to host

Index document

Folder1/index.html

Error document

error.html

Edit redirection rules

☐ Redirect requests

☐ Disable website hosting

Cancel Save

6. Choose **Save**.
7. Add a bucket policy to the website bucket that grants everyone access to the objects in the bucket. When you configure a bucket as a website, you must make the objects that you want to serve publicly readable. To do so, you write a bucket policy that grants everyone `s3:GetObject` permission. The following example bucket policy grants everyone access to the objects in the `example-bucket` bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicReadGetObject",
      "Effect": "Allow",
      "Principal": "*",
      "Action": [
        "s3:GetObject"
      ],
      "Resource": [
        "arn:aws:s3:::example-bucket/*"
      ]
    }
  ]
}
```

For information about adding a bucket policy, see [How Do I Add an S3 Bucket Policy? \(p. 123\)](#). For more information about website permissions, see [Permissions Required for Website](#) in the *Amazon Simple Storage Service Developer Guide*.

Note


If you choose **Disable website hosting**, Amazon S3 removes the website configuration from the bucket, so that the bucket is no longer accessible from the website endpoint. However, the bucket is still available at the REST endpoint. For a list of Amazon S3 endpoints, see [Amazon S3 Regions and Endpoints](#) in the *Amazon Web Services General Reference*.

How Do I Redirect Requests to an S3 Bucket Hosted Website to Another Host?

You can redirect all requests to your S3 bucket hosted static website to another host.

To redirect all requests to an S3 bucket's website endpoint to another host

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to redirect all requests from.



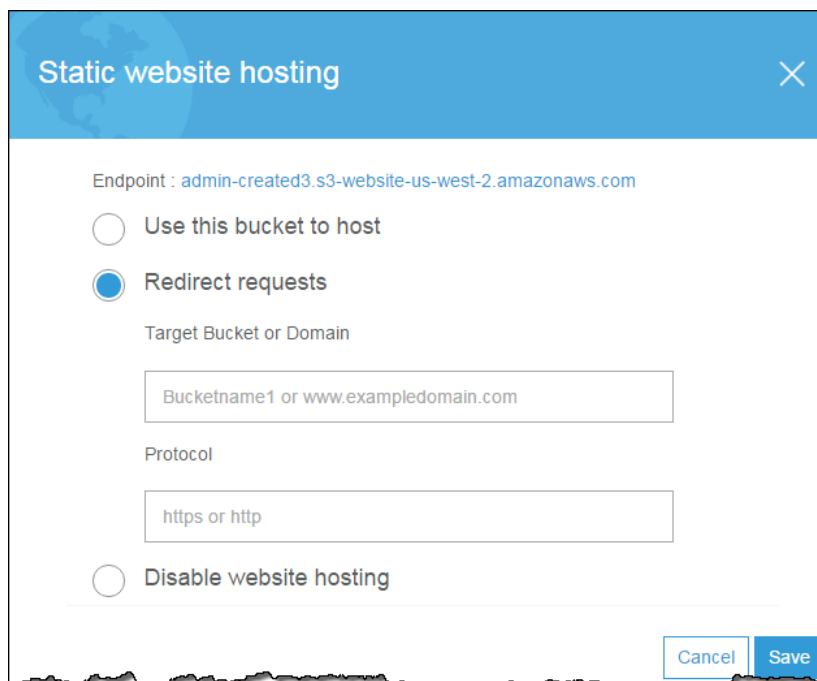
3. Choose **Properties**.



4. Choose **Static website hosting**.



5. Choose **Redirect requests**.



- a. For **Target bucket or domain**, type the name of the bucket or the domain name where you want requests to be redirected. To redirect requests to another bucket, type the name of the target bucket. For example, if you are redirecting to a root domain address, you would type **www.example.com**. For more information, see [Configure a Bucket for Website Hosting](#) in the *Amazon Simple Storage Service Developer Guide*.
 - b. For **Protocol**, type the protocol (http, https) for the redirected requests. If no protocol is specified, the protocol of the original request is used. If you redirect all requests, any request made to the bucket's website endpoint will be redirected to the specified host name.
6. Choose **Save**.

Advanced Settings for S3 Bucket Properties

This section describes how to configure advanced S3 bucket property settings for cross-region replication, event notification, and transfer acceleration.

Topics

- [How Do I Set Up a Destination to Receive Event Notifications?](#) (p. 25)
- [How Do I Enable and Configure Event Notifications for an S3 Bucket?](#) (p. 27)
- [How Do I Enable Transfer Acceleration for an S3 Bucket?](#) (p. 31)

How Do I Set Up a Destination to Receive Event Notifications?

Before you can enable event notifications for your bucket you must set up one of the following destination types:

An Amazon SNS topic

Amazon Simple Notification Service (Amazon SNS) is a web service that coordinates and manages the delivery or sending of messages to subscribing endpoints or clients. You can use the Amazon SNS console to create an Amazon SNS topic that your notifications can be sent to. The Amazon SNS topic must be in the same region as your Amazon S3 bucket. For information about creating an Amazon SNS topic, see [Getting Started](#) in the *Amazon Simple Notification Service Developer Guide*.

Before you can use the Amazon SNS topic that you create as an event notification destination, you need the following:

- The Amazon Resource Name (ARN) for the Amazon SNS topic
- A valid Amazon SNS topic subscription (the topic subscribers are notified when a message is published to your Amazon SNS topic)
- A permissions policy that you set up in the Amazon SNS console (as shown in the following example)

```
{
  "Version": "2012-10-17",
  "Id": "__example_policy_ID",
  "Statement": [
    {
      "Sid": "example-statement-ID",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "SNS:Publish",
      "Resource": "arn:aws:sns:region:account-number:topic-name",
      "Condition": {
        "ArnEquals": {
          "aws:SourceArn": "arn:aws:s3::bucket-name"
        }
      }
    }
  ]
}
```

An Amazon SQS queue

You can use the Amazon SQS console to create an Amazon SQS queue that your notifications can be sent to. The Amazon SQS queue must be in the same region as your Amazon S3 bucket. For information about creating an Amazon SQS queue, see [Getting Started with Amazon SQS](#) in the *Amazon Simple Queue Service Developer Guide*.

Before you can use the Amazon SQS queue as an event notification destination, you need the following:

- The Amazon Resource Name (ARN) for the Amazon SQS topic
- A permissions policy that you set up in the Amazon SQS console (as shown in the following example)

```
{
  "Version": "2012-10-17",
  "Id": "__example_policy_ID",
  "Statement": [
    {
      "Sid": "example-statement-ID",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "SQS:*",
      "Resource": "arn:aws:sqs:region:account-number:queue-name",
      "Condition": {
        "ArnEquals": {
          "aws:SourceArn": "arn:aws:s3::bucket-name"
        }
      }
    }
  ]
}
```

```
}  
  }  
    }  
      }  
        ]  
      }  
    }
```

A Lambda function

You can use the AWS Lambda console to create a Lambda function. The Lambda function must be in the same region as your S3 bucket. For information about creating a Lambda function, see the [AWS Lambda Developer Guide](#).

Before you can use the Lambda function as an event notification destination, you must have the name or the ARN of a Lambda function to set up the Lambda function as a event notification destination.

For information about using Lambda with Amazon S3, see [Using AWS Lambda: with Amazon S3](#) in the *AWS Lambda Developer Guide*.

How Do I Enable and Configure Event Notifications for an S3 Bucket?

You can enable certain Amazon S3 bucket events to send a notification message to a destination whenever the events occur. This section explains how to use the Amazon S3 console to enable event notifications. For information about using event notifications with the the AWS SDKs, and the Amazon S3 REST APIs, see [Configuring Notifications for Amazon S3 Events](#) in the *Amazon Simple Storage Service Developer Guide*.

Topics

- [Amazon S3 Event Notification Types and Destinations \(p. 27\)](#)
- [Enabling and Configuring Event Notifications \(p. 28\)](#)
- [More Info \(p. 31\)](#)

Amazon S3 Event Notification Types and Destinations

When configuring event notifications for a bucket you must specify the type of events you want to be notified of and the destination where you want the notifications sent.

Amazon S3 can send notifications for the following types of events:

- **An object created event** – You choose **ObjectCreated (All)** when configuring your events in the console to enable notifications for anytime an object is created in your bucket. Or, you can select one or more of the specific object-creation actions to trigger event notifications. These actions are **Put**, **Post**, **Copy**, and **CompleteMultiPartUpload**.
- **An object delete event** – You select **ObjectDelete (All)** when configuring your events in the console to enable notification for anytime an object is deleted. Or, you can select **Delete** to trigger event notifications when an unversioned object is deleted or a versioned object is permanently deleted. You select **Delete Marker Created** to trigger event notifications when a delete marker is created for a versioned object.
- **Restore object events** – When configuring events in the console to enable notifications for the restoration of objects stored in the GLACIER storage class. Select **Restore from Glacier initiated** to be

notified of when the restore is initiated. Select **Restore from Glacier completed** to be notified when restoration of an object is complete.

- **A Reduced Redundancy Storage (RRS) object lost event** – You select **RRSObjectLost** to be notified when Amazon S3 detects that an object of the RRS storage class has been lost.

Event notification messages can be sent to the following types of destinations:


- **An Amazon Simple Notification Service (Amazon SNS) topic** – A web service that coordinates and manages the delivery or sending of messages to subscribing endpoints or clients.
- **An Amazon Simple Queue Service (Amazon SQS) queue** – Offers reliable and scalable hosted queues for storing messages as they travel between computer.
- **A Lambda function** – AWS Lambda is a compute service where you can upload your code and the service can run the code on your behalf using the AWS infrastructure. You package up and upload your custom code to AWS Lambda when you create a Lambda function

Enabling and Configuring Event Notifications

Before you can enable event notifications for your bucket you must set up one of these destination types. For more information, see [How Do I Set Up a Destination to Receive Event Notifications?](#) (p. 25).

To enable and configure event notifications for an S3 bucket


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to enable events for.



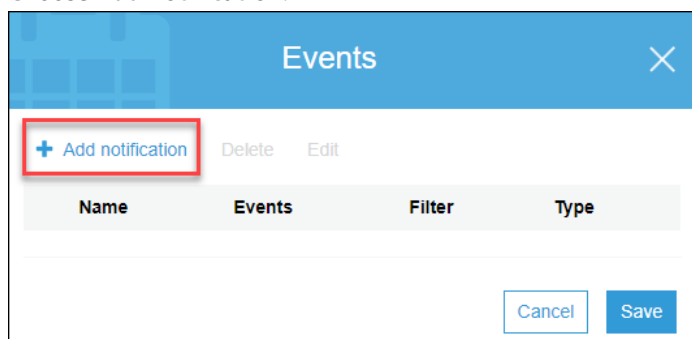
3. Choose **Properties**.



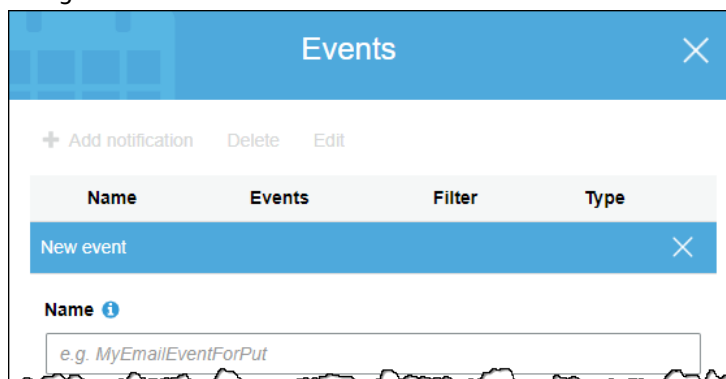
4. Under Advanced settings, choose **Events**.



5. Choose **Add notification**.



6. In **Name**, type a descriptive name for your event configuration. If you do not enter a name, a GUID is autogenerated and used for the name.



7. Under **Events**, select one or more of the type of event occurrences that you want to receive notifications for. When the event occurs a notification is sent to a destination that you choose in **Step 9**. For a description of the event types, see [Amazon S3 Event Notification Types and Destinations](#) (p. 27).



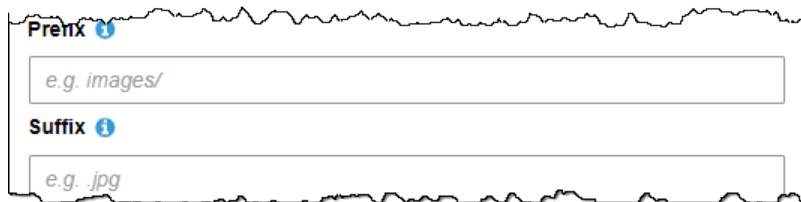
For information about deleting versioned objects, see [Deleting Object Versions](#). For information about object versioning, see [Object Versioning](#) and [Using Versioning](#).

Note

When you delete the last object from a folder Amazon S3 can generate an object creation event. The Amazon S3 console displays a folder under the following circumstances: 1) when a zero byte object has a trailing slash (/) in its name (in this case there is an actual Amazon S3 object of 0 bytes that represents a folder), and 2) if the object has a slash (/) within its name (in this case there isn't an actual object representing the folder). When there are multiple objects with the same prefix with a trailing slash (/) as part of their names, those objects are shown as being part of a folder. The name of the folder is formed from the characters preceding the trailing slash (/). When you delete all the objects listed under that folder, there is no actual object available to represent the empty folder. Under such circumstance the Amazon S3 console creates a zero byte object to represent that folder. If

you enabled event notification for creation of objects, the zero byte object creation action that is taken by the console will trigger an object creation event.

8. Type an object name **Prefix** and/or a **Suffix** to filter the event notifications by the prefix and/or suffix. For example, you can set up a filter so that you are sent a notification only when files are added to an image folder (for example, objects with the name prefix `images/`). For more information, see [Configuring Notifications with Object Key Name Filtering](#).



The screenshot shows two input fields. The first is labeled "Prefix" with an information icon and contains the placeholder text "e.g. images/". The second is labeled "Suffix" with an information icon and contains the placeholder text "e.g. .jpg".

9. Select the type of destination to have the event notifications sent to. For a description of the destinations, see [Amazon S3 Event Notification Types and Destinations](#) (p. 27).



The screenshot shows a dropdown menu labeled "Send to" with an information icon. The menu is open, showing a list of notification destinations. Below the dropdown is a "Cancel" button and a "Save" button.

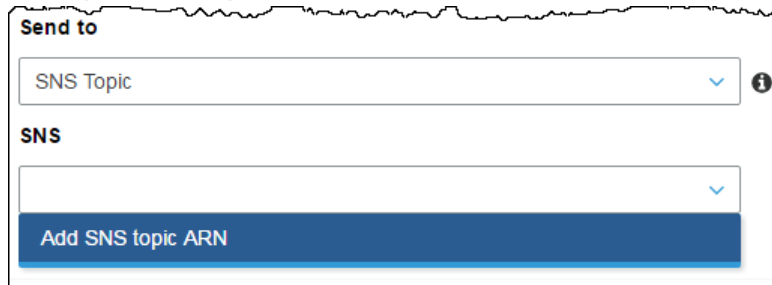
- a. If you select the **SNS Topic** destination type.

- i. In the **SNS topic** box, type the name or select from the menu, the Amazon SNS topic that will receive notifications from Amazon S3. For information about the Amazon SNS topic format, see [SNS FAQ](#).



The screenshot shows the "Send to" dropdown menu with "SNS Topic" selected. Below it is a text input field labeled "SNS" with a dropdown arrow.

- ii. (Optional) You can also select **Add SNS topic ARN** from the menu and type the **ARN** of the SNS topic in **SNS topic ARN**.



The screenshot shows the "Send to" dropdown menu with "Add SNS topic ARN" selected. Below it is a text input field labeled "SNS" with a dropdown arrow.

- b. If you select the **SQS queue** destination type, do the following:
 - i. In **SQS queue**, type or choose a name from the menu of the Amazon SQS queue that you want to receive notifications from Amazon S3. For information about Amazon SQS, see [What is Amazon Simple Queue Service?](#) in the *Amazon Simple Queue Service Developer Guide*.
 - ii. (Optional) You can also select **Add SQS topic ARN** from the menu and type the ARN of the SQS queue in **SQS queue ARN**.
- c. If you select the **Lambda Function** destination type, do the following:
 - i. In **Lambda Function**, type or choose the name of the Lambda function that you want to receive notifications from Amazon S3.
 - ii. If you don't have any Lambda functions in the region that contains your bucket, you'll be prompted to enter a Lambda function ARN. In **Lambda Function ARN**, type the ARN of the Lambda function that you want to receive notifications from Amazon S3.
 - iii. (Optional) You can also choose **Add Lambda function ARN** from the menu and type the ARN of the Lambda function in **Lambda function ARN**.

For information about using Lambda with Amazon S3, see [Using AWS Lambda: with Amazon S3](#) in the *AWS Lambda Developer Guide*.

10. Choose **Save**. Amazon S3 will send a test message to the event notification destination.

More Info


- [How Do I Restore an S3 Object That Has Been Archived to Glacier? \(p. 47\)](#)

How Do I Enable Transfer Acceleration for an S3 Bucket?

Amazon Simple Storage Service (Amazon S3) transfer acceleration enables fast, easy, and secure transfers of files between your client and an S3 bucket over long distances. This topic describes how to enable Amazon S3 transfer acceleration for a bucket. For more information, see [Amazon S3 Transfer Acceleration](#) in the *Amazon Simple Storage Service Developer Guide*.

To enable transfer acceleration for an S3 bucket

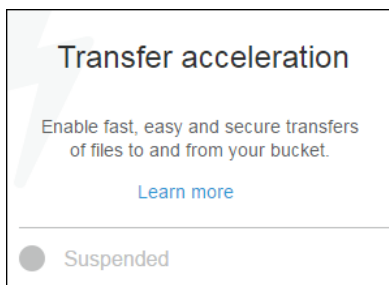
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to enable transfer acceleration for.



3. Choose **Properties**.

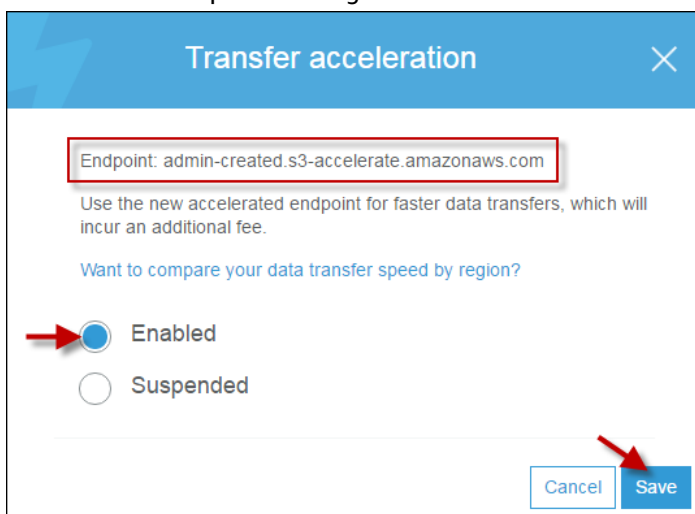


4. Choose **Transfer acceleration**.



5. Choose **Enabled**, and then choose **Save**.

Endpoint displays the endpoint domain name that you use to access accelerated data transfers to and from the bucket that is enabled for transfer acceleration. If you suspend transfer acceleration, the accelerate endpoint no longer works.



6. (Optional) If you want to run the [Amazon S3 Transfer Acceleration Speed Comparison tool](#), which compares accelerated and non-accelerated upload speeds starting with the Region in which the transfer acceleration bucket is enabled, choose the **Want to compare your data transfer speed by region?** option. The Speed Comparison tool uses multipart uploads to transfer a file from your browser to various AWS Regions with and without using Amazon S3 transfer acceleration.

More Info

[How Do I View the Properties for an S3 Bucket? \(p. 11\)](#)

Uploading, Downloading, and Managing Objects

Amazon S3 is cloud storage for the Internet. To upload your data (photos, videos, documents etc.), you first create a bucket in one of the AWS Regions. You can then upload an unlimited number of data objects to the bucket.

The data that you store in Amazon S3 consists of objects. Every object resides within a bucket that you create in a specific AWS Region. Every object that you store in Amazon S3 resides in a bucket.

Objects stored in a region never leave the region unless you explicitly transfer them to another region. For example, objects stored in the EU (Ireland) region never leave it. The objects stored in an AWS region physically remain in that region. Amazon S3 does not keep copies of objects or move them to any other region. However, you can access the objects from anywhere, as long as you have necessary permissions to do so.

Before you can upload an object into Amazon S3, you must have write permissions to a bucket.

Objects can be any file type: images, backups, data, movies, etc. The maximum size of file you can upload by using the Amazon S3 console is 78GB. You can have an unlimited number of objects in a bucket.

The following topics explain how to use the Amazon S3 console to upload, delete, and manage objects.

Topics

- [How Do I Upload Files and Folders to an S3 Bucket? \(p. 33\)](#)
- [How Do I Download an Object from an S3 Bucket? \(p. 42\)](#)
- [How Do I Delete Objects from an S3 Bucket? \(p. 44\)](#)
- [How Do I Undelete a Deleted S3 Object? \(p. 46\)](#)
- [How Do I Restore an S3 Object That Has Been Archived to Glacier? \(p. 47\)](#)
- [How Do I Lock an S3 Object? \(p. 52\)](#)
- [How Do I See an Overview of an Object? \(p. 55\)](#)
- [How Do I See the Versions of an S3 Object? \(p. 57\)](#)
- [How Do I View the Properties of an Object? \(p. 58\)](#)
- [How Do I Add Encryption to an S3 Object? \(p. 60\)](#)
- [How Do I Add Metadata to an S3 Object? \(p. 63\)](#)
- [How Do I Add Tags to an S3 Object? \(p. 68\)](#)
- [How do I use folders in an S3 Bucket? \(p. 71\)](#)

How Do I Upload Files and Folders to an S3 Bucket?

This topic explains how to use the AWS Management Console to upload one or more files or entire folders to an Amazon S3 bucket. Before you can upload files and folders to an Amazon S3 bucket, you need write permissions for the bucket. For more information about access permissions, see [Setting Bucket and Object Access Permissions \(p. 110\)](#). For information about uploading files programmatically, see [Uploading Objects](#) in the *Amazon Simple Storage Service Developer Guide*.

When you upload a file to Amazon S3, it is stored as an S3 object. Objects consist of the file data and metadata that describes the object. You can have an unlimited number of objects in a bucket.

You can upload any file type—images, backups, data, movies, etc.—into an S3 bucket. The maximum size of a file that you can upload by using the Amazon S3 console is 78 GB.

You can upload files by dragging and dropping or by pointing and clicking. To upload folders, you *must* drag and drop them. Drag and drop functionality is supported *only* for the Chrome and Firefox browsers. For information about which Chrome and Firefox browser versions are supported, see [Which Browsers are Supported for Use with the AWS Management Console?](#)

When you upload a folder, Amazon S3 uploads all of the files and subfolders from the specified folder to your bucket. It then assigns an object key name that is a combination of the uploaded file name and the folder name. For example, if you upload a folder called `/images` that contains two files, `sample1.jpg` and `sample2.jpg`, Amazon S3 uploads the files and then assigns the corresponding key names, `images/sample1.jpg` and `images/sample2.jpg`. The key names include the folder name as a prefix. The Amazon S3 console displays only the part of the key name that follows the last `/`. For example, within an `images` folder the `images/sample1.jpg` and `images/sample2.jpg` objects are displayed as `sample1.jpg` and a `sample2.jpg`.

If you upload individual files and you have a folder open in the Amazon S3 console, when Amazon S3 uploads the files, it includes the name of the open folder as the prefix of the key names. For example, if you have a folder named `backup` open in the Amazon S3 console and you upload a file named `sample1.jpg`, the key name is `backup/sample1.jpg`. However, the object is displayed in the console as `sample1.jpg` in the `backup` folder.

If you upload individual files and you do not have a folder open in the Amazon S3 console, when Amazon S3 uploads the files, it assigns only the file name as the key name. For example, if you upload a file named `sample1.jpg`, the key name is `sample1.jpg`. For more information on key names, see [Object Key and Metadata](#) in the *Amazon Simple Storage Service Developer Guide*.

If you upload an object with a key name that already exists in a versioning-enabled bucket, Amazon S3 creates another version of the object instead of replacing the existing object. For more information about versioning, see [How Do I Enable or Suspend Versioning for an S3 Bucket?](#) (p. 12).

Topics


- [Uploading Files and Folders by Using Drag and Drop](#) (p. 34)
- [Uploading Files by Pointing and Clicking](#) (p. 39)
- [More Info](#) (p. 41)

Uploading Files and Folders by Using Drag and Drop

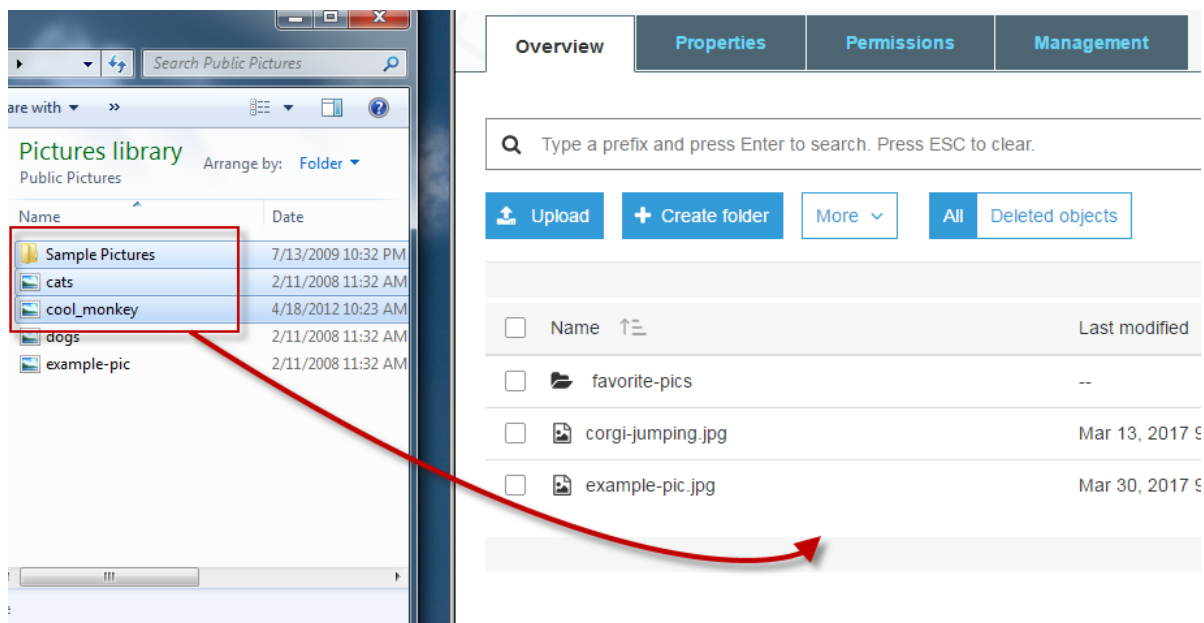
If you are using the Chrome or Firefox browsers, you can choose the folders and files to upload, and then drag and drop them into the destination bucket. Dragging and dropping is the *only* way that you can upload folders.

To upload folders and files to an S3 bucket by using drag and drop

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to upload your folders or files to.

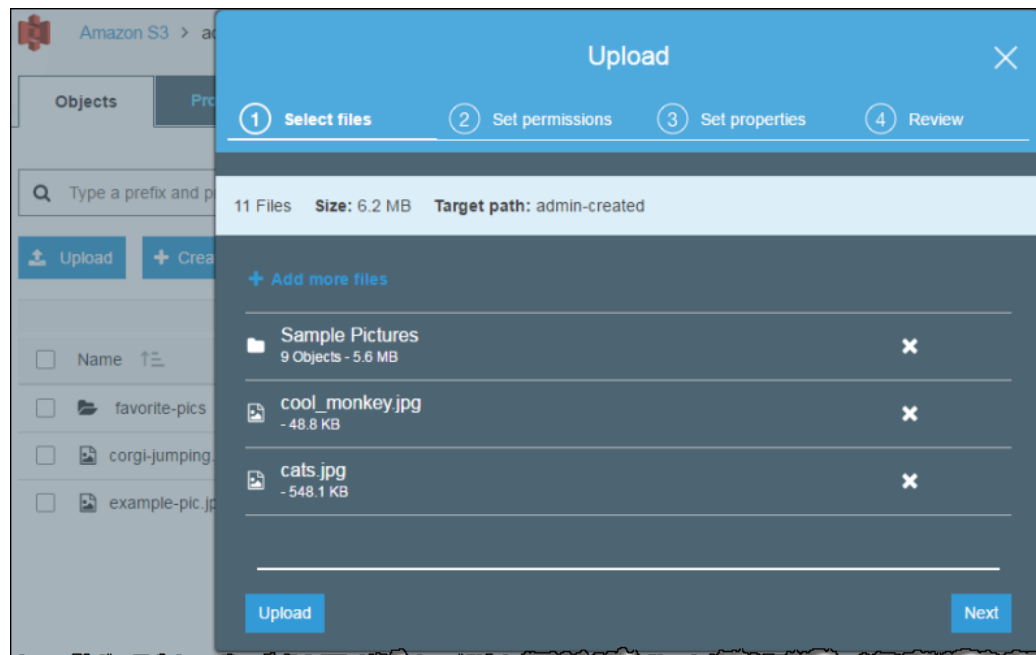


3. In a window other than the console window, select the files and folders that you want to upload. Then drag and drop your selections into the console window that lists the objects in the destination bucket.



The files you chose are listed in the **Upload** dialog box.

4. In the **Upload** dialog box, do one of the following:
 - a. Drag and drop more files and folders to the console window that displays the **Upload** dialog box. To add more files, you can also choose **Add more files**. This option works *only* for files, not folders.
 - b. To immediately upload the listed files and folders without granting or removing permissions for specific users or setting public permissions for all of the files that you're uploading, choose **Upload**. For information about object access permissions, see [How Do I Set Permissions on an Object?](#) (p. 116).
 - c. To set permissions or properties for the files that you are uploading, choose **Next**.



5. On the **Set Permissions** page, under **Manage users** you can change the permissions for the AWS account owner. The *owner* refers to the AWS account root user, and not an AWS Identity and Access Management (IAM) user. For more information about the root user, see [The AWS Account Root User](#).

Choose **Add account** to grant access to another AWS account. For more information about granting permissions to another AWS account, see [How Do I Set ACL Bucket Permissions?](#) (p. 119).

Under **Manage public permissions** you can grant read access to your objects to the general public (everyone in the world), for all of the files that you're uploading. Granting public read access is applicable to a small subset of use cases such as when buckets are used for websites. We recommend that you do not change the default setting of **Do not grant public read access to this object(s)**. You can always make changes to object permissions after you upload the object. For information about object access permissions, see [How Do I Set Permissions on an Object?](#) (p. 116).

When you're done configuring permissions, choose **Next**.

The screenshot shows the 'Upload' console window with the 'Set permissions' step selected. The progress bar at the top indicates four steps: 1. Select files, 2. Set permissions (active), 3. Set properties, and 4. Review. Below the progress bar, it shows '1 Files', 'Size: 1.3 MB', and 'Target path: admin-created'. The 'Manage users' section includes a table with columns for User ID, Objects, and Object permissions. A user 'johndoe(Owner)' is listed with 'Read' and 'Write' permissions checked. Below this is the 'Access for other AWS account' section with an 'Add account' button. The 'Manage public permissions' section has a dropdown menu set to 'Do not grant public read access to this object(s) (Recommended)'. At the bottom are 'Upload', 'Previous', and 'Next' buttons.

User ID	Objects	Object permissions
johndoe(Owner)	<input checked="" type="checkbox"/> Read <input checked="" type="checkbox"/> Write	<input checked="" type="checkbox"/> Read <input checked="" type="checkbox"/> Write

Access for other AWS account [+ Add account](#)

Account Objects Object permissions

Manage public permissions

Do not grant public read access to this object(s) (Recommended)

[Upload](#) [Previous](#) [Next](#)

6. On the **Set Properties** page, choose the storage class and encryption method to use for the files that you are uploading. You can also add or modify metadata.
 - a. Choose a storage class for the files you're uploading. For more information about storage classes, see [Storage Classes](#) in the *Amazon Simple Storage Service Developer Guide*.

The screenshot shows the 'Upload' console window with the 'Set properties' step selected. The progress bar at the top indicates four steps: 1. Select files, 2. Set permissions, 3. Set properties (active), and 4. Review. Below the progress bar, there is a table for selecting a storage class. The table has columns for Storage class, Designed for, Availability Zones, Min storage duration, Min billable object size, Monitoring and automation fees, and Retrieval fees. The 'Standard' class is selected. Below the table is the 'Encryption' section with a text description and three radio button options: None, Amazon S3 master-key, and AWS KMS master-key. At the bottom are 'Upload', 'Previous', and 'Next' buttons.

Storage class	Designed for	Availability Zones	Min storage duration	Min billable object size	Monitoring and automation fees	Retrieval fees
<input checked="" type="radio"/> Standard	Frequently accessed data	≥ 3	-	-	-	-
<input type="radio"/> Intelligent-Tiering	Long-lived data with changing or unknown access patterns	≥ 3	30 days	-	Per-object fees apply	-
<input type="radio"/> Standard-IA	Long-lived, infrequently accessed data	≥ 3	30 days	128KB	-	Per-GB fees apply
<input type="radio"/> One Zone-IA	Long-lived, infrequently accessed, non-critical data	≥ 1	30 days	128KB	-	Per-GB fees apply
<input type="radio"/> Glacier	Data archiving with retrieval times ranging from minutes to hours	≥ 3	90 days	-	-	Per-GB fees apply
<input type="radio"/> Reduced Redundancy (Not recommended)	Frequently accessed, non-critical data	≥ 3	-	-	-	-

Encryption

Protect data at rest by using Amazon S3 master-key or by using AWS KMS master-key.

☒ None ☐ Amazon S3 master-key ☐ AWS KMS master-key

[Upload](#) [Previous](#) [Next](#)

- b. Choose the type of encryption for the files that you're uploading. If you don't want to encrypt them, choose **None**.

- i. To encrypt the uploaded files using keys that are managed by Amazon S3, choose **Amazon S3 master-key**. For more information, see [Protecting Data with Amazon S3-Managed Encryption Keys Classes](#) in the *Amazon Simple Storage Service Developer Guide*.
- ii. To encrypt the uploaded files using the AWS Key Management Service (AWS KMS), choose **AWS KMS master-key**. Then choose a master key from the list of AWS KMS master keys.

Note

To encrypt objects in a bucket, you can use only keys that are available in the same AWS Region as the bucket.

You can give an external account the ability to use an object that is protected by an AWS KMS key. To do this, select **Custom KMS ARN** from the list and enter the Amazon Resource Name (ARN) for the external account. Administrators of an external account that have usage permissions to an object protected by your AWS KMS key can further restrict access by creating a resource-level IAM policy.

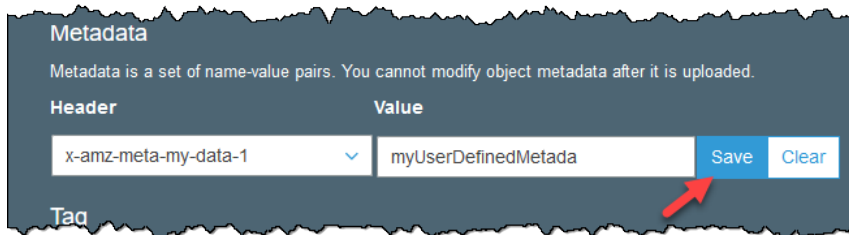
For more information about creating an AWS KMS key, see [Creating Keys](#) in the *AWS Key Management Service Developer Guide*. For more information about protecting data with AWS KMS, see [Protecting Data with AWS KMS-Managed Key](#) in the *Amazon Simple Storage Service Developer Guide*.

- c. Metadata for Amazon S3 objects is represented by a name-value (key-value) pair. There are two kinds of metadata: system-defined metadata and user-defined metadata.

If you want to add Amazon S3 system-defined metadata to all of the objects you are uploading, for **Header**, select a header. You can select common HTTP headers, such as **Content-Type** and **Content-Disposition**. Type a value for the header, and then choose **Save**. For a list of system-defined metadata and information about whether you can add the value, see [System-Defined Metadata](#) in the *Amazon Simple Storage Service Developer Guide*.

- d. Any metadata starting with prefix `x-amz-meta-` is treated as user-defined metadata. User-defined metadata is stored with the object, and is returned when you download the object.

To add user-defined metadata to all of the objects that you are uploading, type `x-amz-meta-` plus a custom metadata name in the **Header** field. Type a value for the header, and then choose **Save**. Both the keys and their values must conform to US-ASCII standards. User-defined metadata can be as large as 2 KB. For more information about user-defined metadata, see [User-Defined Metadata](#) in the *Amazon Simple Storage Service Developer Guide*.



Metadata

Metadata is a set of name-value pairs. You cannot modify object metadata after it is uploaded.

Header	Value
x-amz-meta-my-data-1	myUserDefinedMetada

Save **Clear**

Tag

- e. Object tagging gives you a way to categorize storage. Each tag is a key-value pair. Key and tag values are case sensitive. You can have up to 10 tags per object.

To add tags to all of the objects that you are uploading, type a tag name in the **Key** field. Type a value for the tag, and then choose **Save**. A tag key can be up to 128 Unicode characters in length and tag values can be up to 255 Unicode characters in length. For more information about object tags, see [Object Tagging](#) in the *Amazon Simple Storage Service Developer Guide*.

Tag



Add tags to search, organize and manage access

Key	Value
pet	dog


Save Clear

Upload Previous Next

7. Choose **Next**.
8. On the **Upload** review page, verify that your settings are correct, and then choose **Upload**. To make changes, choose **Previous**.
9. To see the progress of the upload, choose **In progress** at the bottom of the browser window.



To see a history of your uploads and other operations, choose **Success**.




Uploading Files by Pointing and Clicking


This procedure explains how to upload files into an S3 bucket by choosing **Upload**.

To upload files to an S3 bucket by pointing and clicking

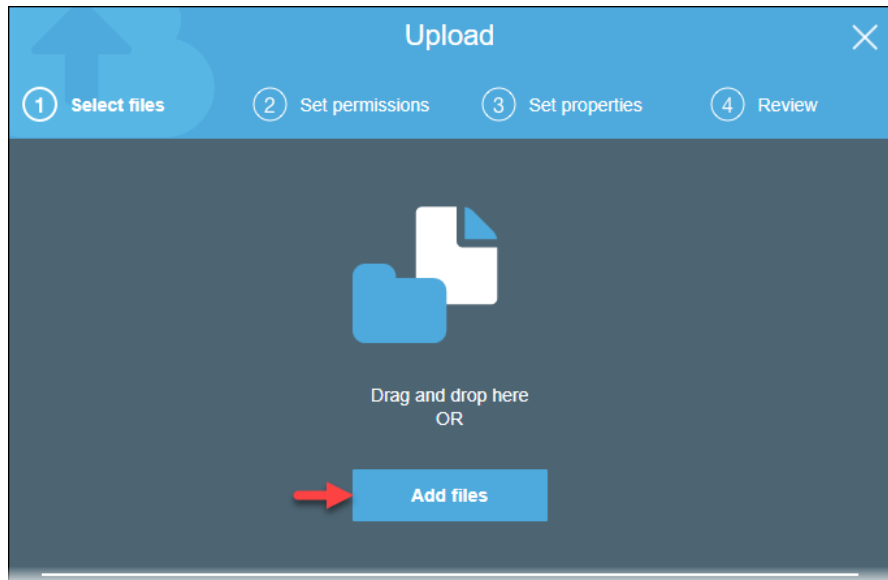
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to upload your files to.




3. Choose **Upload**.




4. In the **Upload** dialog box, choose **Add files**.



5. Choose one or more files to upload, and then choose **Open**.



6. After you see the files that you chose listed in the **Upload** dialog box, do one of the following:
 - a. To add more files, choose **Add more files**.
 - b. To immediately upload the listed files, choose **Upload**.
 - c. To set permissions or properties for the files that you are uploading, choose **Next**.



7. To set permissions and properties, start with **Step 5** of [Uploading Files and Folders by Using Drag and Drop](#) (p. 34).

More Info

- [How Do I Set Permissions on an Object?](#) (p. 116).
- [How Do I Download an Object from an S3 Bucket?](#) (p. 42)


How Do I Download an Object from an S3 Bucket?

This section explains how to use the Amazon S3 console to download objects from an S3 bucket.

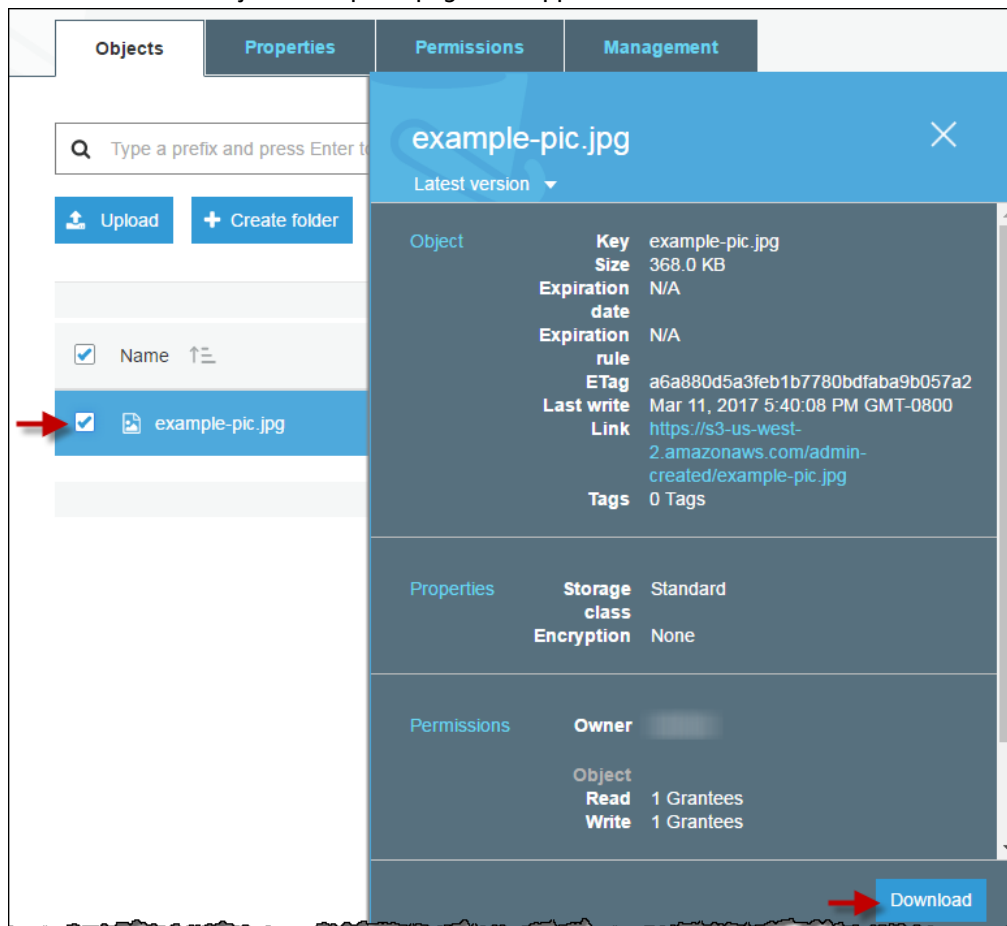
Data transfer fees apply when you download objects. For information about Amazon S3 features, and pricing, see [Amazon S3](#).

To download an object from an S3 bucket

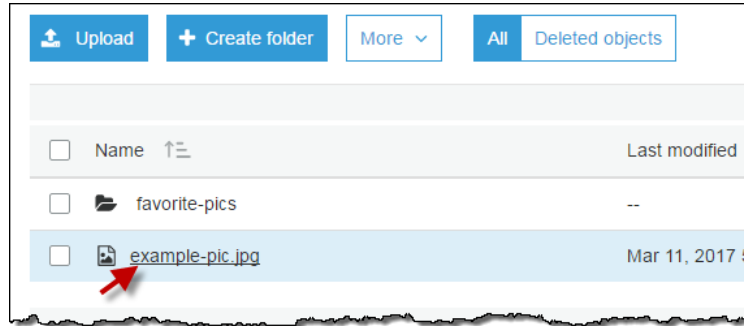
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to download an object from.




3. You can download an object from an S3 bucket in any of the following ways:
 - In the **Name** list, select the check box next to the object you want to download, and then choose **Download** on the object description page that appears.




- Choose the name of the object that you want to download.




On the **Overview** page, choose **Download**.



- Choose the name of the object that you want to download and then choose **Download as** on the **Overview** page.



- Choose the name of the object that you want to download. Choose **Latest version** and then choose the download icon.



Related Topics

- [How Do I Upload Files and Folders to an S3 Bucket? \(p. 33\)](#)


How Do I Delete Objects from an S3 Bucket?

This section explains how to use the Amazon S3 console to delete objects. Because all objects in your S3 bucket incur storage costs, you should delete objects that you no longer need. If you are collecting log files, for example, it's a good idea to delete them when they're no longer needed. You can set up a lifecycle rule to automatically delete objects such as log files.


For information about Amazon S3 features and pricing, see [Amazon S3](#).

To delete objects from an S3 bucket

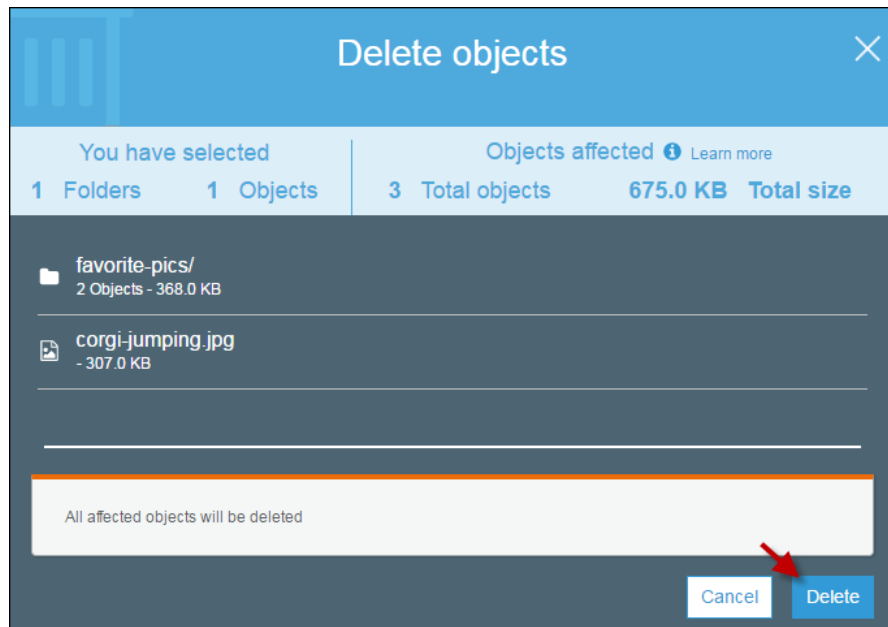
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to delete an object from.



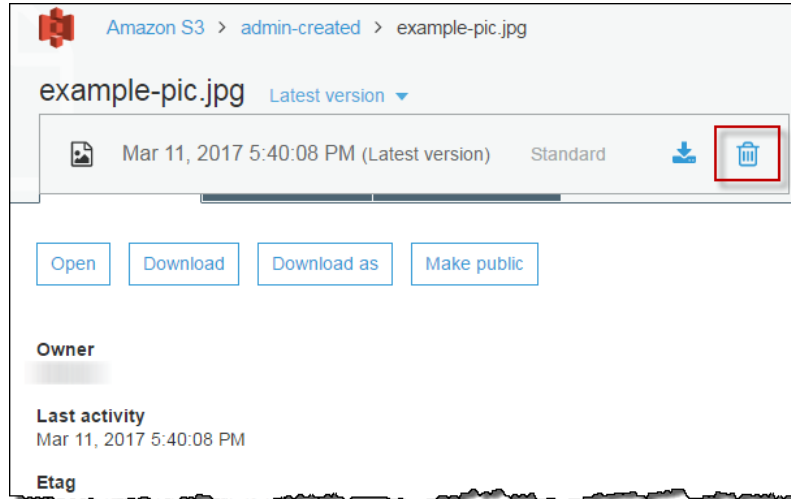
3. You can delete objects from an S3 bucket in any of the following ways:
 - In the **Name** list, select the check box next to the objects and folders that you want to delete, choose **More**, and then choose **Delete**.



In the **Delete objects** dialog box, verify that the names of the objects and folders you selected for deletion are listed and then choose **Delete**.



- Choose the name of the object that you want to delete, choose **Latest version**, and then choose the trash can icon.



More Info

- [How Do I Undelete a Deleted S3 Object? \(p. 46\)](#)
- [How Do I Create a Lifecycle Policy for an S3 Bucket? \(p. 78\)](#)

How Do I Undelete a Deleted S3 Object?


This section explains how to use the Amazon S3 console to recover (undelete) deleted objects.

To be able to undelete a deleted object, you must have had versioning enabled on the bucket that contains the object before the object was deleted. For information about enabling versioning, see [How Do I Enable or Suspend Versioning for an S3 Bucket? \(p. 12\)](#).


When you delete an object in a versioning-enabled bucket, all versions remain in the bucket and Amazon S3 creates a delete marker for the object. To undelete the object, you must delete this delete marker. For more information about versioning and delete markers, see [Object Versioning](#) in the *Amazon Simple Storage Service Developer Guide*.

To recover deleted objects from an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want.






3. To see a list of the versions of the objects in the bucket, choose **Show**. You'll be able to see the delete markers for deleted objects.



4. To undelete an object, you must delete the delete marker. Select the check box next to the delete marker of the object to recover, and then choose **delete** from the **More** menu.



<input type="checkbox"/>	Name	Version ID
	amazon-reindeer.jpg	
<input type="checkbox"/>	 Sep 1, 2017 5:43:05 PM (Latest version)	na_axXxRr.xXDaWfcp4idCkIpXsyv6m9
<input type="checkbox"/>	 Sep 1, 2017 5:42:50 PM	cM5luTOwME0WUDRZuAqv7vbhm7ZI...
	amazon-spheres.jpg	
<input checked="" type="checkbox"/>	 Sep 1, 2017 5:46:31 PM (Delete marker)	XfdtMN.2X.yHhbNlceyAoM_mlkDA3Nzi
<input type="checkbox"/>	 Sep 1, 2017 5:43:07 PM	1yAY8OBXQoaELJ0L393xCg.CqjbGe...
<input type="checkbox"/>	 Sep 1, 2017 5:42:52 PM	Ai4hSgfCjJB902ygpjQErUQPbVj7HMur
	screen-shot3.png	
<input type="checkbox"/>	 Apr 15, 2016 4:04:18 PM (Latest version)	null

5. Choose **Hide**, you'll see the undeleted object listed.

More Info

- [How Do I See the Versions of an S3 Object? \(p. 57\)](#)
- [How Do I Enable or Suspend Versioning for an S3 Bucket? \(p. 12\)](#)
- [Using Versioning in the Amazon Simple Storage Service Developer Guide](#)

How Do I Restore an S3 Object That Has Been Archived to Glacier?

This section explains how to use the Amazon S3 console to restore an object that has been archived by using the `GLACIER` storage class. Objects in the `GLACIER` storage class are not immediately accessible. To access an object in this class, you must restore a temporary copy of it to its S3 bucket. You can access objects that have been archived to Glacier only by using Amazon S3. You can't use the Glacier console, AWS command line interface (CLI), or APIs to see the S3 archived objects. For information about when to use the `GLACIER` storage class for objects, see [Storage Classes](#) and [Object Lifecycle Management](#) in the *Amazon Simple Storage Service Developer Guide*.

Glacier charges a retrieval fee for retrieving objects stored with the `GLACIER` storage class. For retrieval pricing information, see [Glacier Pricing](#).

When you restore an archive, you pay for both the archive and the restored copy. Because there is a storage cost for the copy, restore objects only for the duration you need them. If you want a permanent copy of the object, create a copy of it in your S3 bucket. For information about Amazon S3 features and pricing, see [Amazon S3](#).

After restoring an object, you can download it from the **Overview** page. For more information, see [How Do I See an Overview of an Object? \(p. 55\)](#).

Topics

- [Archive Retrieval Options \(p. 48\)](#)

- [Restoring an Archived S3 Object](#) (p. 48)
- [Upgrade an In-Progress Restore](#) (p. 50)
- [Checking Archive Restore Status and Expiration Date](#) (p. 51)

Archive Retrieval Options

You restore archived objects using one of the following retrieval types:

- **Expedited** - Expedited retrievals allow you to quickly access your data when occasional urgent requests for a subset of archives are required. For all but the largest archived objects (250 MB+), data accessed using Expedited retrievals are typically made available within 1–5 minutes. Provisioned Capacity ensures that retrieval capacity for Expedited retrievals is available when you need it. For more information, see [Provisioned Capacity](#).
- **Standard** - Standard retrievals allow you to access any of your archived objects within several hours. Standard retrievals typically complete within 3–5 hours. This is the default option for retrieval requests that do not specify the retrieval option.
- **Bulk** - Bulk retrievals are Glacier's lowest-cost retrieval option, enabling you to retrieve large amounts, even petabytes, of data inexpensively in a day. Bulk retrievals typically complete within 5–12 hours.


For more information about retrieval options, see [Restoring Archived Objects](#) in the *Amazon Simple Storage Service Developer Guide*.

Restoring an Archived S3 Object


This topic explains how to use the Amazon S3 console to restore an object that has been archived to Glacier.

To restore archived S3 objects

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the objects that you want to restore.



3. In the **Name** list, select the objects that you want to restore, choose **Actions**, and then choose **Restore from Glacier**.



4. In the **Initiate restore** dialog box, type the number of days that you want your archived data to be accessible.
5. Choose one of the following retrieval options from the **Retrieval options** menu.
 - Choose **Bulk retrieval** or **Standard retrieval**, and then choose **Restore**.
 - Choose **Expedited retrieval**.

A screenshot of the 'Restore objects from Glacier' dialog box. The title bar says 'Restore objects from Glacier' with a close button. Below the title bar, it shows 'Selection: 3 Objects, 0 Folders', 'Total size: 1.7 MB', and 'Total objects: 3'. The main section is titled 'Number of days the restored copy is available' and explains that the restored copy in the Reduced Redundancy Storage (RRS) is automatically deleted after the specified number of days. There is a text input field with '10' and the label 'days'. Below this, it says 'Available until approximately 2018-12-05'. The next section is 'Restore tier', which explains that Glacier charges request fees and per GB retrieval fees, which vary based on the tier selected. There are three radio button options: 'Bulk retrieval' (selected) with the note 'Typically within 5-12 hours', 'Standard retrieval' with the note 'Typically within 3 - 5 hours', and 'Expedited retrieval' with the note 'Typically within 1 - 5 minutes when retrieving less than 250MB'. At the bottom right, there are 'Cancel' and 'Restore' buttons.

6. If you have provisioned capacity, choose **Restore** to start a provisioned retrieval. If you have provisioned capacity, all of your expedited retrievals are served by your provisioned capacity. For more information about provisioned capacity, see [Provisioned Capacity](#).

- If you don't have provisioned capacity and you don't want to buy it, choose **Restore**.
- If you don't have provisioned capacity, but you want to buy it, choose **Add capacity unit**, and then choose **Buy**. When you get the **Purchase succeeded** message, choose **Restore** to start provisioned retrieval.

☒ Expedited retrieval
Typically within 1 - 5 minutes when retrieving less than 250MB

Purchased capacity units: 0

[Add 1 capacity unit](#)

Purchase 1 provisioned capacity unit.

You will be immediately charged for each provisioned capacity unit and the purchase is not refundable. See [S3 pricing](#).

Provisioned capacity ensures that retrieval capacity for expedited retrievals is available when you need it. Each unit of capacity provides that at least three expedited retrievals can be performed every five minutes and provides up to 150 MB/s of retrieval throughput.

Once purchased, provisioned capacity units will be available for your use in the current region for one month from the date of purchase.

[Purchase](#)

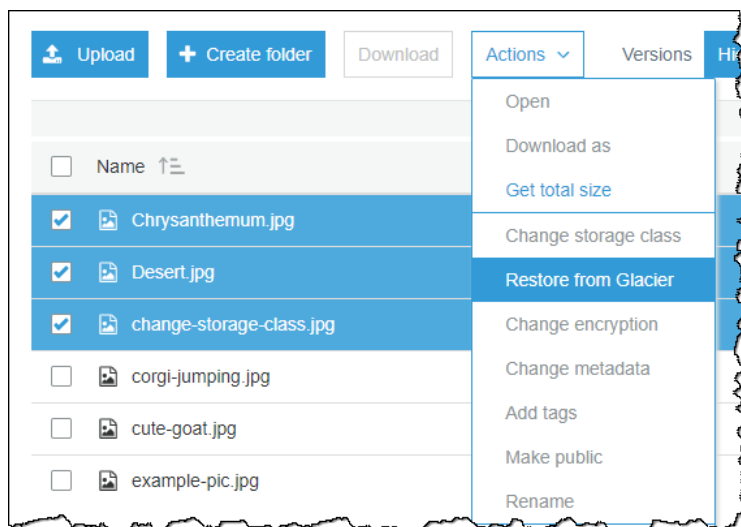
[Cancel](#) [Restore](#)

Upgrade an In-Progress Restore

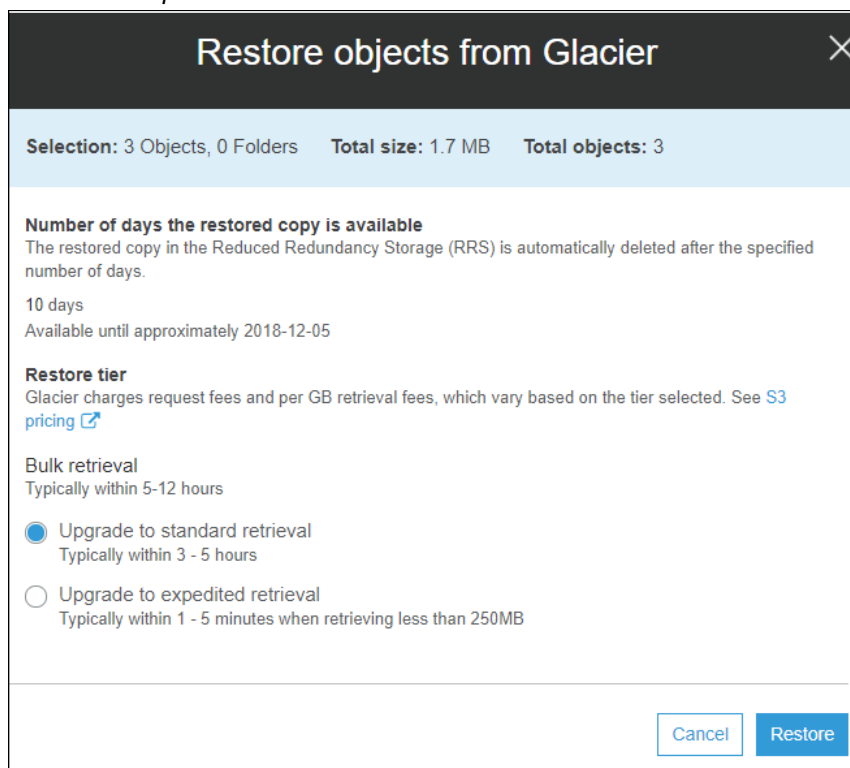
You can upgrade the speed of your restoration while it is in progress.

To upgrade an in-progress restore to a faster tier

1. In the **Name** list, select one or more of the objects that you are restoring, choose **Actions**, and then choose **Restore from Glacier**. For information about checking the restoration status of an object, see [Checking Archive Restore Status and Expiration Date](#) (p. 51).



2. Choose the tier that you want to upgrade to and then choose **Restore**. For more information about upgrading to a faster restore tier, see [Restoring Archived Objects](#) in the *Amazon Simple Storage Service Developer Guide*.




Checking Archive Restore Status and Expiration Date

To check the progress of the restoration, see the object overview panel. For information about the overview panel, see [How Do I See an Overview of an Object?](#) (p. 55).

The **Overview** section shows that restoration is **In progress**.

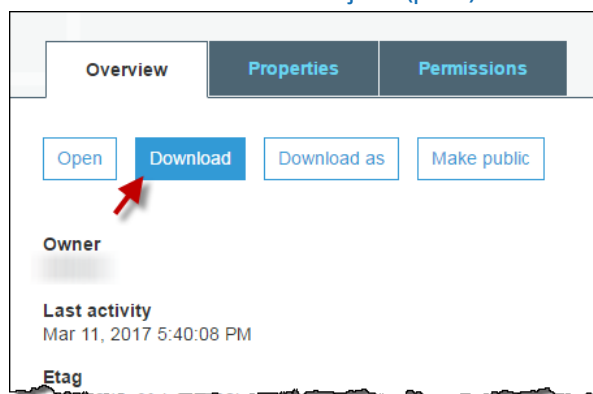
When the temporary copy of the object is available, the object's **Overview** section shows the **Restoration expiry date**. This is when Amazon S3 will remove the restored copy of your archive.



Restored objects are stored only for the number of days that you specify. If you want a permanent copy of the object, create a copy of it in your Amazon S3 bucket.

Amazon S3 calculates the expiry date by adding the number of days that you specify to the time you request to restore the object, and then rounding to the next day at midnight UTC. This calculation applies to the initial restoration of the object and to any extensions to availability that you request. For example, if an object was restored on 10/15/2012 10:30 AM UTC and the number of days that you specified is 3, then the object is available until 10/19/2012 00:00 UTC. If, on 10/16/2012 11:00 AM UTC you change the number of days that you want it to be accessible to 1, then Amazon S3 makes the restored object available until 10/18/2012 00:00 UTC.

After restoring an object, you can download it from the **Overview** page. For more information, see [How Do I See an Overview of an Object?](#) (p. 55).



More Info

- [How Do I Create a Lifecycle Policy for an S3 Bucket?](#) (p. 78)
- [How Do I Undelete a Deleted S3 Object?](#) (p. 46)

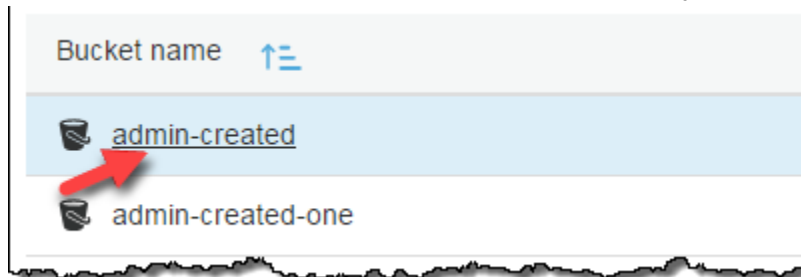
How Do I Lock an S3 Object?

Amazon S3 object lock enables you to store objects using a "Write Once Read Many" (WORM) model. When using object lock you can prevent an object from being deleted or overwritten for a fixed amount of time or indefinitely. For information about object locking using the AWS CLI, AWS SDKs, and the Amazon S3 REST APIs, see [Introduction to Amazon S3 Object Lock](#) in the *Amazon Simple Storage Service Developer Guide*.

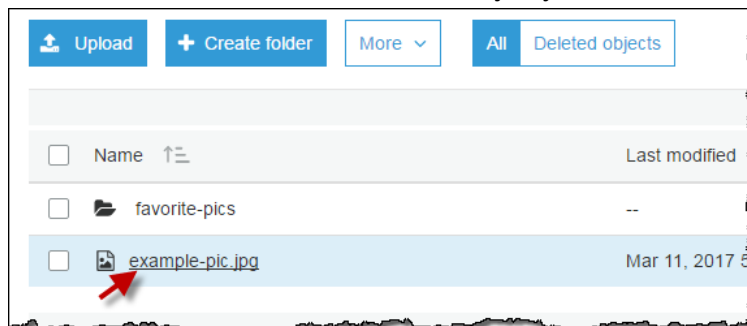
Before you lock any objects, you have to enable a bucket to use Amazon S3 object lock. You enable object lock when you create a bucket. Once you enable object lock on a bucket you can lock objects in that bucket. Once you create a bucket with object lock enabled, you can't disable object lock or suspend versioning for the bucket. For information on creating a bucket with object lock enabled, see [How Do I Create an S3 Bucket?](#) (p. 3).

To lock an S3 object

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want.



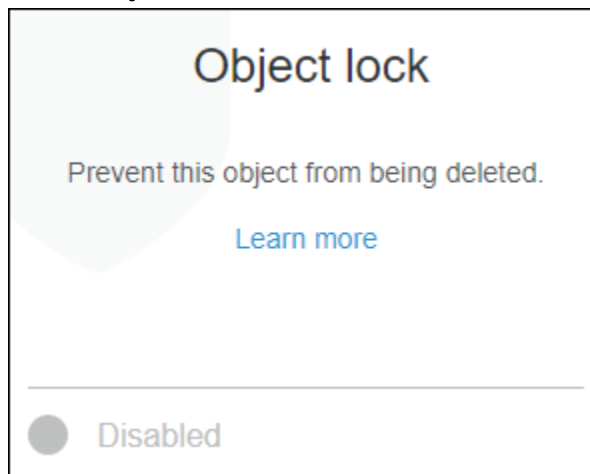
3. In the **Name** list, choose the name of the object you want to lock.



4. Choose **Properties**.



5. Choose **Object Lock**.



6. Choose a retention mode. You can change the **Retain until date**. You can also choose to enable a legal hold. For more information about retention modes and a legal hold, see [Amazon S3 Object Lock Overview](#) in the *Amazon Simple Storage Service Developer Guide*

Object lock

Prevent objects from being deleted in order to help ensure data integrity and regulatory compliance. [Learn more](#)

Retention mode

☒ **Enable governance mode**
Governance mode can be disabled by AWS accounts that have specific IAM permissions.

☐ **Enable compliance mode**
Compliance mode cannot be disabled by any user, including the root account.

☐ **Disable**

Retain until date

2018-11-27

Legal hold

Legal hold prevents an object from being deleted regardless of its retain until date. Legal hold can be applied and removed by AWS accounts that have specific IAM permissions.

☐ **Enable**

☒ **Disable**

Cancel

Save

7. Choose **Save**.

More Info


- [Setting Bucket and Object Access Permissions](#) (p. 110)

How Do I See an Overview of an Object?

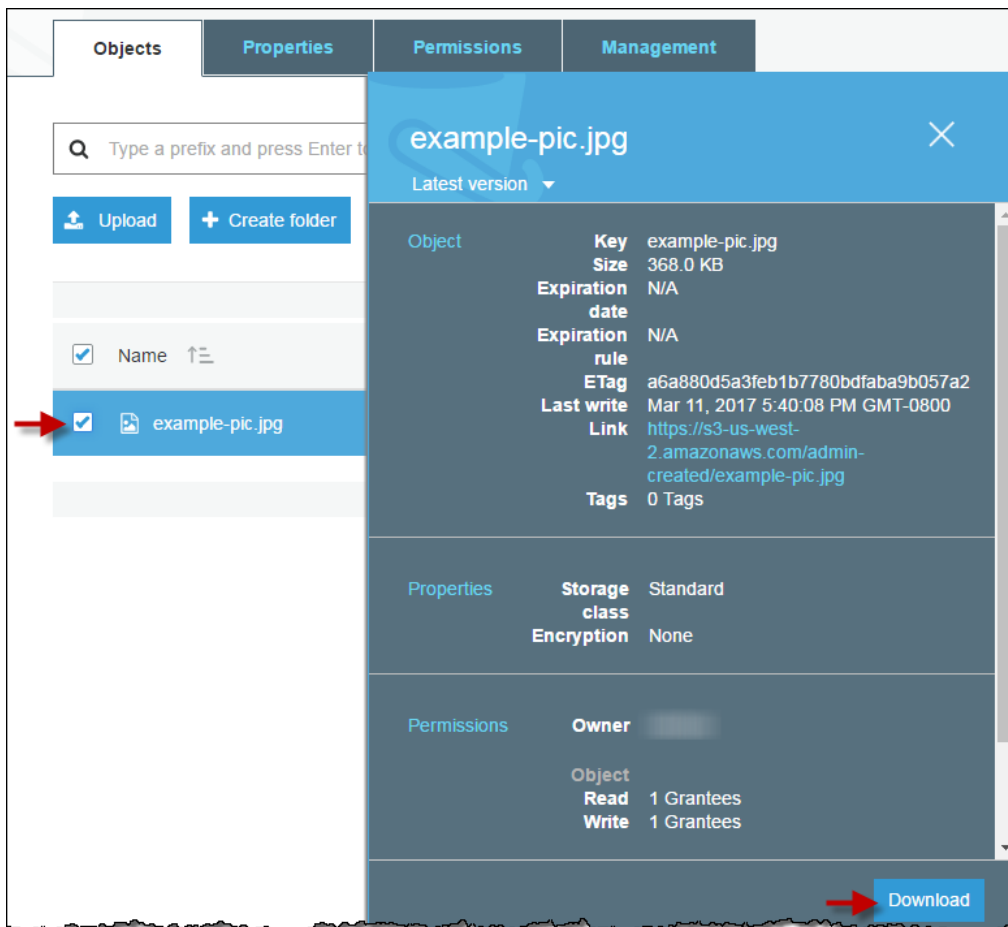
This section explains how to use the Amazon S3 console to view the object overview panel. This panel provides an overview of all the essential information for an object in one place.

To see the overview panel for an object


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the object.



3. In the **Name** list, select the check box next to the name of the object for which you want an overview.



4. To download the object, choose **Download** in the object overview panel. To copy the path of the object to the clipboard, choose **Copy Path**.



5. If versioning is enabled on the bucket, choose **Latest versions** to list the versions of the object. You can then choose the download icon to download an object version, or choose the trash can icon to delete an object version.



Important

You can undelete an object only if it was deleted as the latest (current) version. You can't undelete a previous version of an object that was deleted. For more information, see [Object Versioning](#) and [Using Versioning](#) in the *Amazon Simple Storage Service Developer Guide*.

More Info

- [How Do I See the Versions of an S3 Object? \(p. 57\)](#)

How Do I See the Versions of an S3 Object?

This section explains how to use the Amazon S3 console to see the different versions of an object.

A versioning-enabled bucket can have many versions of the same object; one current (latest) version and zero or more noncurrent (previous) versions. Amazon S3 assigns each object a unique version ID. For information about enabling versioning, see [How Do I Enable or Suspend Versioning for an S3 Bucket?](#) (p. 12).


If a bucket is versioning-enabled, Amazon S3 creates another version of an object under the following conditions:

- If you upload an object that has the same name as an object that already exists in the bucket, Amazon S3 creates another version of the object instead of replacing the existing object.
- If you update any object properties after you upload the object to the bucket, such as changing the storage details or other metadata, Amazon S3 creates a new object version in the bucket.

For more information about versioning support in Amazon S3, see [Object Versioning](#) and [Using Versioning](#) in the *Amazon Simple Storage Service Developer Guide*.

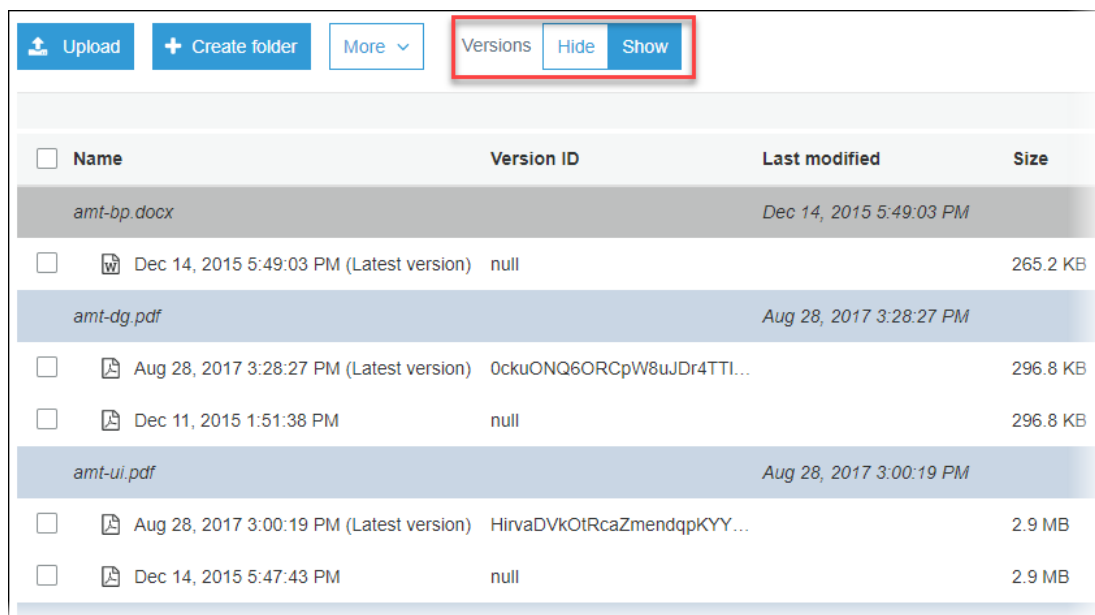
To see multiple versions of an object

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the object.



3. To see a list of the versions of the objects in the bucket, choose **Show**. For each object version, the console shows a unique version ID, the date and time the object version was created, and other properties. (Objects stored in your bucket before you set the versioning state have a version ID of **null**.)

To list the objects without the versions, choose **Hide**.



<input type="checkbox"/>	Name	Version ID	Last modified	Size
<input type="checkbox"/>	amt-bp.docx		Dec 14, 2015 5:49:03 PM	
<input type="checkbox"/>	amt-dg.pdf		Aug 28, 2017 3:28:27 PM	
<input type="checkbox"/>	amt-ui.pdf		Aug 28, 2017 3:00:19 PM	
<input type="checkbox"/>	Dec 14, 2015 5:49:03 PM (Latest version)	null		265.2 KB
<input type="checkbox"/>	Aug 28, 2017 3:28:27 PM (Latest version)	0ckuONQ6ORCpW8uJDr4TTI...		296.8 KB
<input type="checkbox"/>	Dec 11, 2015 1:51:38 PM	null		296.8 KB
<input type="checkbox"/>	Aug 28, 2017 3:00:19 PM (Latest version)	HirvaDVkOtRcaZmendqpKYY...		2.9 MB
<input type="checkbox"/>	Dec 14, 2015 5:47:43 PM	null		2.9 MB

You also can view, download, and delete object versions in the object overview panel. For more information, see [How Do I See an Overview of an Object? \(p. 55\)](#).

Important

You can undelete an object only if it was deleted as the latest (current) version. You can't undelete a previous version of an object that was deleted. For more information, see [Object Versioning](#) and [Using Versioning](#) in the *Amazon Simple Storage Service Developer Guide*.

More Info


- [How Do I Enable or Suspend Versioning for an S3 Bucket? \(p. 12\)](#)
- [How Do I Create a Lifecycle Policy for an S3 Bucket? \(p. 78\)](#)

How Do I View the Properties of an Object?

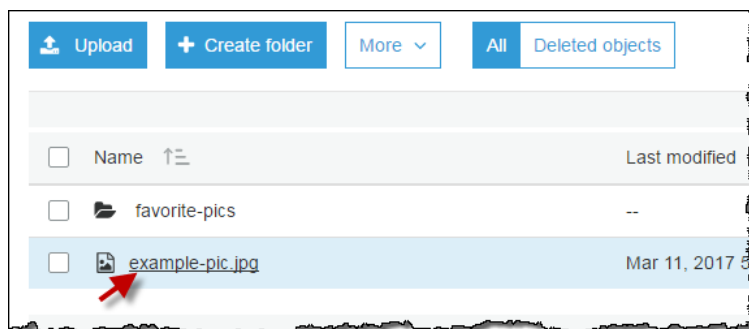
This section explains how to use the console to view the properties of an object.

To view the properties of an object


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the object.



3. In the **Name** list, choose the name of the object you want to view the properties for.



4. Choose **Properties**.



5. On the **Properties** page, you can configure the following properties for the object.

- a. **Storage class** – Each object in Amazon S3 has a storage class associated with it. The storage class that you choose to use depends on how frequently you access the object. The default storage class for S3 objects is STANDARD. You choose which storage class to use when you upload an object. For more information about storage classes, see [Storage Classes](#) in the *Amazon Simple Storage Service Developer Guide*.

To change the storage class after you upload an object, choose **Storage class**. Choose the storage class that you want, and then choose **Save**.

- b. **Encryption** – You can encrypt your S3 objects. For more information, see [How Do I Add Encryption to an S3 Object?](#) (p. 60).
- c. **Metadata** – Each object in Amazon S3 has a set of name-value pairs that represents its metadata. For information on adding metadata to an S3 object, see [How Do I Add Metadata to an S3 Object?](#) (p. 63).


- d. **Tags** – You can add tags to an S3 object. For more information, see [How Do I Add Tags to an S3 Object?](#) (p. 68).

How Do I Add Encryption to an S3 Object?

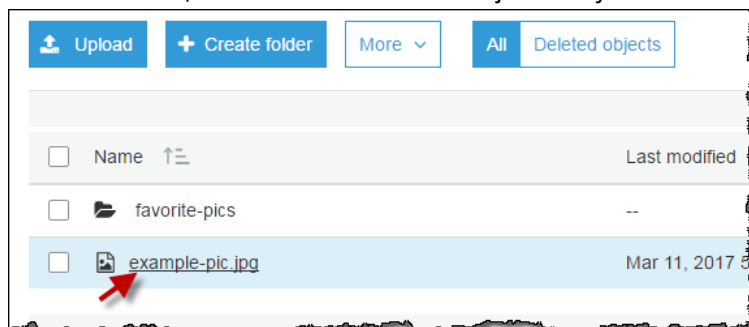
This topic describes how to set or change the type of encryption an object is using.

To add encryption to an object


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the object.




3. In the **Name** list, choose the name of the object that you want to add encryption to.



4. Choose **Properties**, and then choose **Encryption**.



5. Select **AES-256** or **AWS-KMS**.
- a. To encrypt your object using keys that are managed by Amazon S3, select **AES-256**. For more information about using Amazon S3 server-side encryption to encrypt your data, see [Protecting Data with Amazon S3-Managed Encryption Keys Classes](#) in the *Amazon Simple Storage Service Developer Guide*.



- b. To encrypt your object using AWS Key Management Service (AWS KMS), choose **AWS-KMS**, choose a master key from the list of the AWS KMS master keys that you have created, and then choose **Save**.

Note

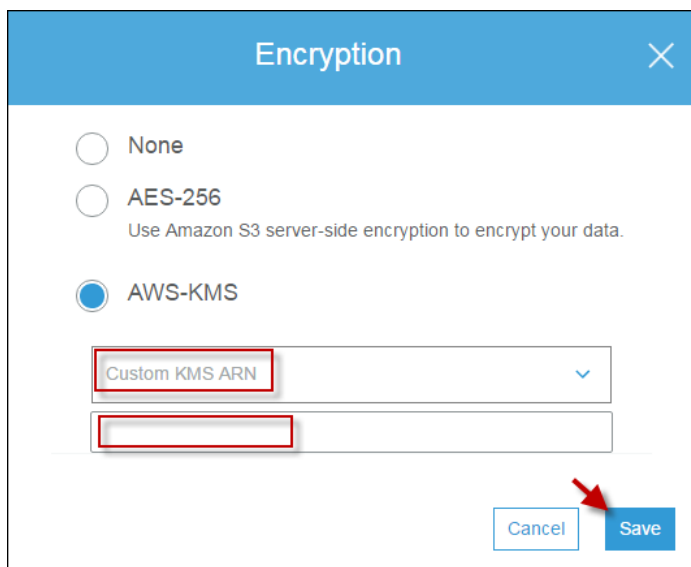
To encrypt objects in the bucket, you can use only keys that are enabled in the same AWS Region as the bucket.

For more information about creating an AWS KMS key, see [Creating Keys](#) in the *AWS Key Management Service Developer Guide*. For more information, see [Protecting Data with AWS KMS–Managed Key](#) in the *Amazon Simple Storage Service Developer Guide*.

The screenshot shows the 'Encryption' dialog box in the AWS S3 console. It has a blue header with the title 'Encryption' and a close button. Below the header, there are three radio buttons: 'None', 'AES-256' (with a subtext 'Use Amazon S3 server-side encryption to encrypt your data.'), and 'AWS-KMS'. The 'AWS-KMS' radio button is selected, indicated by a red arrow. Below the radio buttons, there is a dropdown menu labeled 'Select a key'. The dropdown is open, showing a search bar with the placeholder 'Type to search' and a magnifying glass icon. Below the search bar, there is a list of keys, including 'arn:aws:kms:us-west-2:...' and 'Custom KMS ARN'. A red arrow points to the 'Custom KMS ARN' option. To the right of the dropdown, there is a 'Save' button.

You can give an external account the ability to use an object that is protected by an AWS KMS key. To do this, select **Custom KMS ARN** from the list, type the Amazon Resource Name (ARN) for the external account, and then choose **Save**. Administrators of an external account that have usage permissions to an object protected by your AWS KMS key can further restrict access by creating a resource-level AWS Identity and Access Management (IAM) policy.

This screenshot is identical to the one above, showing the 'Encryption' dialog box with the 'AWS-KMS' radio button selected and the 'Custom KMS ARN' option highlighted in the dropdown menu. A red arrow points to the 'Custom KMS ARN' option.



More Info

- [How Do I View the Properties of an Object? \(p. 58\)](#)
- [Uploading, Downloading, and Managing Objects \(p. 33\)](#)

How Do I Add Metadata to an S3 Object?

Each object in Amazon Simple Storage Service (Amazon S3) has a set of name-value pairs that provides metadata about the object. *Metadata* is additional information about the object. Some metadata is set by Amazon S3 when you upload the object, for example, `Date` and `Content-Length`. You can also set some metadata when you upload the object, or you can add it later. This section explains how to use the Amazon S3 console to add metadata to an S3 object.

Object metadata is a set of name-value (key-value) pairs. For example, the metadata for content length, `Content-Length`, is the name (key) and the size of the object in bytes (value). For more information about object metadata, see [Object Metadata](#) in the *Amazon Simple Storage Service Developer Guide*.

There are two kinds of metadata for an S3 object, Amazon S3 system metadata and user-defined metadata:

- **System metadata**—There are two categories of system metadata. Metadata such as the `Last-Modified` date is controlled by the system. Only Amazon S3 can modify the value. There is also system metadata that you control, for example, the storage class configured for the object.
- **User-defined metadata**—You can define your own custom metadata, called user-defined metadata. You can assign user-defined metadata to an object when you upload the object or after the object has been uploaded. User-defined metadata is stored with the object and is returned when you download the object. Amazon S3 does not process user-defined metadata.

The following topics describe how to add metadata to an object.

Topics

- [Adding System-Defined Metadata to an S3 Object \(p. 64\)](#)


- [Adding User-Defined Metadata to an S3 Object \(p. 66\)](#)

Adding System-Defined Metadata to an S3 Object


You can configure some system metadata for an S3 object. For a list of system-defined metadata and whether you can modify their values, see [System-Defined Metadata](#) in the *Amazon Simple Storage Service Developer Guide*.

To add system metadata to an object


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the object.



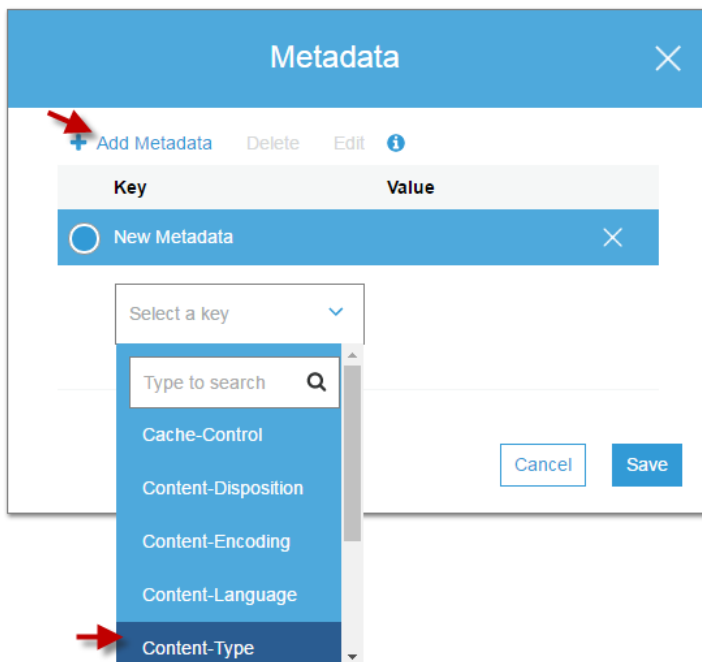
3. In the **Name** list, choose the name of the object that you want to add metadata to.




4. Choose **Properties**, and then choose **Metadata**.




5. Choose **Add Metadata**, and then choose a key from the **Select a key** menu.



6. Depending on which key you chose, choose a value from the **Select a value** menu or type a value.



7. Choose **Save**.




Adding User-Defined Metadata to an S3 Object

You can assign user-defined metadata to an object. User-defined metadata must begin with the prefix "x-amz-meta-", otherwise Amazon S3 will not set the key value pair as you define it. You define custom metadata by adding a name that you choose to the x-amz-meta- key. This creates a custom key. For example, if you add the custom name alt-name, the metadata key would be x-amz-meta-alt-name.


User-defined metadata can be as large as 2 KB. Both keys and their values must conform to US-ASCII standards. For more information, see [User-Defined Metadata](#) in the *Amazon Simple Storage Service Developer Guide*.

To add user-defined metadata to an object

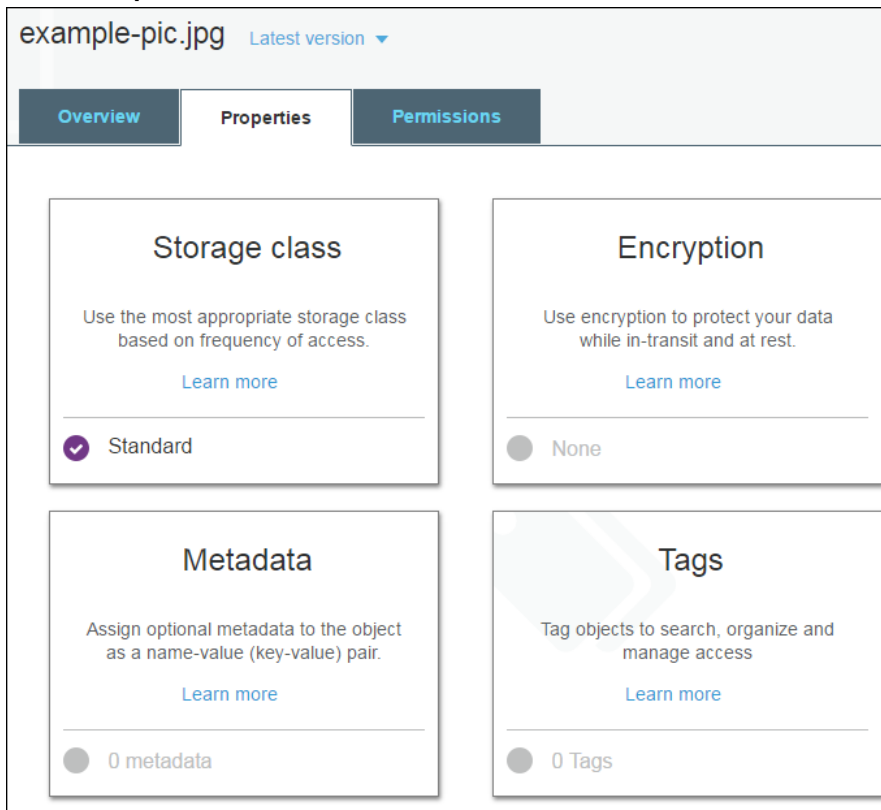
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the object.




3. In the **Name** list, choose the name of the object that you want to add metadata to.




4. Choose **Properties**, and then choose **Metadata**.



5. Choose **Add Metadata**, and then choose the **x-amz-meta-** key from the **Select a key** menu. Any metadata starting with the prefix **x-amz-meta-** is user-defined metadata.



6. Type a custom name following the `x-amz-meta-` key. For example, for the custom name `alt-name`, the metadata key would be `x-amz-meta-alt-name`. Enter a value for the custom key, and then choose **Save**.



- [How Do I View the Properties of an Object? \(p. 58\)](#)
- [Uploading, Downloading, and Managing Objects \(p. 33\)](#)

How Do I Add Tags to an S3 Object?

Object tagging gives you a way to categorize storage. This topic explains how to use the console to add tags to an S3 object after the object has been uploaded. For information about adding tags to an object when the object is being uploaded, see [How Do I Upload Files and Folders to an S3 Bucket? \(p. 33\)](#).


Each tag is a key-value pair that adheres to the following rules:

- You can associate up to 10 tags with an object. Tags associated with an object must have unique tag keys.
- A tag key can be up to 128 Unicode characters in length and tag values can be up to 255 Unicode characters in length.
- Key and tag values are case sensitive.

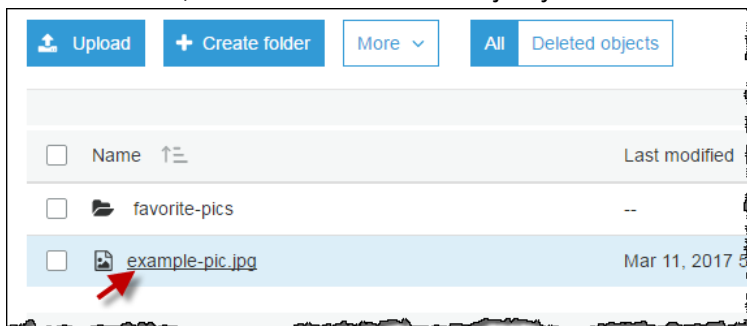
For more information about object tags, see [Object Tagging](#) in the *Amazon Simple Storage Service Developer Guide*.

To add tags to an object

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that contains the object.



3. In the **Name** list, choose the name of the object you want to add tags to.



4. Choose **Properties**.

example-pic.jpg Latest version ▾

Overview Properties Permissions

Storage class

Use the most appropriate storage class based on frequency of access.

[Learn more](#)

✓ Standard

Encryption

Use encryption to protect your data while in-transit and at rest.

[Learn more](#)

○ None

Metadata

Assign optional metadata to the object as a name-value (key-value) pair.

[Learn more](#)

○ 0 metadata

Tags

Tag objects to search, organize and manage access

[Learn more](#)

○ 0 Tags

5. Choose **Tags** and then choose **Add Tag**.

Tags ✕

Key	Value
+ Add Tag	

To enable replication of object tags IAM policies used for Cross-Region Replication must be updated if they were created prior to the introduction of Object tagging.

Cancel Save

6. Each tag is a key-value pair. Type a **Key** and a **Value**. Then choose **Add Tag** to add another tag or choose **Save**.

You can enter up to 10 tags for an object.

Key	Value
dog	corgi X

+ Add Tag

To enable replication of object tags IAM policies used for Cross-Region Replication must be updated if they were created prior to the introduction of Object tagging.

Cancel Save

More Info

- [How Do I View the Properties of an Object? \(p. 58\)](#)
- [Uploading, Downloading, and Managing Objects \(p. 33\)](#)

How do I use folders in an S3 Bucket?

In Amazon S3, buckets and objects are the primary resources, where objects are stored in buckets. Amazon S3 has a flat structure with no hierarchy like you would see in a file system. However, for the sake of organizational simplicity, the Amazon S3 console supports the folder concept as a means of grouping objects. Amazon S3 does this by using a shared name prefix for objects (that is, objects that have names that begin with a common string). Object names are also referred to as key names.

For example, you can create a folder in the console called `photos`, and store an object named `myphoto.jpg` in it. The object is then stored with the key name `photos/myphoto.jpg`, where `photos/` is the prefix.

Here are two more examples:

- If you have three objects in your bucket—`logs/date1.txt`, `logs/date2.txt`, and `logs/date3.txt`—the console will show a folder named `logs`. If you open the folder in the console, you will see three objects: `date1.txt`, `date2.txt`, and `date3.txt`.
- If you have an object named `photos/2017/example.jpg`, the console will show you a folder named `photos` containing the folder `2017` and the object `example.jpg`.

Topics

- [Creating a Folder \(p. 72\)](#)
- [How Do I Delete Folders from an S3 Bucket? \(p. 73\)](#)
- [Making Folders Public \(p. 75\)](#)

You can have folders within folders, but not buckets within buckets. You can upload and copy objects directly into a folder. Folders can be created, deleted, and made public, but they cannot be renamed. Objects can be copied from one folder to another.

Important

The Amazon S3 console treats all objects that have a forward slash "/" character as the last (trailing) character in the key name as a folder, for example `examplekeyname/`. You cannot upload an object with a key name with a trailing "/" character by using the Amazon S3 console. However, objects named with a trailing "/" can be uploaded with the Amazon S3 API by using the AWS CLI, the AWS SDKs, or REST API.


An object named with a trailing "/" displays as a folder in the Amazon S3 console. The Amazon S3 console does not display the content and metadata for such an object. When copying an object named with a trailing "/" by using the Amazon S3 console, a new folder is created in the destination location but the object's data and metadata are not copied.

Creating a Folder

This section describes how to use the Amazon S3 console to create a folder.

To create a folder

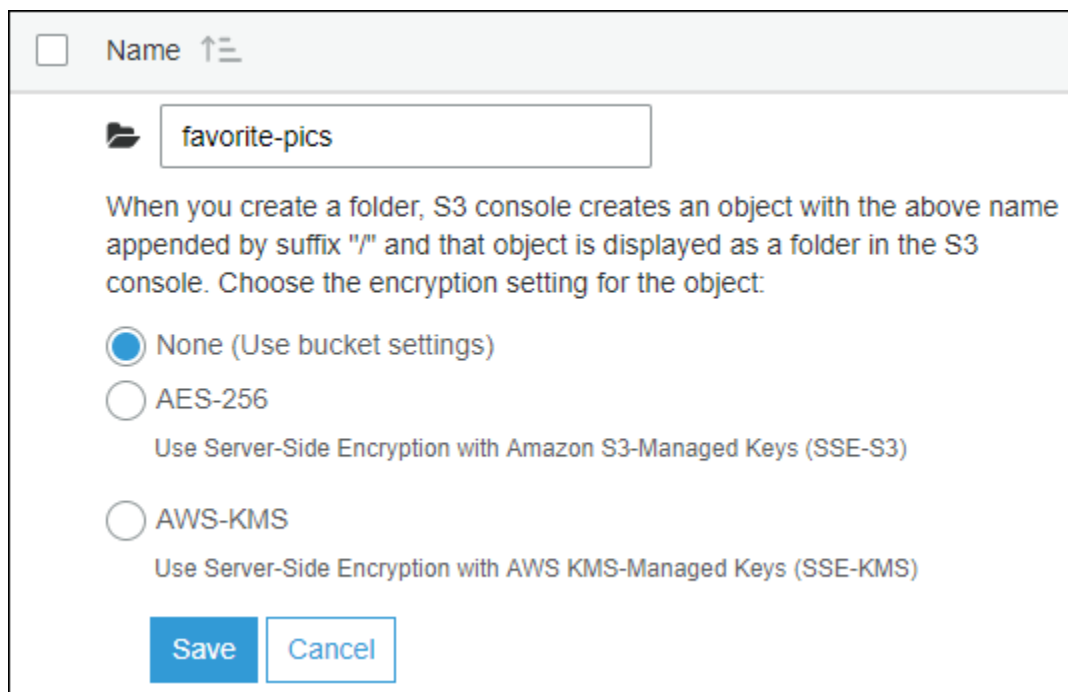
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to create a folder in.



3. Choose **Create folder**.



4. Type a name for the folder, for example, favorite-pics. Choose the encryption setting for the folder object and then choose **Save**.




How Do I Delete Folders from an S3 Bucket?

This section explains how to use the Amazon S3 console to delete folders from an S3 bucket.


For information about Amazon S3 features and pricing, see [Amazon S3](#).

To delete folders from an S3 bucket


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to delete folders from.



3. In the **Name** list, select the check box next to the folders and objects that you want to delete, choose **More**, and then choose **Delete**.



In the **Delete objects** dialog box, verify that the names of the folders you selected for deletion are listed and then choose **Delete**.



Related Topics

- [How Do I Delete Objects from an S3 Bucket? \(p. 44\)](#)

Making Folders Public

You can make a folder public, which means that all of the objects that appear within a public folder in the Amazon S3 console are available for viewing or downloading to anyone on the internet. However, as described in [How do I use folders in an S3 Bucket? \(p. 71\)](#), the S3 console supports a folder concept as a way to group objects. If you use a web browser to view a folder that you made public, you will get an access denied error because the folder is just a naming prefix for an object or group of objects.

Important

Use caution when making a folder public because anyone in the world with access to the internet can view or download the objects in your public folder. Also, it is easy to make a folder public, but you cannot make a folder private after you make it public. To make the objects in a public folder private, you must set permissions on each individual object so that the object has no public access. For more information about how to set an object's permissions, see [How Do I Set Permissions on an Object? \(p. 116\)](#).

More Info

- [How Do I Delete Folders from an S3 Bucket? \(p. 73\)](#)
- [How Do I Set ACL Bucket Permissions? \(p. 119\)](#)

Introduction to Amazon S3 Batch Operations

Amazon S3 batch operations performs large-scale batch operations on Amazon S3 objects. You can use Amazon S3 batch operations to copy objects, set object tags or access control lists (ACLs), initiate object restores from Amazon S3 Glacier, or invoke an AWS Lambda function to perform custom actions using your objects. You can perform these operations on a custom list of objects, or you can use an Amazon S3 inventory report to make generating even the largest lists of objects easy. Batch operations use the same Amazon S3 APIs that you already use, so you'll find the interface familiar. For information about performing batch operations using the AWS CLI, AWS SDKs, and the Amazon S3 REST APIs, see [Performing Batch Operations](#) in the *Amazon Simple Storage Service Developer Guide*.

The following topics explain how to use the Amazon S3 console to configure and run batch operations.

Topics

- [Creating a Batch Operations Job](#) (p. 76)
- [Managing Batch Operations Jobs](#) (p. 77)

Creating a Batch Operations Job

This section describes how to create a Amazon S3 batch operations job. For information about performing batch operations using the AWS CLI, AWS SDKs, and the Amazon S3 REST APIs, see [Performing Batch Operations](#) in the *Amazon Simple Storage Service Developer Guide*.

To create a batch job

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Choose **Batch Operations** on the navigation pane of the Amazon S3 console.
3. Choose **Create job**.

Amazon S3

Buckets

Batch operations

Public access settings for this account

The Amazon S3 batch operations feature is currently available in Preview and is not suitable for production workloads. For more information, see the [AWS Service Terms](#).

S3 batch operations

A job is used to execute batch operations on a list of S3 objects. The list of S3 objects is contained in a manifest object, which can be an S3 inventory report or a list of objects that you generate. After the total number of objects listed in the manifest has been confirmed, the job status will update to *Awaiting confirmation*, and you must **Confirm and run** the job within 14 days. Job events are published to [CloudWatch Events](#). Jobs are deleted 30 days after they finish. [Learn more](#)

Search by job ID or description

Status: All

Region: US East (N. Virginia)

+ Create job Clone job Update priority Confirm and run Cancel job

Job ID	Description	Operation	Date created	Status	Total objects	% Complete	Total failed (rate)	Priority
--------	-------------	-----------	--------------	--------	---------------	------------	---------------------	----------

4. Choose the **Region** where you want to create your job.
5. Under **Manifest format** choose the type of manifest object to use.

- If you choose **S3 Inventory report**, enter the path to the manifest.json object that Amazon S3 generated as part of the inventory report, and optionally the version ID for the manifest object if you want to use a version other than the most recent.
 - If you choose **CSV**, enter the path to a CSV-formatted manifest object. The manifest object must follow the format described in the console. You can optionally include the version ID for the manifest object if you want to use a version other than the most recent.
6. Under **Operation** choose the operation that you want to perform on all objects listed in the manifest. Fill out the information for the operation you chose and then choose **Next**.
 7. Fill out the information for **Configure additional options** and then choose **Next**.
 8. For **Review**, verify the settings. If you need to make changes, choose **Previous**. Otherwise, choose **Create Job**.

More Info

- [The Basics: Amazon S3 Batch Operations Jobs](#) in the *Amazon Simple Storage Service Developer Guide*
- [Creating a Batch Operations Job](#) in the *Amazon Simple Storage Service Developer Guide*
- [Operations](#) in the *Amazon Simple Storage Service Developer Guide*

Managing Batch Operations Jobs

Amazon S3 provides a set of tools to help you manage your batch operations jobs after you create them. For more information about managing batch operations, see [Managing Batch Operations Jobs](#) in the *Amazon Simple Storage Service Developer Guide*.

Storage Management

This section explains how to configure Amazon S3 storage management tools.

Topics

- [How Do I Create a Lifecycle Policy for an S3 Bucket? \(p. 78\)](#)
- [How Do I Add a Cross-Region Replication \(CRR\) Rule to an S3 Bucket? \(p. 81\)](#)
- [How Do I Manage the Cross-Region Replication Rules for an S3 Bucket? \(p. 95\)](#)
- [How Do I Configure Storage Class Analysis? \(p. 97\)](#)
- [How Do I Configure Amazon S3 Inventory? \(p. 101\)](#)
- [How Do I Configure Request Metrics for an S3 Bucket? \(p. 105\)](#)
- [How Do I Configure a Request Metrics Filter? \(p. 107\)](#)

How Do I Create a Lifecycle Policy for an S3 Bucket?


You can use lifecycle policies to define actions you want Amazon S3 to take during an object's lifetime (for example, transition objects to another storage class, archive them, or delete them after a specified period of time).

You can define a lifecycle policy for all objects or a subset of objects in the bucket by using a shared prefix (that is, objects that have names that begin with a common string).


A versioning-enabled bucket can have many versions of the same object, one current version and zero or more noncurrent (previous) versions. Using a lifecycle policy, you can define actions specific to current and noncurrent object versions. For more information, see [Object Lifecycle Management](#) and [Object Versioning](#) and [Using Versioning](#) in the *Amazon Simple Storage Service Developer Guide*.

To create a lifecycle policy

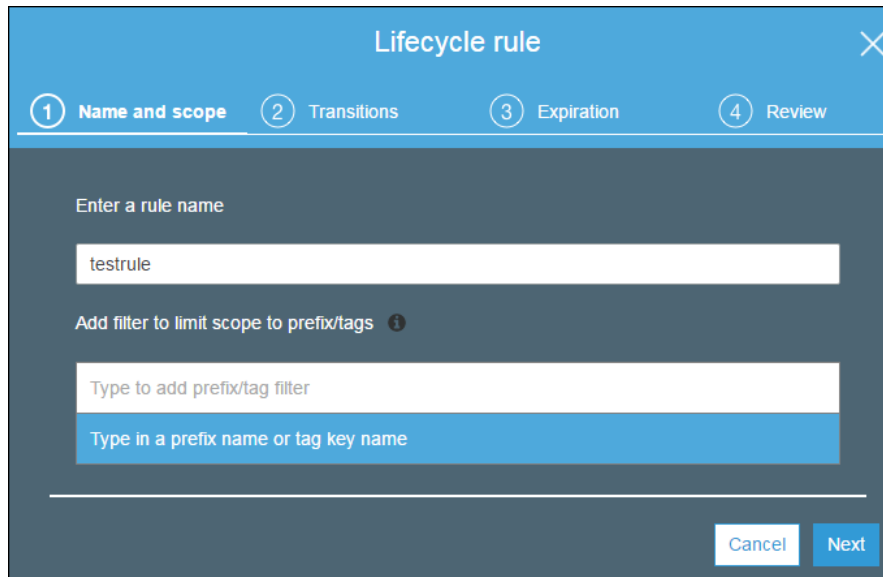
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to create a lifecycle policy for.



3. Choose the **Management** tab, and then choose **Add lifecycle rule**.



4. In the **Lifecycle rule** dialog box, type a name for your rule to help identify the rule later. The name must be unique within the bucket. Configure the rule as follows:
 - To apply this lifecycle rule to all objects with a specified name prefix (that is, objects with names that begin with a common string), type a prefix in the box, choose the prefix from the drop-down list, and then press **Enter**. For more information about object name prefixes, see [Object Keys](#) in the *Amazon Simple Storage Service Developer Guide*.
 - To apply this lifecycle rule to all objects with one or more object tags, type a tag in the box, choose the tag from the drop-down list, and then press **Enter**. Repeat the procedure to add another tag. You can combine a prefix and tags. For more information about object tags, see [Object Tagging](#) in the *Amazon Simple Storage Service Developer Guide*.
 - To apply this lifecycle rule to all objects in the bucket, choose **Next**.

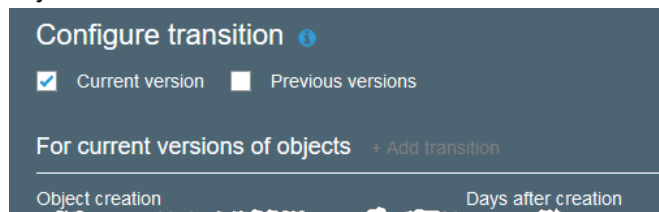


5. You configure lifecycle rules by defining rules to transition objects to the Standard-IA, One Zone-IA, and Glacier storage classes. For more information, see [Storage Classes](#) in the *Amazon Simple Storage Service Developer Guide*.

You can define transitions for current or previous object versions, or for both current and previous versions. Versioning enables you to keep multiple versions of an object in one bucket. For more information about versioning, see [How Do I Enable or Suspend Versioning for an S3 Bucket?](#) (p. 12).

- a. Select **Current version** to define transitions that are applied to the current version of the object.

Select **Previous versions** to define transitions that are applied to all previous versions of the object.



- b. Choose **Add transitions** and specify one of the following transitions:
 - Choose **Transition to Standard-IA after**, and then type the number of days after the creation of an object that you want the transition to be applied (for example, 30 days).

- Choose **Transition to One Zone-IA after**, and then type the number of days after the creation of an object that you want the transition to be applied (for example, 30 days).
- Choose **Transition to Glacier after**, and then type the number of days after the creation of an object that you want the transition to be applied (for example, 100 days).

Important

Amazon S3 stores the archived objects in Glacier. However, these are Amazon S3 objects, and you can access them only by using the Amazon S3 console or the Amazon S3 API. You cannot access the archived objects through the Glacier console or the Glacier API. For more information, see [Transitioning Objects: General Considerations](#).



6. When you are done configuring transitions, choose **Next**.

7. For this example, select both **Current version** and **Previous versions**.
8. Select **Expire current version of object**, and then enter the number of days after object creation to delete the object (for example, 395 days). If you select this expire option, you cannot select the option to clean up expired delete markers.
9. Select **Permanently delete previous versions**, and then enter the number of days after an object becomes a previous version to permanently delete the object (for example, 465 days).
10. It is a recommended best practice to always select **Clean up incomplete multipart uploads**. For example, type **7** for the number of days after the multipart upload initiation date that you want to end and clean up any multipart uploads that have not completed. For more information about

multipart uploads, see [Multipart Upload Overview](#) in the Amazon Simple Storage Service Developer Guide.

11. Choose **Next**.

Lifecycle rule

✓ Name and scope ✓ Transitions **3 Expiration** 4 Review

Configure expiration

☒ Current version ☒ Previous versions

☒ Expire current version of object ⓘ

After days from object creation

☒ Permanently delete previous versions ⓘ

After days from becoming a previous version

Clean up expired object delete markers and incomplete multipart uploads

☐ Clean up expired object delete markers ⓘ

You cannot enable clean up expired object delete markers if you enable Expiration.

☒ Clean up incomplete multipart uploads ⓘ

After Days from start of upload

[Previous](#) [Next](#)

12. For **Review**, verify the settings for your rule. If you need to make changes, choose **Previous**. Otherwise, choose **Save**.

13. If the rule does not contain any errors, it is listed on the **Lifecycle** page and is enabled.

+ Add lifecycle rule	Edit	Delete	More ▾
Lifecycle rule	Applied to	Actions for current version	Actions for previous version(s)
✓ testrule	Whole bucket	Standard-IA / Expire	Standard-IA / Amazon Glacier / Permanently Delete

How Do I Add a Cross-Region Replication (CRR) Rule to an S3 Bucket?

Cross-region replication is the automatic, asynchronous copying of objects across buckets in different AWS Regions. Cross-region replication replicates newly created objects and object updates from a source bucket to a destination bucket in a different AWS Region. For more information about cross-region replication concepts and how to use CRR with the AWS CLI, AWS SDKs, and the Amazon S3 REST APIs, see [Cross-Region Replication](#) in the *Amazon Simple Storage Service Developer Guide*.

Cross-region replication requires that the source and destination buckets be in different AWS Regions, and versioning must be enabled on both the source and destination buckets. To review the full list of

requirements, see [Requirements for Cross-Region Replication](#) in the *Amazon Simple Storage Service Developer Guide*. For more information about versioning, see [How Do I Enable or Suspend Versioning for an S3 Bucket?](#) (p. 12).

The object replicas in the destination bucket are exact replicas of the objects in the source bucket. They have the same key names and the same metadata—for example, creation time, owner, user-defined metadata, version ID, access control list (ACL), and storage class. Optionally, you can explicitly specify a different storage class for object replicas. And regardless of who owns the source bucket or the source object, you can choose to change replica ownership to the AWS account that owns the destination bucket. For more information, see [CRR: Change Replica Owner](#) in the *Amazon Simple Storage Service Developer Guide*.

The time it takes for Amazon S3 to replicate an object depends on the object size. It can take up to several hours to replicate a large-sized object.

Note about replication and lifecycle rules

Metadata for an object remains identical between original objects and replica objects. Lifecycle rules abide by the creation time of the original object, and not by when the replicated object becomes available in the destination bucket. However, lifecycle does not act on objects that are pending replication until replication is complete.

You use the Amazon S3 console to add replication rules to the source bucket. Replication rules define which source bucket objects to replicate and the destination bucket where the replicated objects are stored. You can create a rule to replicate all the objects in a bucket or a subset of objects with a specific key name prefix, one or more object tags, or both. A destination bucket can be in the same AWS account as the source bucket, or it can be in a different account. The destination bucket must always be in a different Region than the source bucket.

If the destination bucket is in a different account from the source bucket, you must add a bucket policy to the destination bucket to grant the owner of the source bucket account permission to replicate objects in the destination bucket. The Amazon S3 console builds this required bucket policy for you to copy and add to the destination bucket in the other account.

When you add a replication rule to a bucket, the rule is enabled by default, so it starts working as soon as you save it.


Topics

- [Adding a CRR Rule When the Destination Bucket Is in the Same AWS Account](#) (p. 82)
- [Adding a CRR Rule When the Destination Bucket Is in a Different AWS Account](#) (p. 89)
- [More Info](#) (p. 95)


Adding a CRR Rule When the Destination Bucket Is in the Same AWS Account

Follow these steps to configure a cross-region replication rule when the destination bucket is in the same AWS account as the source bucket.

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want.




3. Choose **Management**, choose **Replication**, and then choose **Add rule**.



4. In the **Replication rule** wizard, under **Set source**, you have the following options for setting the replication source:
- To replicate the whole bucket, choose **Entire bucket** **bucket-name**.
 - To replicate all objects that have the same prefix (for example, all objects that have names that begin with the string `pictures`), choose **Prefix or tags**. Enter a prefix in the box, choose the prefix from the drop-down list, and then press **Enter**. If you enter a prefix that is the name of a folder, you must use `/` (forward slash) as the last character (for example, `pictures/`). For more information about prefixes, see [Object Keys](#) in the *Amazon Simple Storage Service Developer Guide*.
 - To replicate all objects with one or more object tags, enter a tag in the box, choose the tag from the drop-down list, and then press **Enter**. Enter a tag value and then press **Enter**. Repeat the procedure to add another tag. You can combine a prefix and tags. For more information about object tags, see [Object Tagging](#) in the *Amazon Simple Storage Service Developer Guide*.

The screenshot shows the 'Replication rule' configuration window in the Amazon S3 console. The window has a blue header with the title 'Replication rule' and a close button (X). Below the header is a progress bar with four steps: 1. Set source, 2. Set destination, 3. Configure options, and 4. Review. Step 1 is currently active. The 'Set source' section has two radio button options: 'Entire bucket' (which is selected) and 'Prefix or tags'. The 'Entire bucket' option is accompanied by a lock icon and the text 'admin-created'. The 'Prefix or tags' option is accompanied by an information icon. Below these options are two text input fields. The first field is labeled 'Type to add prefix/tag filter' and the second field is labeled 'Type in a prefix name or tag key name'. The bottom of the window has a decorative torn paper effect.

You might get the following warning about your CRR rule being created with a new schema.

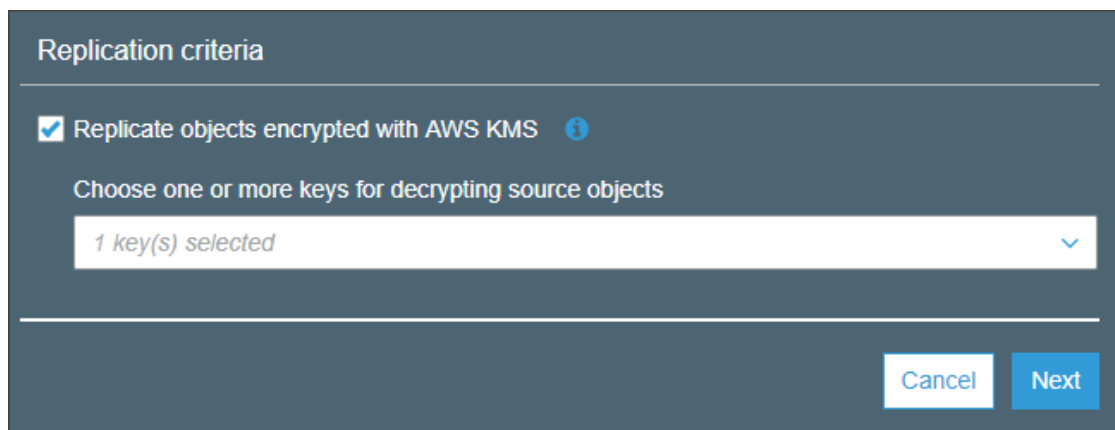
 **This action upgrades your CRR schema permanently**

Cross-region replication (CRR) now has a new schema that supports replication based on prefixes, one or more object tags or a combination of the two. As part of the new schema, you can set overlapping rules with priorities. The new schema does not support delete marker replication, which would prevent any delete actions from replicating. [Learn more](#)

The new schema supports prefix and tag filtering and the prioritization of rules. For more information about the new schema, see [Replication Configuration Backward Compatibility](#) in the *Amazon Simple Storage Service Developer Guide*. The developer guide describes the XML used with the Amazon S3 API that works behind the user interface. In the developer guide, the new schema is described as *replication configuration XML V2*.

5. To replicate objects in the source bucket that are encrypted with AWS KMS, under **Replication criteria**, select **Replicate objects encrypted with AWS KMS**. Under **Choose one or more keys for decrypting source objects** are the source AWS KMS key or keys that you allow cross-region replication to use. All source keys are included by default. You can choose to narrow the key selection.

Objects encrypted by AWS KMS keys that you do not select are not replicated by cross-region replication. A key or a group of keys is chosen for you, but you can choose the keys if you want. For information about using AWS KMS with cross-region replication, see [CRR: Replicating Objects Created with Server-Side Encryption \(SSE\) Using AWS KMS-Managed Encryption Key](#) in the *Amazon Simple Storage Service Developer Guide*.



Replication criteria

☒ Replicate objects encrypted with AWS KMS ⓘ

Choose one or more keys for decrypting source objects

1 key(s) selected ▼

Cancel Next

Important

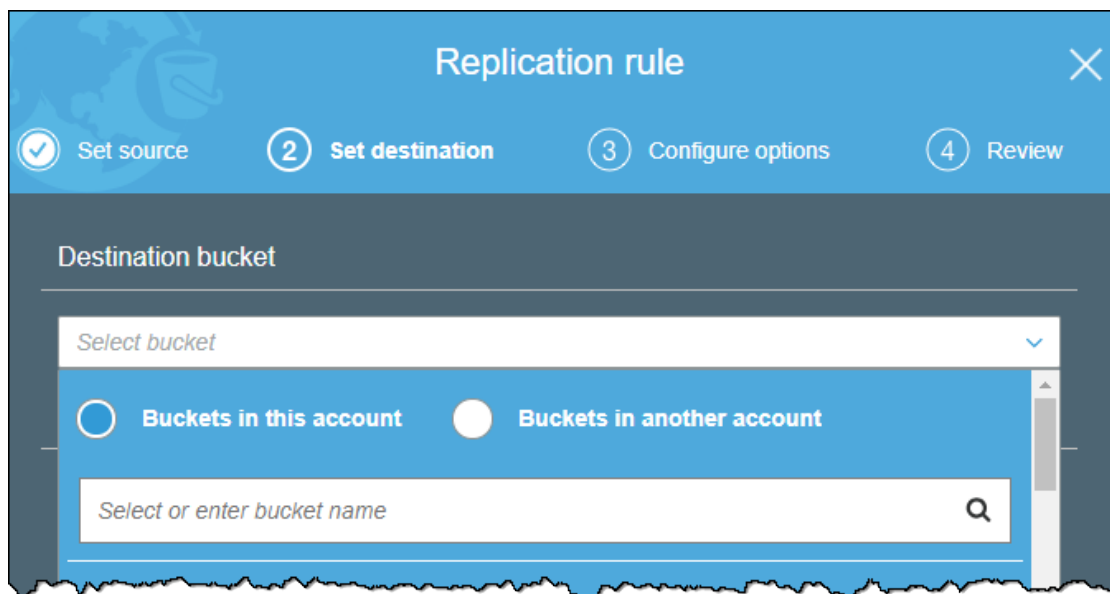
When you replicate objects that are encrypted with AWS KMS, the AWS KMS request rate doubles in the source Region and increases in the destination Region by the same amount. These increased call rates to AWS KMS are due to the way that data is re-encrypted using the customer master key (CMK) that you define for the cross-region replication destination Region. AWS KMS has a request rate limit that is per calling account per Region. For information about the limit defaults, see [AWS KMS Limits - Requests per second: varies](#). If your current Amazon S3 PUT object request rate during cross-region replication is more than half the default AWS KMS rate limit for your account, we recommend that you request an increase to your AWS KMS request rate limit. To request an increase, create a case in the AWS Support Center at [Contact Us](#). For example, suppose that your current PUT object request rate is 1,000 requests per second and you use AWS KMS to encrypt your objects. In this case, we recommend that you ask AWS Support to increase your AWS KMS rate limit to 2,500 requests per second, in both your source and destination Regions, to ensure that there is no throttling by AWS KMS.

To see your PUT object request rate in the source bucket, view `PutRequests` in the Amazon CloudWatch request metrics for Amazon S3. For information about viewing CloudWatch metrics, see [How Do I Configure Request Metrics for an S3 Bucket? \(p. 105\)](#).

Choose **Next**.

6. To choose a destination bucket from the account that you're currently using, on the **Set destination** page, under **Destination bucket**, choose **Buckets in this account**. Type the name of the destination bucket for the replication, or choose a name in the drop-down list. If you don't see the bucket that you want in the list, confirm that the bucket exists and that it's in a different Region than the source bucket.

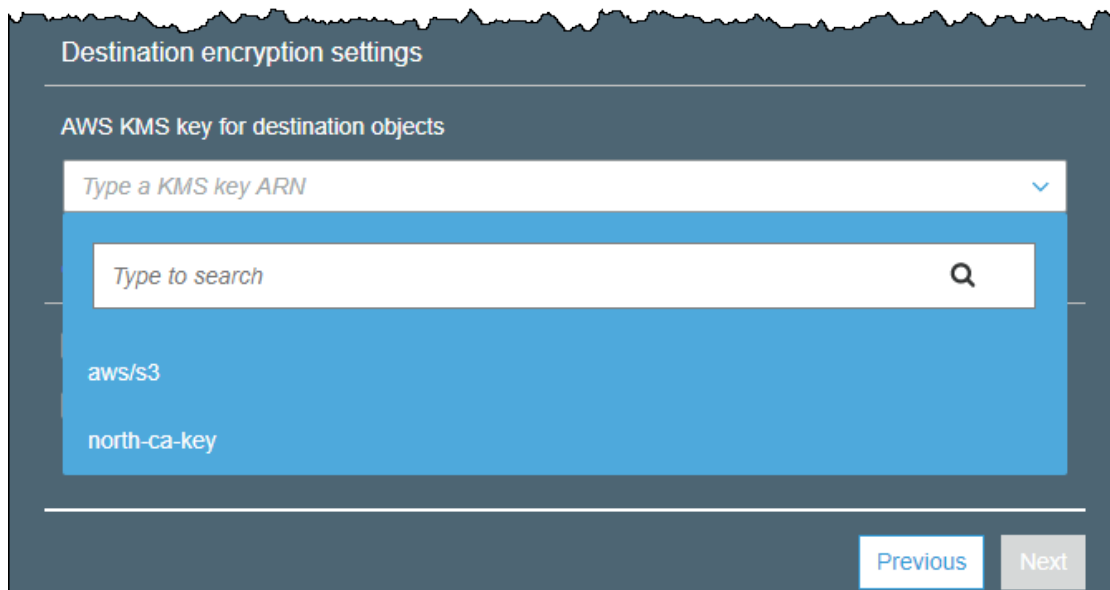
If you want to choose a destination bucket from a different AWS account, see [Adding a CRR Rule When the Destination Bucket Is in a Different AWS Account \(p. 89\)](#).



If versioning is not enabled on the destination bucket, you get a warning message that contains an **Enable versioning** button. Choose this button to enable versioning on the bucket.

7. If you chose to replicate objects encrypted with AWS KMS, under **Destination encryption settings**, type the Amazon Resource Name (ARN) of the AWS KMS key to use to encrypt the replicas in the destination bucket. You can find the ARN for your AWS KMS key in the IAM console, under **Encryption keys**. Or, you can choose a key name from the drop-down list.

For more information about creating an AWS KMS key, see [Creating Keys](#) in the *AWS Key Management Service Developer Guide*.



8. If you want to replicate your data into a specific storage class in the destination bucket, on the **Set destination** page, under **Options**, select **Change the storage class for the replicated object(s)**. Then choose the storage class that you want to use for the replicated objects in the destination bucket. If you don't select this option, the storage class for replicated objects is the same class as the original objects.

For information about **Change object ownership to the destination bucket owner**, see [Adding a CRR Rule When the Destination Bucket Is in a Different AWS Account](#) (p. 89).

Choose **Next**.

9. Set up an AWS Identity and Access Management (IAM) role that Amazon S3 can assume to perform cross-region replication of objects on your behalf.

To set up an IAM role, on the **Configure options** page, under **Select role**, do one of the following:

- We highly recommend that you choose **Create new role** to have Amazon S3 create a new IAM role for you. When you save the rule, a new policy is generated for the IAM role that matches the source and destination buckets that you choose. The name of the generated role is based on the bucket names and uses the following naming convention: **replication_role_for_source-bucket_to_destination-bucket**.
- You can choose to use an existing IAM role. If you do, you must choose a role that grants Amazon S3 the necessary permissions for replication. Replication fails if this role does not grant Amazon S3 sufficient permissions to follow your replication rule.

Important

When you add a replication rule to a bucket, you must have the `iam:PassRole` permission to be able to pass the IAM role that grants Amazon S3 replication permissions. For more information, see [Granting a User Permissions to Pass a Role to an AWS Service](#) in the *IAM User Guide*.

The screenshot shows the 'Replication rule' configuration window in the Amazon S3 console. The window has a blue header with the title 'Replication rule' and a close button (X). Below the header is a progress bar with four steps: 'Set source' (checked), 'Set destination' (checked), '3 Configure options' (active), and '4 Review'. The main content area is dark blue and contains two sections: 'IAM role' and 'Rule name'. Under 'IAM role', there is a dropdown menu with 'Create new role' selected. Under 'Rule name', there is a text input field with the placeholder text 'Type a name'.

Under **Rule name**, type a name for your rule to help identify the rule later. The name is required and must be unique within the bucket.

10. If the bucket has existing replication rules, you are instructed to set a priority for the rule. You must set a priority for the rule to avoid conflicts caused by objects that are included in the scope of more than one rule. In the case of overlapping rules, Amazon S3 uses the rule priority to determine which rule to apply. The higher the number, the higher the priority. For more information about rule priority, see [Replication Configuration Overview](#) in the *Amazon Simple Storage Service Developer Guide*.

Under **Status**, **Enabled** is selected by default. An enabled rule starts to work as soon as you save it. If you want to enable the rule later, select **Disabled**.

Choose **Next**.

Priority

Edit

Set the replication rule priority for objects that are included in multiple rules. You can change the priority after creating the rule.

☒ Highest (2)
 ☐ Lowest (0)

Status

☒ Enabled
 ☐ Disabled

Previous

Next

- On the **Review** page, review your replication rule. If it looks correct, choose **Save**. Otherwise, choose **Previous** to edit the rule before saving it.
- After you save your rule, you can edit, enable, disable, or delete your rule on the **Replication** page.

Lifecycle

Replication

Analytics

Metrics

Inventory

Cross-region replication updated successfully.

Source	Destination	Permissions
Scope All contents in the bucket	Bucket ca-example-bucket	IAM role s3crr_role_for_admin-created_to_ca-example-bucket
Region US West (Oregon)	Region US West (N. California)	Bucket policy Copy
Edit global settings		

+ Add rule

Edit priorities

Edit

Delete

Actions ▾

Viewing 1 to 1 of 1

Rule name	Scope ⓘ	Storage Class ⓘ	Replicated object owner ⓘ	KMS-Encrypted objects ⓘ	Status ⓘ	Priority ⓘ
<input type="radio"/> PictureRule	prefix: pictures, tag: cat manx	One Zone-IA	Same as source bucket	Do not replicate	Disabled	1

Adding a CRR Rule When the Destination Bucket Is in a Different AWS Account

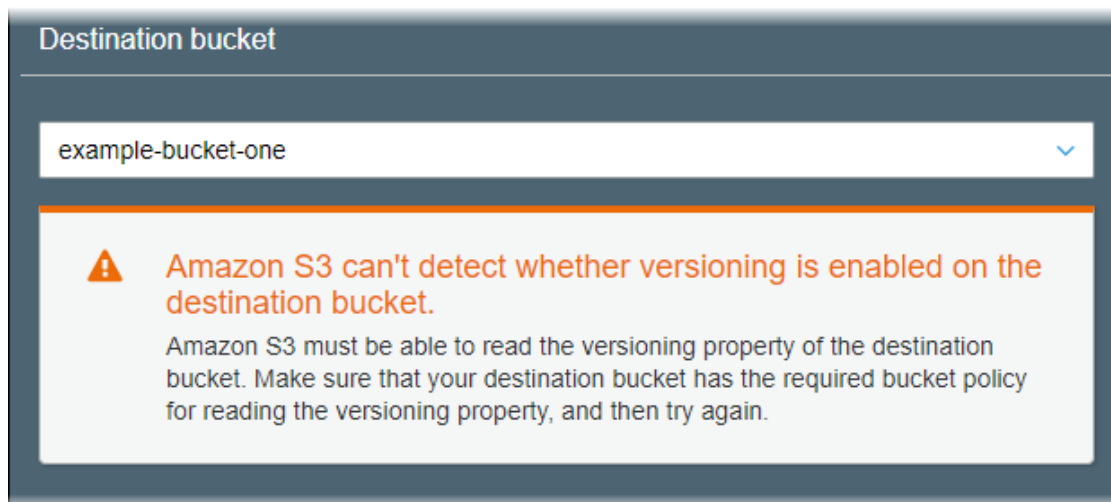
Follow these steps to configure a cross-region replication rule when the destination bucket is in a different AWS account than the source bucket.

1. If you have never created a cross-region replication rule before, start with [Adding a CRR Rule When the Destination Bucket Is in the Same AWS Account](#) (p. 82).

On the **Replication rule** wizard **Set destination** page, under **Destination bucket**, choose **Buckets in another account**. Then type the name of the destination bucket and the account ID from a different AWS account. Choose **Save**.


The screenshot shows the 'Replication rule' wizard in the Amazon S3 console, specifically the 'Set destination' step. The wizard has four steps: 'Set source' (completed), 'Set destination' (current), 'Configure options', and 'Review'. The 'Destination bucket' section is active, showing a dropdown menu with 'Select bucket' and a blue panel with two radio buttons: 'Buckets in this account' (selected) and 'Buckets in another account' (unselected). Below the radio buttons, there are input fields for 'Account ID' (containing '111122223333') and 'Bucket name' (containing 'example-bucket-one'). A 'Save' button is at the bottom of the panel.

After you save the destination bucket name and account ID, you might get a warning message indicating that you must add a bucket policy to the destination bucket so that Amazon S3 can verify whether versioning is enabled on the bucket. You can copy the bucket policy from the **Permissions** page, and then add the policy to the destination bucket in the other account. For information about adding a bucket policy to an S3 bucket, see [How Do I Add an S3 Bucket Policy?](#) (p. 123).



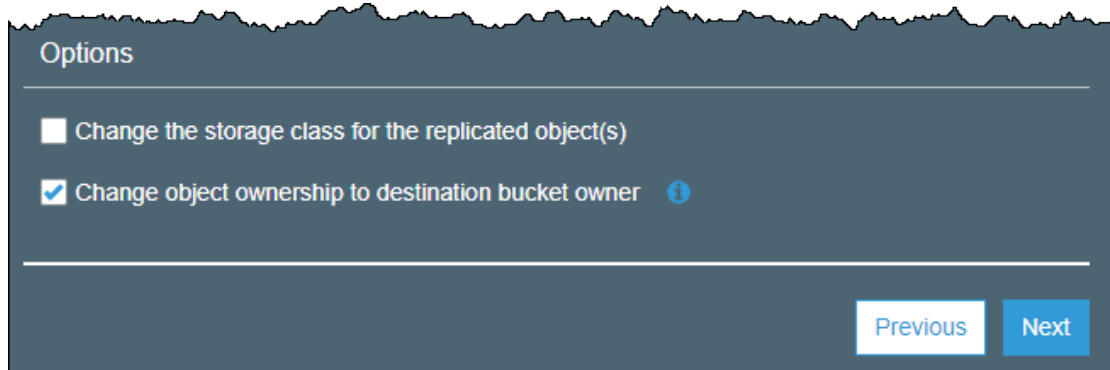
2. If you chose to replicate objects encrypted with AWS KMS, under **Destination encryption settings**, type the Amazon Resource Name (ARN) AWS KMS key to use to encrypt the replicas in the destination bucket.

For more information about creating an AWS KMS key, see [Creating Keys](#) in the *AWS Key Management Service Developer Guide*.



3. On the **Set destination** page, under **Options**:
 - To replicate your data into a specific storage class in the destination bucket, select **Change the storage class for the replicated object(s)**. Then choose the storage class that you want to use for the replicated objects in the destination bucket. If you don't select this option, the storage class for replicated objects is the same class as the original objects.
 - To change the object ownership of the replica objects to the destination bucket owner, select **Change object ownership to destination owner**. This option enables you to separate object ownership of the replicated data from the source. If asked, type the account ID of the destination bucket.

When you select this option, regardless of who owns the source bucket or the source object, the AWS account that owns the destination bucket is granted full permission to replica objects. For more information, see [CRR: Change Replica Owner](#) in the *Amazon Simple Storage Service Developer Guide*.



Options

☐ Change the storage class for the replicated object(s)

☒ Change object ownership to destination bucket owner ⓘ

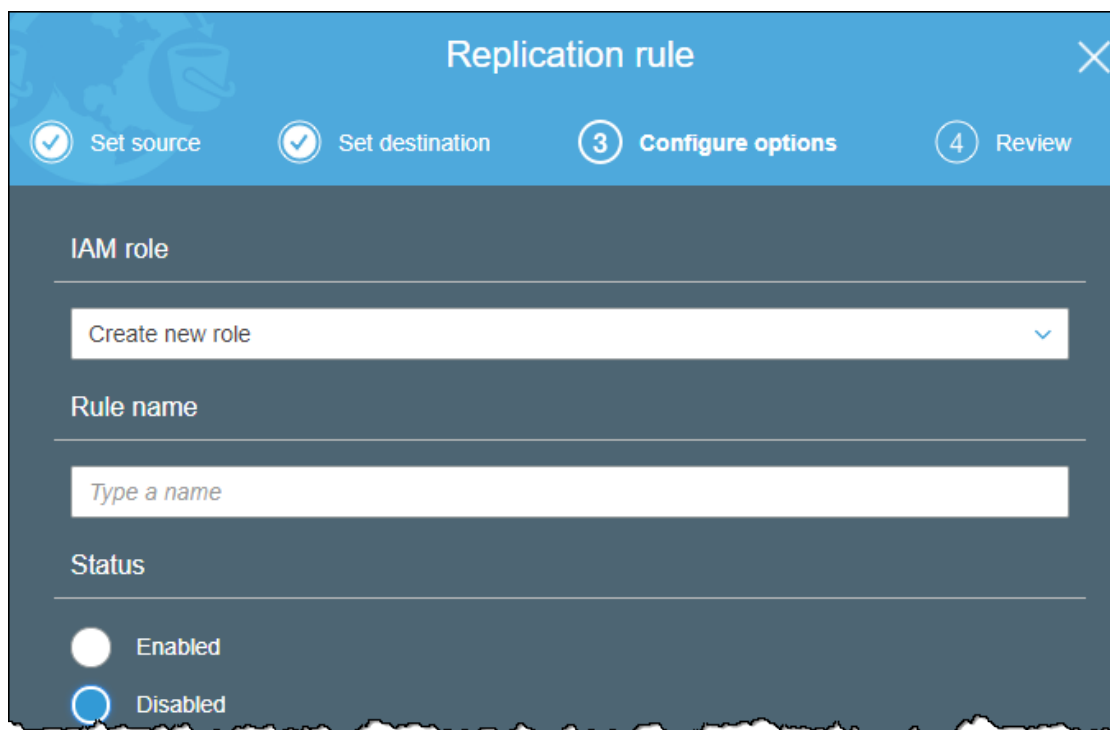
Previous Next

Choose **Next**.

4. Set up an AWS Identity and Access Management (IAM) role that Amazon S3 can assume to perform cross-region replication of objects on your behalf.

To set up an IAM role, on the **Configure options** page, under **Select role**, do one of the following:

- We highly recommend that you choose **Create new role** to have Amazon S3 create a new IAM role for you. When you save the rule, a new policy is generated for the IAM role that matches the source and destination buckets that you choose. The name of the generated role is based on the bucket names and uses the following naming convention: **replication_role_for_source-bucket_to_destination-bucket**.
- You can choose to use an existing IAM role. If you do, you must choose a role that allows Amazon S3 to replicate objects from the source bucket to the destination bucket on your behalf.



Replication rule

Set source Set destination 3 Configure options 4 Review

IAM role

Create new role

Rule name

Type a name

Status

Enabled Disabled

5. A bucket policy is provided on the **Configure options** page that you can copy and add to the destination bucket in the other account. For information about adding a bucket policy to an S3 bucket, see [How Do I Add an S3 Bucket Policy? \(p. 123\)](#).

Bucket policy

Amazon S3 requires you to put the following policy in the destination bucket.

Bucket policy [Copy](#)

```
1 {
2   "Version": "2008-10-17",
3   "Id": "S3-Console-Replication-Policy",
4   "Statement": [
5     {
6       "Sid": "S3ReplicationPolicyStmt1",
7       "Effect": "Allow",
8       "Principal": {
9         "AWS": "arn:aws:iam::444455556666 :root"
10      },
11      "Action": [
12        "s3:GetBucketVersioning",
13        "s3:PutBucketVersioning",
14        "s3:ReplicateObject",
15        "s3:ReplicateDelete"
16      ],
17      "Resource": [
```

[Previous](#) [Next](#)

6. If you chose to replicate objects encrypted with AWS KMS, an AWS KMS key policy is provided on the **Configure options** page. You can copy this policy to add to the key policy for the AWS KMS key customer master key (CMK) that you are using. The key policy grants the source bucket owner permission to use the key. For information about updating the key policy, see [Grant the Source Bucket Owner Permission to Encrypt Using the AWS KMS Key \(p. 94\)](#).

KMS key policy

Before saving KMS encryption settings, ensure that all producers and consumers have access to the specified encryption key. You can choose Copy to copy and paste policy to the KMS key. [Learn more](#)

KMS key policy [Copy](#)

```
1 {
2   "Sid": "Enable cross account encrypt access for S3 Cross Region Replication",
3   "Effect": "Allow",
4   "Principal": {
5     "AWS": "arn:aws:iam::[account-id] :root"
6   },
7   "Action": [
8     "kms:Encrypt"
9   ],
10  "Resource": "*"
11 }
```

[Previous](#) [Next](#)

7. On the **Review** page, review your replication rule. If it looks correct, choose **Save**. Otherwise, choose **Previous** to edit the rule before saving it.
8. After you save your rule, you can edit, enable, disable, or delete your rule on the **Replication** page.

The screenshot shows the Amazon S3 console's Replication page. At the top, there are tabs for Lifecycle, Replication (selected), Analytics, Metrics, and Inventory. A green success message states: "Cross-region replication updated successfully." Below this is an orange warning box that reads: "The CRR rule is saved, but additional settings are required in the destination account." The warning text continues: "Sign in to the destination account to enable versioning and add the required bucket policies. In the destination bucket, on the Management tab, choose Replication, and then choose **More > Receive objects**." A "Signout" button is located below the warning. The main content area displays the details of the replication rule "XAccountRule" in a table format:

Source	Destination	Permissions
Scope All contents in the bucket	Bucket example-bucket-one	IAM role s3crr_role_for_crrtag-destination_to_example-bucket-one
Region US West (N. California)	Region US West (Oregon)	Bucket policy Copy

Below the table are buttons: "+ Add rule", "Edit priorities", "Edit", "Delete", and "Actions". At the bottom, a table lists the rule's configuration:

Rule name	Scope	Storage Class	Replicated object owner	KMS-Encrypted objects	Status	Priority
XAccountRule	Entire bucket	Same as source	Destination bucket owner	Do not replicate	Disabled	1

9. Follow the instructions given on the Replication page under the warning message **The CRR rule is saved, but additional settings are required in the destination account**. Sign out of the AWS account that you are currently in, and then sign in to the destination account.

Important

Cross-region replication fails until you sign in to the destination account and complete the following steps.

10. After you sign in to the destination account, choose the **Management** tab, choose **Replication**, and then choose **Receive objects** from the **Actions** menu.
11. From the Receive objects page, you can perform the following:
 - Enable versioning on the destination bucket.
 - Apply the bucket policy provided by Amazon S3 to the destination bucket.
 - Copy the AWS KMS key policy that you need to update the AWS KMS CMK key that is being used to encrypt the replica objects in the destination bucket. For information about updating the key policy, see [Grant the Source Bucket Owner Permission to Encrypt Using the AWS KMS Key \(p. 94\)](#).

Receive objects

To use this bucket as the destination for replicated objects, type the your source account ID. Then configure the policy settings that are required to receive objects in this bucket.

Type the source account ID to customize

Versioning

If you don't have versioning enabled for this bucket, S3 enables it for you.

► **Bucket policy** Copy Apply settings

▼ **KMS policy** Copy

```
1 {
2   "Sid": "Enable cross account encrypt access for S3 Cross Region Replication",
3   "Effect": "Allow",
4   "Principal": {
5     "AWS": "SOURCE_ACCOUNT_ID"
6   },
7   "Action": [
8     "kms:Encrypt"
9   ],
10  "Resource": "*"
11 }
```

Done

Grant the Source Bucket Owner Permission to Encrypt Using the AWS KMS Key

You must grant permissions to the account of the source bucket owner to encrypt using your AWS KMS key with a key policy. The following procedure describes how to use the AWS Identity and Access Management (IAM) console to modify the key policy for the AWS KMS customer master key (CMK) that is being used to encrypt the replica objects in the destination bucket.

To grant permissions to encrypt using your AWS KMS key

1. Sign in to the AWS Management Console using the AWS account that owns the AWS KMS CMK. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the left navigation pane, choose **Encryption keys**.
3. For **Region**, choose the appropriate AWS Region. Do not use the region selector in the navigation bar (upper-right corner).

4. Choose the alias of the CMK that you want to encrypt with.
5. In the **Key Policy** section of the page, choose **Switch to policy view**.
6. Using the **Key Policy** editor, insert the key policy provided by Amazon S3 into the existing key policy, and then choose **Save Changes**. You might want to add the policy to the end of the existing policy.

For more information about creating and editing AWS KMS CMKs, see [Getting Started](#) in the *AWS Key Management Service Developer Guide*.

More Info

- [How Do I Manage the Cross-Region Replication Rules for an S3 Bucket?](#) (p. 95)
- [How Do I Enable or Suspend Versioning for an S3 Bucket?](#) (p. 12)
- [Cross-Region Replication](#) in the *Amazon Simple Storage Service Developer Guide*

How Do I Manage the Cross-Region Replication Rules for an S3 Bucket?


Cross-region replication is the automatic, asynchronous copying of objects across buckets in different AWS Regions. It replicates newly created objects and object updates from a source bucket to a destination bucket in a different Region.

You use the Amazon S3 console to add replication rules to the source bucket. Replication rules define the source bucket objects to replicate and the destination bucket where the replicated objects are stored. For more information about cross-region replication, see [Cross-Region Replication](#) in the *Amazon Simple Storage Service Developer Guide*.


You can manage replication rules on the **Replication** page. You can add, view, enable, disable, delete, and change the priority of the replication rules. For information about adding replication rules to a bucket, see [How Do I Add a Cross-Region Replication \(CRR\) Rule to an S3 Bucket?](#) (p. 81).

To manage the cross-region replication rules for an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want.



3. Choose **Management**, and then choose **Replication**.



4. You change the replication rules in the following ways.

- To change settings that affect all the replication rules in the bucket, choose **Edit global settings**.

Source	Destination	Permissions
Bucket admin-created	Bucket ca-example-bucket	IAM role s3crr_role_for_admin-created_to_ca-example-bucket
Region US West (Oregon)	Region US West (N. California)	Bucket policy Copy

[Edit global settings](#)

You can change the destination bucket, and the IAM role. If needed, you can copy the required bucket policy for cross-account destination buckets.

Source	Destination	Permissions
Bucket admin-created	Bucket <input type="text" value="ca-example-bucket"/>	IAM role <input type="text" value="s3crr_role_for_admin-..."/>
Region US West (Oregon)		Bucket policy Copy

[Cancel](#) [Save](#)

- To change a replication rule, select the rule and choose **Edit**, which starts the Replication wizard to help you make the change. For information about using the wizard, see [How Do I Add a Cross-Region Replication \(CRR\) Rule to an S3 Bucket?](#) (p. 81).

[+ Add rule](#) [Edit priorities](#) [Edit](#) [Delete](#) [Actions](#)

Scope	Storage Class	Replicated object owner
<input checked="" type="radio"/> Prefix : pictures	Standard-IA	Same as source bucket
<input type="radio"/> Prefix : examples	Standard	Same as source bucket

- To enable or disable a replication rule, select the rule, choose **More**, and in the drop-down list, choose **Enable rule** or **Disable rule**. You can also disable, enable, or delete all the rules in the bucket from the **More** drop-down list.


[+ Add rule](#) [Edit priorities](#) [Edit](#) [Delete](#) [Actions](#)

Scope	Storage Class	Replicated object owner
<input checked="" type="radio"/> Prefix : pictures	Standard-IA	Same as source bucket
<input type="radio"/> Prefix : examples	Standard	Same as source bucket

[Enable rule](#)
[Disable rule](#)
[Upgrade schema](#)
[Enable all rules](#)
[Disable all rules](#)
[Delete all rules](#)
[Receive objects...](#)

- To change the rule priorities, choose **Edit priorities**. You can then change the priority for each rule under the **Priority** column heading. Choose **Save** to save your changes.

You set rule priorities to avoid conflicts caused by objects that are included in the scope of more than one rule. In the case of overlapping rules, Amazon S3 uses the rule priority to determine which rule to apply. The higher the number, the higher the priority. For more information about rule priority, see [Replication Configuration Overview](#) in the *Amazon Simple Storage Service Developer Guide*.



More Info

- [How Do I Add a Cross-Region Replication \(CRR\) Rule to an S3 Bucket? \(p. 81\)](#)
- [Cross-Region Replication](#) in the *Amazon Simple Storage Service Developer Guide*

How Do I Configure Storage Class Analysis?

By using the Amazon S3 analytics storage class analysis tool, you can analyze storage access patterns to help you decide when to transition the right data to the right storage class. Storage class analysis observes data access patterns to help you determine when to transition less frequently accessed STANDARD storage to the STANDARD_IA (IA, for infrequent access) storage class. For more information about STANDARD_IA, see the [Amazon S3 FAQ](#) and [Storage Classes](#) in the *Amazon Simple Storage Service Developer Guide*.


Important

Storage class analysis does not give recommendations for transitions to the ONEZONE_IA or GLACIER storage classes.


For more information about analytics, see [Amazon S3 Analytics – Storage Class Analysis](#) in the *Amazon Simple Storage Service Developer Guide*.

To configure storage class analysis


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket for which you want to configure storage class analysis.



3. Choose the **Management** tab, and then choose **Analytics**.



4. Choose **Add**.



5. Type a name for the filter. If you want to analyze the whole bucket, leave the **Prefix / tags** field empty.

A screenshot of the 'Add filter' dialog box. It has a title bar 'Add filter'. Below the title bar, there is a 'Filter name' field with a placeholder text 'Enter a name for this filter'. Below that, there is a 'Prefix / tags to monitor (optional)' field with a placeholder text 'Leave empty for entire bucket'. At the bottom, there is a section titled 'Export data (optional)' with a plus icon. At the very bottom, there are two buttons: 'Save' and 'Cancel'.

6. In the **Prefix / tags** field, type text for the prefix or tag for the objects that you want to analyze, or choose from the dropdown list that appears when you start typing.

Add filter

Filter name

testfilter

Prefix / tags to monitor (optional)

images

prefix images (press enter)

tag images

► Export data (optional)

Save Cancel

7. If you chose **tag**, enter a value for the tag. You can enter one prefix and multiple tags.

Add filter

Filter name

testfilter

Prefix / tags to monitor (optional)

prefix videos X

tag dog | corgi X

tag dog | bulldog X

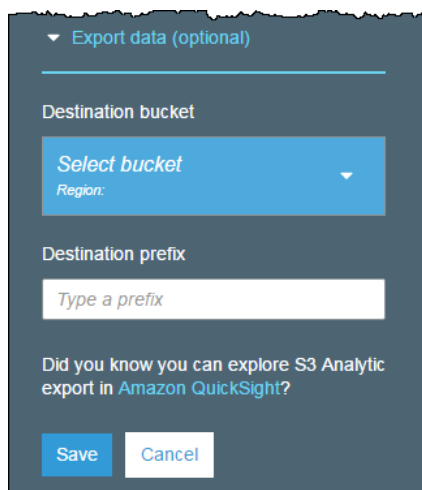
Leave empty for entire bucket

Type in a prefix name or tag key name

► Export data (optional)

Save Cancel

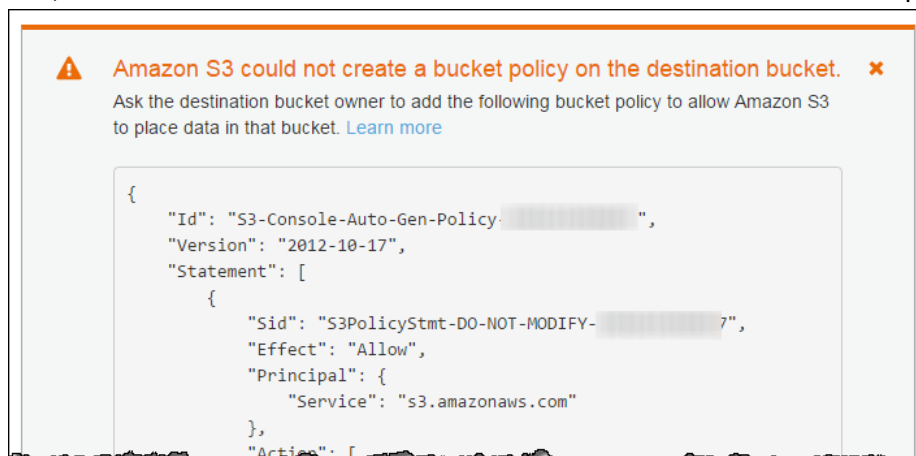
8. Optionally, you can choose **Export data** to export analysis reports to a comma-separated values (.csv) flat file. Choose a destination bucket where the file can be stored. You can type a prefix for the destination bucket. The destination bucket must be in the same AWS Region as the bucket for which you are setting up the analysis. The destination bucket can be in a different AWS account.



9. Choose **Save**.

Amazon S3 creates a bucket policy on the destination bucket that grants Amazon S3 write permission. This allow it to write the export data to the bucket.

If an error occurs when you try to create the bucket policy, you'll be given instructions on how to fix it. For example, if you chose a destination bucket in another AWS account and do not have permissions to read and write to the bucket policy, you'll see the following message. You must have the destination bucket owner add the displayed bucket policy to the destination bucket. If the policy is not added to the destination bucket you won't get the export data because Amazon S3 doesn't have permission to write to the destination bucket. If the source bucket is owned by a different account than that of the current user, then the correct account ID of the source bucket must be substituted in the policy.



For information about the exported data and how the filter works, see [Amazon S3 Analytics – Storage Class Analysis](#) in the *Amazon Simple Storage Service Developer Guide*.

More Info


[Storage Management \(p. 78\)](#)

How Do I Configure Amazon S3 Inventory?


Amazon S3 inventory provides a flat file list of your objects and metadata, which is a scheduled alternative to the Amazon S3 synchronous `List` API operation. Amazon S3 inventory provides comma-separated values (CSV) or [Apache optimized row columnar \(ORC\)](#) or [Apache Parquet \(Parquet\)](#) output files that list your objects and their corresponding metadata on a daily or weekly basis for an S3 bucket or for objects that share a prefix (objects that have names that begin with the same string). For more information, see [Amazon S3 Inventory](#) in the *Amazon Simple Storage Service Developer Guide*.

To configure inventory

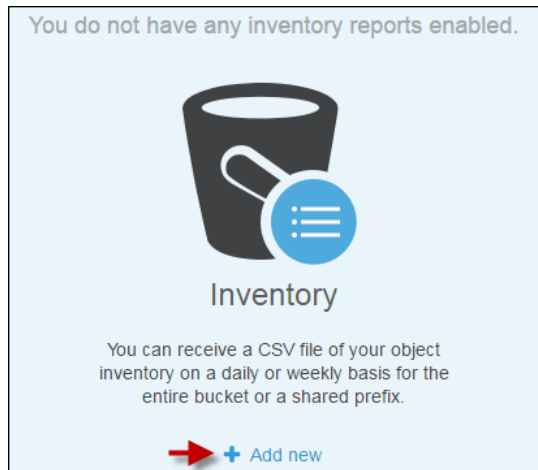
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket for which you want to configure Amazon S3 inventory.



3. Choose the **Management** tab, and then choose **Inventory**.



4. Choose **Add new** if you do not have any inventory reports enabled.



5. Type a name for the inventory and set it up as follows:
 - Optionally, add a prefix for your filter to inventory only objects whose names begin with the same string.
 - Choose the destination bucket where you want reports to be saved. The destination bucket must be in the same AWS Region as the bucket for which you are setting up the inventory. The destination bucket can be in a different AWS account.
 - Optionally, choose a prefix for the destination bucket.

- Choose how frequently to generate the inventory.

The screenshot shows the top section of the Amazon S3 Inventory configuration form. It includes buttons for '+ Add new', 'Edit', and 'Delete'. Below these are five main sections: 'Inventory name' with a text input 'Enter inventory name'; 'Filters' with a text input 'Filter by prefix (optional)'; 'Destination bucket' with a dropdown menu 'Select bucket'; 'Destination prefix' with a text input 'Type prefix (optional)'; and 'Frequency' with a dropdown menu 'Daily'.

6. Under **Advanced settings**, you can set the following:

- Choose either the CSV or ORC output file format for your inventory. For more information about these formats, see [Amazon S3 Inventory](#) in the *Amazon Simple Storage Service Developer Guide*.

The screenshot shows the 'Advanced settings' section of the Amazon S3 Inventory configuration form. It includes the following options:

- Output format:** Radio buttons for CSV (selected), Apache ORC, and Apache Parquet. A note for CSV says: 'Choose this format for listing 1 million or fewer objects, or if you plan to analyze S3 inventory with tools like Excel.' A 'Learn more' link is also present.
- Object versions:** A dropdown menu set to 'Current version only'.
- Optional fields:** A list of checkboxes: Size (checked), Last modified date (unchecked), Storage class (checked), Etag (unchecked), Multipart upload (unchecked), Replication status (unchecked), Encryption status (checked), and All object lock configurations (3 selected) (checked). Under 'All object lock configurations', there are three sub-checkboxes: Retention mode (checked), Retain until date (checked), and Legal hold status (checked).
- Encryption:** Radio buttons for None (selected), AES-256, and AWS-KMS.

 At the bottom are 'Cancel' and 'Save' buttons.

- To include all versions of the objects in the inventory, choose **Include all versions** in the **Object versions** list. By default, the inventory includes only the current versions of the objects.
- For **Optional fields**, select one or more of the following to add to the inventory report:
 - **Size** – Object size in bytes.
 - **Last modified date** – Object creation date or the last modified date, whichever is the latest.
 - **Storage class** – Storage class used for storing the object.
 - **ETag** – The entity tag is a hash of the object. The ETag reflects changes only to the contents of an object, and not its metadata. The ETag may or may not be an MD5 digest of the object data. Whether it is depends on how the object was created and how it is encrypted.
 - **Multipart upload** – Specifies that the object was uploaded as a multipart upload. For more information, see [Multipart Upload Overview](#) in the *Amazon Simple Storage Service Developer Guide*.
 - **Replication status** – The cross-region replication status of the object. For more information, see [How Do I Add a Cross-Region Replication \(CRR\) Rule to an S3 Bucket?](#) (p. 81)

- **Encryption status** – The server-side encryption used to encrypt the object. For more information, see [Protecting Data Using Server-Side Encryption](#) in the *Amazon Simple Storage Service Developer Guide*.
- **Object lock configurations** – The object lock status of the object, including the following settings:
 - **Retention mode** – The level of protection applied to the object, either *Governance* or *Compliance*.
 - **Retain until date** – The date until which the locked object cannot be deleted.
 - **Legal hold status** – The Legal hold status of the locked object.

For information about object lock, see [Amazon S3 Object Lock Overview](#) in the *Amazon Simple Storage Service Developer Guide*.

For more information about the contents of an inventory report, see [What's Included in an Amazon S3 Inventory?](#) in the *Amazon Simple Storage Service Developer Guide*.

- d. For **Encryption**, choose a server-side encryption option to encrypt the inventory report, or choose **None**:
- **None** – Do not encrypt the inventory report.
 - **AES-256** – Encrypt the inventory report using server-side encryption with Amazon S3-managed keys (SSE-S3). Amazon S3 server-side encryption uses 256-bit Advanced Encryption Standard (AES-256). For more information, see [Amazon S3-Managed Encryption Keys \(SSE-S3\)](#) in the *Amazon Simple Storage Service Developer Guide*.
 - **AWS-KMS** – Encrypt the report using server-side encryption with AWS KMS-managed keys (SSE-KMS). For more information, see [AWS KMS-Managed Keys \(SSE-KMS\)](#) in the *Amazon Simple Storage Service Developer Guide*.

Note

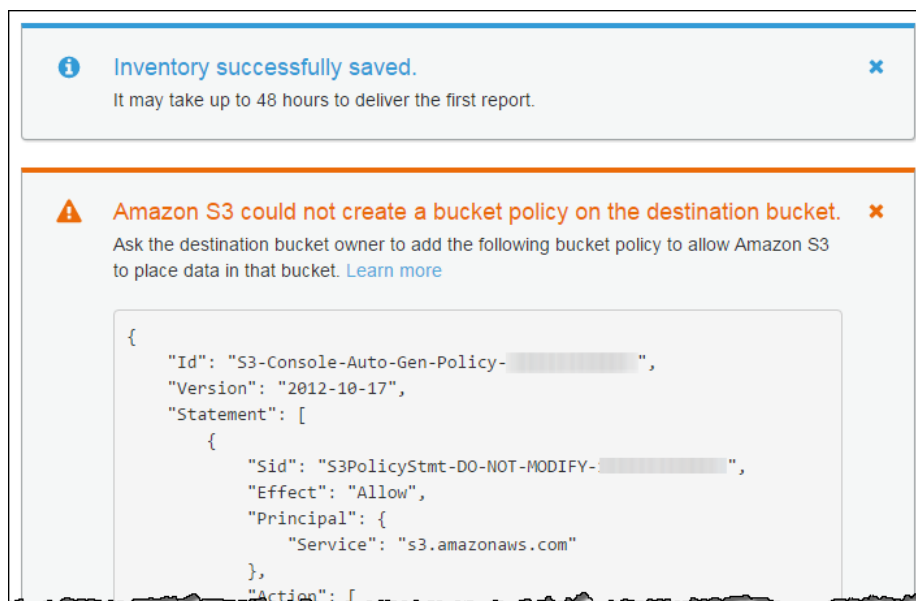
To encrypt the inventory list file with SSE-KMS, you must grant Amazon S3 permission to use the AWS KMS key. For instructions, see [Grant Amazon S3 Permission to Encrypt Using Your AWS KMS Key](#) (p. 104).

7. Choose **Save**.

Destination Bucket Policy

Amazon S3 creates a bucket policy on the destination bucket that grants Amazon S3 write permission. This allows Amazon S3 to write data for the inventory reports to the bucket.

If an error occurs when you try to create the bucket policy, you are given instructions on how to fix it. For example, if you choose a destination bucket in another AWS account and don't have permissions to read and write to the bucket policy, you see the following message.



In this case, the destination bucket owner must add the displayed bucket policy to the destination bucket. If the policy is not added to the destination bucket, you won't get an inventory report because Amazon S3 doesn't have permission to write to the destination bucket. If the source bucket is owned by a different account than that of the current user, the correct account ID of the source bucket must be substituted in the policy.

For more information, see [Amazon S3 Inventory](#) in the *Amazon Simple Storage Service Developer Guide*.

Grant Amazon S3 Permission to Encrypt Using Your AWS KMS Key

You must grant Amazon S3 permission to encrypt using your AWS KMS key with a key policy. The following procedure describes how to use the AWS Identity and Access Management (IAM) console to modify the key policy for the AWS KMS customer master key (CMK) that is being used to encrypt the inventory file.

To grant permissions to encrypt using your AWS KMS key

1. Sign in to the AWS Management Console using the AWS account that owns the AWS KMS CMK. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the left navigation pane, choose **Encryption keys**.
3. For **Region**, choose the appropriate AWS Region. Do not use the region selector in the navigation bar (upper-right corner).
4. Choose the alias of the CMK that you want to encrypt inventory with.
5. In the **Key Policy** section of the page, choose **Switch to policy view**.
6. Using the **Key Policy** editor, insert following key policy into the existing policy and then choose **Save Changes**. You might want to copy the policy to the end of the existing policy.

```
{
  "Sid": "Allow Amazon S3 use of the key",
  "Effect": "Allow",
  "Principal": {
    "Service": "s3.amazonaws.com"
  },
```

```
"Action": [  
    "kms:GenerateDataKey*",  
],  
"Resource": "*"
```

For more information about creating and editing AWS KMS CMKs, see [Getting Started](#) in the *AWS Key Management Service Developer Guide*.

More Info

[Storage Management](#) (p. 78)


How Do I Configure Request Metrics for an S3 Bucket?

There are two types of Amazon CloudWatch (CloudWatch) metrics for Amazon S3: storage metrics and request metrics. Storage metrics are reported once per day and are provided to all customers at no additional cost. Request metrics are available at 1-minute intervals after some latency to process, and metrics are billed at the standard CloudWatch rate. To get request metrics, you must opt into them by configuring them in the console or with the Amazon S3 API.


For more conceptual information about CloudWatch metrics for Amazon S3, see [Monitoring Metrics with Amazon CloudWatch](#) in the *Amazon Simple Storage Service Developer Guide*.

To configure request metrics on a bucket


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that has the objects you want to get request metrics for.



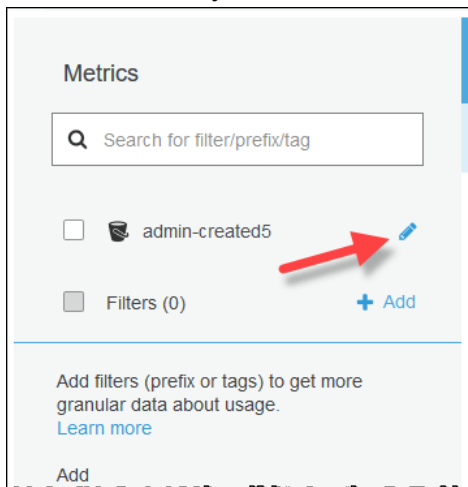
3. Choose the **Management** tab, and then choose **Metrics**.



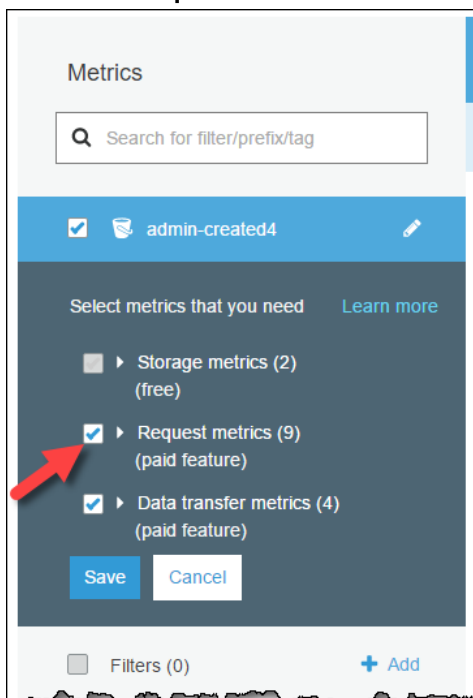
4. Choose **Requests**.



5. From the name of your bucket in the left-side pane, choose the edit icon.



6. Choose the **Request metrics** check box. This also enables Data Transfer metrics.



7. Choose **Save**.

You have now created a metrics configuration for all the objects in an Amazon S3 bucket. About 15 minutes after CloudWatch begins tracking these request metrics, you can see graphs for the metrics in both the Amazon S3 or CloudWatch consoles. You can also define a filter so the metrics are only collected and reported on a subset of objects in the bucket. For more information, see [How Do I Configure a Request Metrics Filter?](#) (p. 107).


How Do I Configure a Request Metrics Filter?

There are two types of CloudWatch metrics for Amazon S3: storage metrics and request metrics. Storage metrics are reported once per day and are provided to all customers at no additional cost. Request metrics are available at 1 minute intervals after some latency to process, and metrics are billed at the standard CloudWatch rate. To get request metrics, you must opt into them by configuring them in the console or with the Amazon S3 API.


For more conceptual information about CloudWatch metrics for Amazon S3, see [Monitoring Metrics with Amazon CloudWatch](#) in the *Amazon Simple Storage Service Developer Guide*.

To filter request metrics on a subset of objects in a bucket


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that has the objects you want to get request metrics for.




3. Choose the **Management** tab, and then choose **Metrics**.



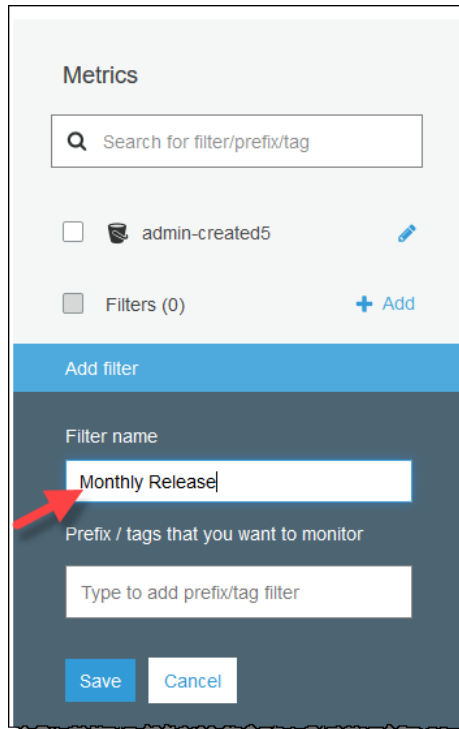
4. Choose **Requests**.



5. From **Filters** in the left-side pane, choose **Add**.

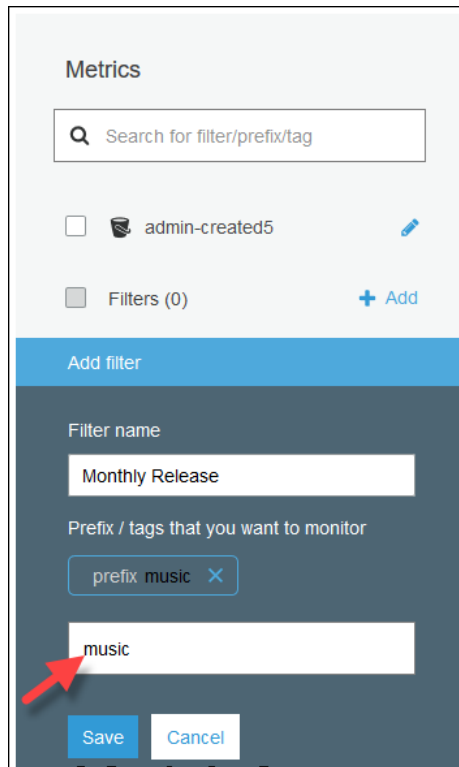


6. Provide a name for this metrics configuration.



The screenshot shows the 'Metrics' section of the Amazon S3 console. At the top, there is a search bar labeled 'Search for filter/prefix/tag'. Below it, there is a list of filters, currently empty, with a '+ Add' button. A modal window titled 'Add filter' is open. Inside the modal, the 'Filter name' field is highlighted with a red arrow and contains the text 'Monthly Release'. Below this, there is a section for 'Prefix / tags that you want to monitor' with a text input field labeled 'Type to add prefix/tag filter'. At the bottom of the modal are 'Save' and 'Cancel' buttons.

7. Provide one or more prefixes or tags, separated by commas, in **Prefix /tags that you want to monitor**. From the drop down, select whether the value you provided is a tag or a prefix.



The screenshot shows the 'Metrics' section of the Amazon S3 console. At the top, there is a search bar labeled 'Search for filter/prefix/tag'. Below it, there is a list of filters, currently empty, with a '+ Add' button. A modal window titled 'Add filter' is open. Inside the modal, the 'Filter name' field contains the text 'Monthly Release'. Below this, there is a section for 'Prefix / tags that you want to monitor'. This section contains a dropdown menu with the text 'prefix music' and a red 'X' icon. Below the dropdown is a text input field containing the text 'music', which is highlighted with a red arrow. At the bottom of the modal are 'Save' and 'Cancel' buttons.

8. Choose **Save**.

You have now created a metrics configuration for request metrics on a subset of the objects in an Amazon S3 bucket. About 15 minutes after CloudWatch begins tracking these request metrics, you can see graphs for the metrics in both the Amazon S3 or CloudWatch consoles. You can also request metrics at the bucket level. For information, see [How Do I Configure Request Metrics for an S3 Bucket? \(p. 105\)](#)

Setting Bucket and Object Access Permissions

This section explains how to use the Amazon Simple Storage Service (Amazon S3) console to grant access permissions to your buckets and objects. It also explains how to use Amazon S3 block public access to prevent the application of any settings that allow public access to data within S3 buckets.

Buckets and objects are Amazon S3 resources. You grant access permissions to your buckets and objects by using resource-based access policies. You can associate an access policy with a resource. An access policy describes who has access to resources. The resource owner is the AWS account that creates the resource. For more information about resource ownership and access policies, see [Overview of Managing Access](#) in the *Amazon Simple Storage Service Developer Guide*.

Bucket access permissions specify which users are allowed access to the objects in a bucket and which types of access they have. *Object access permissions* specify which users are allowed access to the object and which types of access they have. For example, one user might have only read permission, while another might have read and write permissions.

Bucket and object permissions are independent of each other. An object does not inherit the permissions from its bucket. For example, if you create a bucket and grant write access to a user, you will not be able to access that user's objects unless the user explicitly grants you access.

To grant access to your buckets and objects to other AWS accounts and to the general public, you use resource-based access policies called access control lists (ACLs).

A *bucket policy* is a resource-based AWS Identity and Access Management (IAM) policy that grants other AWS accounts or IAM users access to an S3 bucket. Bucket policies supplement, and in many cases, replace ACL-based access policies. For more information on using IAM with Amazon S3, see [Managing Access Permissions to Your Amazon S3 Resources](#) in the *Amazon Simple Storage Service Developer Guide*.

For more in-depth information about managing access permissions, see [Introduction to Managing Access Permissions to Your Amazon S3 Resources](#) in the *Amazon Simple Storage Service Developer Guide*.

This section also explains how to use the Amazon S3 console to add a cross-origin resource sharing (CORS) configuration to an S3 bucket. CORS allows client web applications that are loaded in one domain to interact with resources in another domain.

Topics

- [How Do I Block Public Access to S3 Buckets? \(p. 110\)](#)
- [How Do I Edit Public Access Settings for S3 Buckets? \(p. 112\)](#)
- [How Do I Edit Public Access Settings for All the S3 Buckets in an AWS Account? \(p. 115\)](#)
- [How Do I Set Permissions on an Object? \(p. 116\)](#)
- [How Do I Set ACL Bucket Permissions? \(p. 119\)](#)
- [How Do I Add an S3 Bucket Policy? \(p. 123\)](#)
- [How Do I Allow Cross-Domain Resource Sharing with CORS? \(p. 124\)](#)

How Do I Block Public Access to S3 Buckets?

Amazon S3 block public access prevents the application of any settings that allow public access to data within S3 buckets. You can configure block public access settings for an individual S3 bucket or for all

the buckets in your account. For information about blocking public access using the AWS CLI, AWS SDKs, and the Amazon S3 REST APIs, see [Using Amazon S3 Block Public Access](#) in the *Amazon Simple Storage Service Developer Guide*.

The following topics explain how to use the Amazon S3 console to configure block public access settings:

- [How Do I Edit Public Access Settings for S3 Buckets?](#) (p. 112)
- [How Do I Edit Public Access Settings for All the S3 Buckets in an AWS Account?](#) (p. 115)

The following sections explain viewing bucket access status and searching by access types.

Viewing Access Status


The list buckets view shows whether your bucket is publicly accessible. Amazon S3 labels the permissions for a bucket as follows:

- **Public** – Everyone has access to one or more of the following: List objects, Write objects, Read and write permissions.
- **Objects can be public** – The bucket is not public, but anyone with the appropriate permissions can grant public access to objects.
- **Buckets and objects not public** – The bucket and objects do not have any public access.
- **Only authorized users of this account** – Access is isolated to IAM users and roles in this account and AWS service principals because there is a policy that grants public access.

The access column shows the access status of the listed buckets.

<input type="checkbox"/> Bucket name	Access	Region
<input type="checkbox"/> zzz-example-bucket	Bucket and objects not public	US East (N. Virginia)
<input type="checkbox"/> zz-example-bucket	Only authorized users of this account	US East (N. Virginia)
<input type="checkbox"/> example-bucket-77	Error	US East (N. Virginia)

You can also filter bucket searches by access type. Choose an access type from the drop-down list that is next to the **Search for buckets** bar.



More Info

- [How Do I Edit Public Access Settings for S3 Buckets?](#) (p. 112)

- [How Do I Edit Public Access Settings for All the S3 Buckets in an AWS Account?](#) (p. 115)
- [Setting Bucket and Object Access Permissions](#) (p. 110)

How Do I Edit Public Access Settings for S3 Buckets?

Amazon S3 block public access prevents the application of any settings that allow public access to data within S3 buckets. This section describes how to edit block public access settings for one or more S3 buckets. For information about blocking public access using the AWS CLI, AWS SDKs, and the Amazon S3 REST APIs, see [Using Amazon S3 Block Public Access](#) in the *Amazon Simple Storage Service Developer Guide*.

Topics


- [Editing Public Access Settings for an S3 Bucket](#) (p. 112)
- [Editing Public Access Settings for Multiple S3 Buckets](#) (p. 113)
- [More Info](#) (p. 114)

Editing Public Access Settings for an S3 Bucket


Follow these steps if you need to change the public access settings for a single S3 bucket.

To edit the block public access settings for an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want.



3. Choose **Permissions**.



4. Choose **Edit** to change the public access settings for the bucket.

Manage public access control lists (ACLs)	Manage public bucket policies	Edit
Block new public ACLs and uploading public objects <i>(Recommended)</i> True	Block new public bucket policies <i>(Recommended)</i> True	
Remove public access granted through public ACLs <i>(Recommended)</i> True	Block public and cross-account access if bucket has public policies <i>(Recommended)</i> True	


5. Choose the setting that you want to change, and then choose **Save**. For details about each setting, pause on the **i** icons.

Manage public access control lists (ACLs) for this bucket
Access control lists (ACLs) is an access policy option to grant basic read/write permissions to other AWS accounts.

☒ Block new public ACLs and uploading public objects *(Recommended)* **i**
☒ Remove public access granted through public ACLs *(Recommended)* **i**

Manage public bucket policies for this bucket
Bucket policies use JSON-based access policy language to manage advanced permission to your Amazon S3 resources.

☒ Block new public bucket policies *(Recommended)* **i**
☒ Block public and cross-account access if bucket has public policies *(Recommended)* **i**

 Access status appears as **Only authorized users of this account** for a bucket with a public bucket policy. By choosing this option, users outside this account cannot access that bucket until the public bucket policy is removed.

Refresh Cancel Save

Before you save your changes, you can choose **Refresh** to restore the settings to what they were before you started making changes.

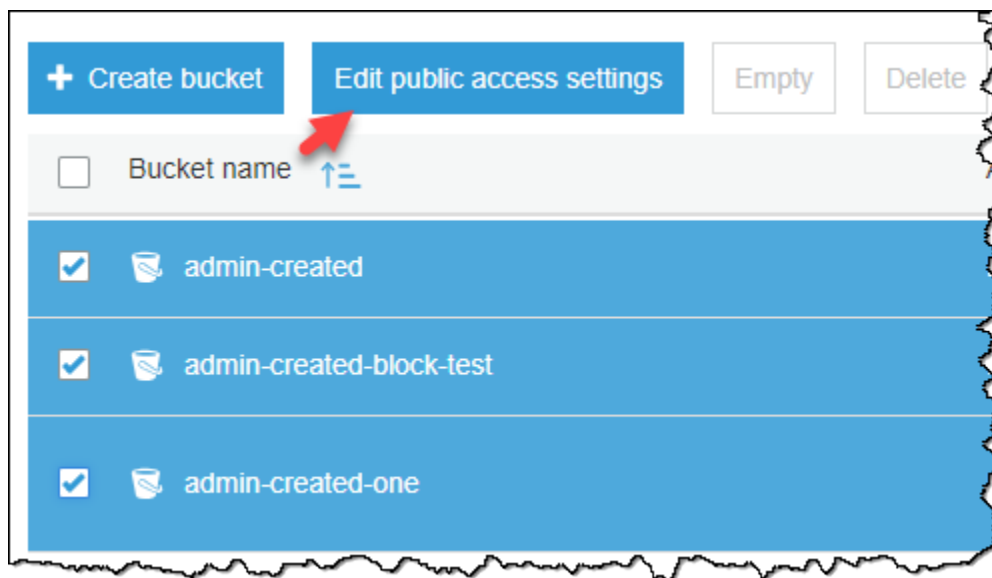
6. When you're asked for confirmation, enter **confirm**. Then choose **Confirm** to save your changes.

Editing Public Access Settings for Multiple S3 Buckets


Follow these steps if you need to change the public access settings for more than one S3 bucket.

To edit the block public access settings for multiple S3 buckets

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the buckets that you want, and then choose **Edit public access settings**.



3. Choose the setting that you want to change, and then choose **Save**. For details about each block public access setting, pause on the **i** icons.



4. When you're asked for confirmation, enter **confirm**. Then choose **Confirm** to save your changes.

You can change block public access settings when you create a bucket. For more information, see [How Do I Create an S3 Bucket?](#) (p. 3).

More Info

- [How Do I Block Public Access to S3 Buckets?](#) (p. 110)
- [How Do I Edit Public Access Settings for All the S3 Buckets in an AWS Account?](#) (p. 115)
- [Setting Bucket and Object Access Permissions](#) (p. 110)

How Do I Edit Public Access Settings for All the S3 Buckets in an AWS Account?

Amazon S3 block public access prevents the application of any settings that allow public access to data within S3 buckets. This section describes how to edit block public access settings for all the S3 buckets in your AWS account. For information about blocking public using the AWS CLI, AWS SDKs, and the Amazon S3 REST APIs, see [Using Amazon S3 Block Public Access](#) in the *Amazon Simple Storage Service Developer Guide*.

To edit block public access settings for all the S3 buckets in an AWS account

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Choose **Public access settings for this account**.



3. Choose **Edit** to change the block public access settings for all the buckets in your AWS account.

Manage public access control lists (ACLs)	Manage public bucket policies	Edit
Block new public ACLs and uploading public objects <i>(Recommended)</i> False	Block new public bucket policies <i>(Recommended)</i> False	
Remove public access granted through public ACLs <i>(Recommended)</i> False	Block public and cross-account access to buckets that have public policies <i>(Recommended)</i> False	

4. Choose the settings you want to change, and then choose **Save**. For details about each setting, pause on the **i** icons.

Manage public access control lists (ACLs) for this account

Access control lists (ACLs) is an access policy option to grant basic read/write permissions to other AWS accounts.

☐ Block new public ACLs and uploading public objects *(Recommended)* ⓘ

☐ Remove public access granted through public ACLs *(Recommended)* ⓘ

Manage public bucket policies for this account

Bucket policies use JSON-based access policy language to manage advanced permission to your Amazon S3 resources.

☐ Block new public bucket policies *(Recommended)* ⓘ

☐ Block public and cross-account access to buckets that have public policies *(Recommended)* ⓘ

Refresh

Cancel

Save

Before you save your changes, you can choose **Refresh** to restore the settings to what they were before you started making changes.

- When you're asked for confirmation, enter **confirm**. Then choose **Confirm** to save your changes.

More Info

- [How Do I Block Public Access to S3 Buckets?](#) (p. 110)
- [How Do I Edit Public Access Settings for S3 Buckets?](#) (p. 112)
- [Setting Bucket and Object Access Permissions](#) (p. 110)

How Do I Set Permissions on an Object?

This section explains how to use the Amazon Simple Storage Service (Amazon S3) console to manage access permissions for an S3 object by using access control lists (ACLs). ACLs are resource-based access policies that grant access permissions to buckets and objects. For more information about managing access permissions with resource-based policies, see [Overview of Managing Access](#) in the *Amazon Simple Storage Service Developer Guide*.

Bucket and object permissions are independent of each other. An object does not inherit the permissions from its bucket. For example, if you create a bucket and grant write access to a user, you can't access that user's objects unless the user explicitly grants you access.


You can grant permissions to other AWS accounts or predefined groups. The user or group that you grant permissions to is called the grantee. By default, the owner, which is the AWS account that created the bucket, has full permissions.

Each permission you grant for a user or a group adds an entry in the ACL that is associated with the object. The ACL lists grants, which identify the grantee and the permission granted. For more information about ACLs, see [Managing Access with ACLs](#) in the *Amazon Simple Storage Service Developer Guide*.


To set permissions for an object

- Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.


2. In the **Bucket name** list, choose the name of the bucket that contains the object.



3. In the **Name** list, choose the name of the object for which you want to set permissions.



4. Choose **Permissions**.



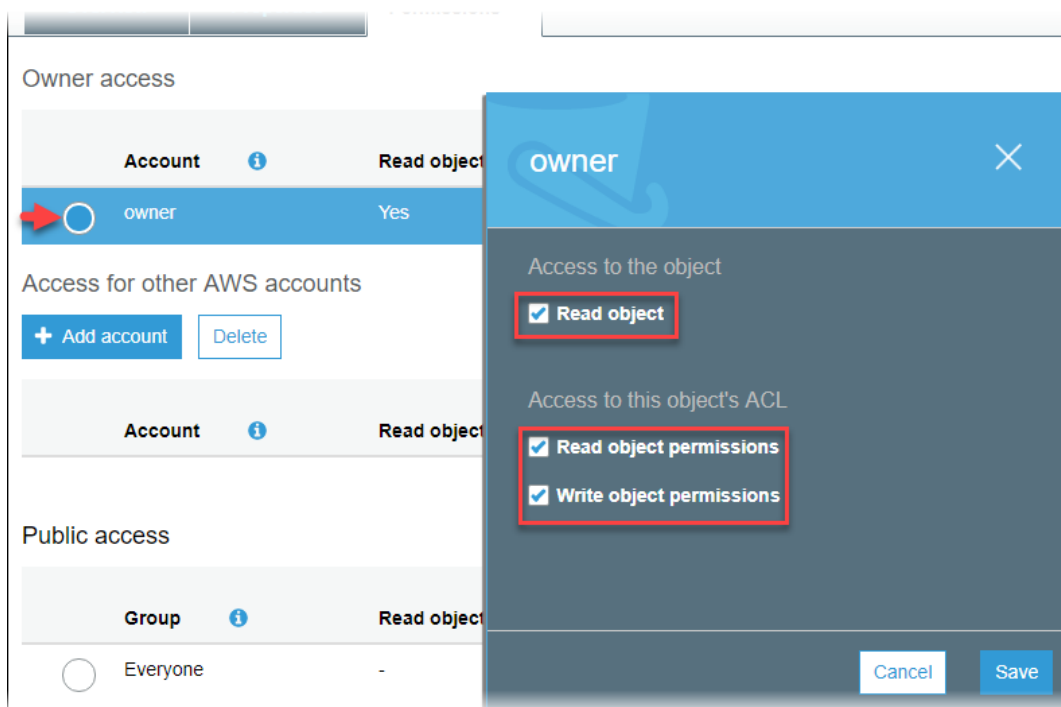
5. You can manage object access permissions for the following:

a. **Owner access**

The *owner* refers to the AWS account root user, and not an AWS Identity and Access Management (IAM) user. For more information about the root user, see [The AWS Account Root User](#).

To make changes to the owner's object access permissions, under **Owner access**, choose the account name, which is the name of the AWS account root user.

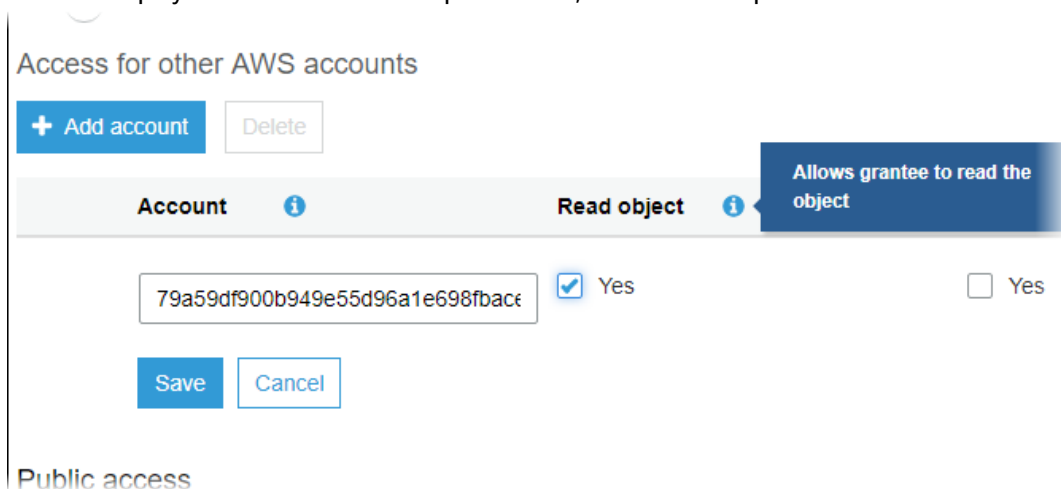
Select the check boxes for the permissions that you want to change, and then choose **Save**.



b. **Access for other AWS accounts**

To grant permissions to an AWS user from a different AWS account, under **Access for other AWS accounts**, choose **Add account**. In the **Enter an ID** field, type the canonical ID of the AWS user that you want to grant object permissions to. For information about finding a canonical ID, see [AWS Account Identifiers](#) in the *Amazon Web Services General Reference*. You can add as many as 99 users.

Select the check boxes for the permissions that you want to grant to the user, and then choose **Save**. To display information about the permissions, choose the Help icons.



c. **Public access**

To grant access to your object to the general public (everyone in the world), under **Public access**, choose **Everyone**. Granting public access permissions means that anyone in the world can access the object.

Select the check boxes for the permissions that you want to grant, and then choose **Save**.

The screenshot shows the 'Permissions' tab in the Amazon S3 console. On the left, under 'Public access', the 'Everyone' group is selected. A modal dialog titled 'Everyone' is open, displaying a warning: 'This object will have public access. Everyone will have access to one or all of the following: read this object, read and write permissions.' Below the warning, under 'Access to the object', the 'Read object' checkbox is checked. Under 'Access to this object's ACL', the 'Read object permissions' and 'Write object permissions' checkboxes are unchecked. At the bottom right of the modal, the 'Save' button is highlighted in blue.

Warning

Use caution when granting the **Everyone** group anonymous access to your S3 objects. When you grant access to this group, anyone in the world can access your object. If you need to grant access to everyone, we highly recommend that you only grant permissions to **Read objects**. We highly recommend that you *do not* grant the **Everyone** group write object permissions. Doing so allows anyone to overwrite the ACL permissions for the object.

You can also set object permissions when you upload objects. For more information about setting permissions when uploading objects, see [How Do I Upload Files and Folders to an S3 Bucket? \(p. 33\)](#).

More Info

- [Setting Bucket and Object Access Permissions \(p. 110\)](#)
- [How Do I Set ACL Bucket Permissions? \(p. 119\)](#)

How Do I Set ACL Bucket Permissions?


This section explains how to use the Amazon Simple Storage Service (Amazon S3) console to manage access permissions for S3 buckets by using access control lists (ACLs). ACLs are resource-based access policies that grant access permissions to buckets and objects. For more information about managing access permissions with resource-based policies, see [Overview of Managing Access](#) in the *Amazon Simple Storage Service Developer Guide*.

You can grant permissions to other AWS account users or to predefined groups. The user or group that you are granting permissions to is called the grantee. By default, the owner, which is the AWS account that created the bucket, has full permissions.


Each permission you grant for a user or group adds an entry in the ACL associated with the bucket. The ACL lists grants, which identify the grantee and the permission granted. For more information about ACLs, see [Managing Access with ACLs](#) in the *Amazon Simple Storage Service Developer Guide*.

To set ACL access permissions for an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to set permissions for.



3. Choose **Permissions**.

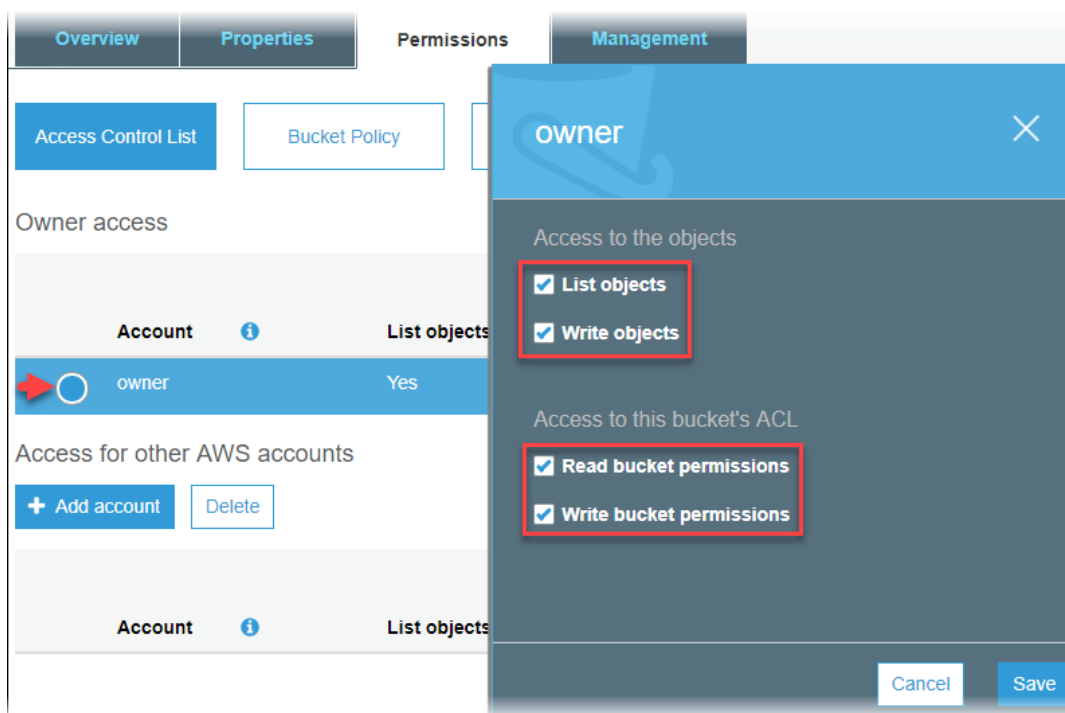


4. You can manage bucket access permissions for the following:
 - a. **Owner access**

The *owner* refers to the AWS account root user, and not an AWS Identity and Access Management (IAM) user. For more information about the root user, see [The AWS Account Root User](#).

To make changes to the owner's bucket access permissions, under **Owner access**, choose the account name, which is the name of the AWS account root user.


Select the check boxes for the permissions that you want to change, and then choose **Save**.



b. **Access for other AWS accounts**

To grant permissions to an AWS user from a different AWS account, under **Access for other AWS accounts**, choose **Add account**. In the **Enter an ID** field, type the canonical ID of the AWS user that you want to grant bucket permissions to. For information about finding a canonical ID, see [AWS Account Identifiers](#) in the *Amazon Web Services General Reference*. You can add as many as 99 users.

Select the check boxes next to the permissions that you want to grant to the user, and then choose **Save**. To display information about the permissions, choose the Help icons.



Warning

When you grant other AWS accounts access to your resources, be aware that the AWS accounts can delegate their permissions to users under their accounts. This is known as *cross-account access*. For information about using cross-account access, see [Creating a Role to Delegate Permissions to an IAM User](#) in the *IAM User Guide*.

c. **Public access**

To grant access to your bucket to the general public (everyone in the world), under **Public access**, choose **Everyone**. Granting public access permissions means that anyone in the world can access the bucket.

Select the check boxes for the permissions that you want to grant, and then choose **Save**.

The screenshot shows the 'Public access' section of the Amazon S3 console. Under 'Public access', the 'Everyone' group is selected. A modal dialog is open, displaying a warning: 'This bucket will have public access. Everyone will have access to one or all of the following: list objects, write objects, read and write permissions.' The dialog also shows checkboxes for 'List objects' (checked), 'Write objects', 'Read bucket permissions', and 'Write bucket permissions'. The 'Save' button is highlighted.

Warning

Use caution when granting the **Everyone** group public access to your S3 bucket. When you grant access to this group, anyone in the world can access your bucket. We highly recommend that you never grant any kind of public write access to your S3 bucket.

d. **S3 log delivery group**

To grant access to Amazon S3 to write server access logs to the bucket, under **S3 log delivery group**, choose **Log Delivery**.

If a bucket is set up as the target bucket to receive access logs, the bucket permissions must allow the **Log Delivery** group write access to the bucket. When you enable server access logging on a bucket, the S3 console grants write access to the **Log Delivery** group for the target bucket that you choose to receive the logs. For more information about server access logging, see [How Do I Enable Server Access Logging for an S3 Bucket? \(p. 16\)](#).

You can also set bucket permissions when you are creating a bucket. For more information on setting permissions when creating a bucket, see [How Do I Create an S3 Bucket? \(p. 3\)](#).

More Info

- [Setting Bucket and Object Access Permissions \(p. 110\)](#)

- [How Do I Set Permissions on an Object?](#) (p. 116)
- [How Do I Add an S3 Bucket Policy?](#) (p. 123)


How Do I Add an S3 Bucket Policy?

This section explains how to use the Amazon Simple Storage Service (Amazon S3) console to add a new bucket policy or edit an existing bucket policy. A bucket policy is a resource-based AWS Identity and Access Management (IAM) policy. You add a bucket policy to a bucket to grant other AWS accounts or IAM users access permissions for the bucket and the objects in it. Object permissions apply only to the objects that the bucket owner creates. For more information about bucket policies, see [Overview of Managing Access](#) in the *Amazon Simple Storage Service Developer Guide*.


For examples of Amazon S3 bucket policies, see [Bucket Policy Examples](#) in the *Amazon Simple Storage Service Developer Guide*.

To create or edit a bucket policy

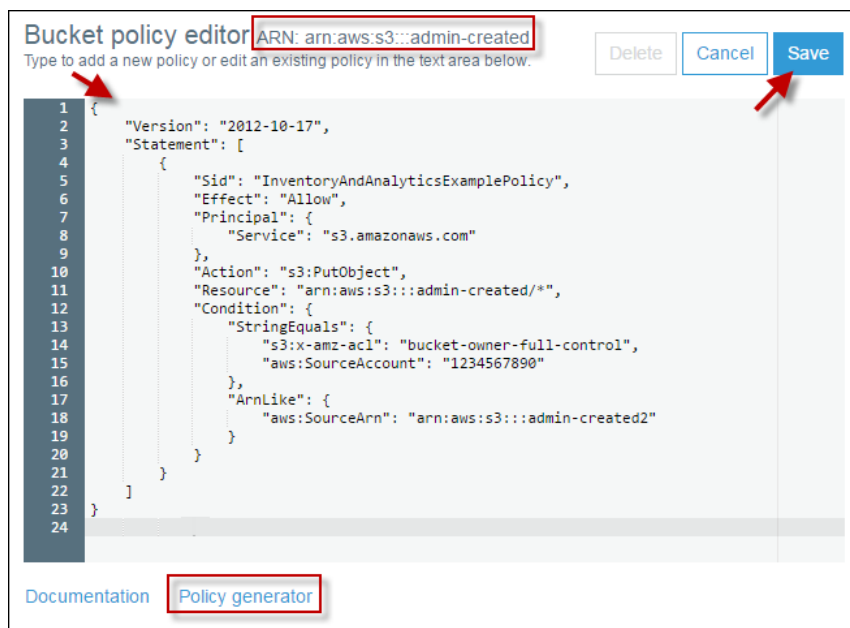
1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to create a bucket policy for or whose bucket policy you want to edit.



3. Choose **Permissions**, and then choose **Bucket Policy**.



4. In the **Bucket policy editor** text box, type or copy and paste a new bucket policy, or edit an existing policy. The bucket policy is a JSON file. The text you type in the editor must be valid JSON.



5. Choose **Save**.

Note

Amazon S3 displays the Amazon Resource Name (ARN) for the bucket next to the **Bucket policy editor** title. For more information about ARNs, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#) in the *Amazon Web Services General Reference*.

Directly below the bucket policy editor text box is a link to the **Policy Generator**, which you can use to create a bucket policy.

More Info

- [Setting Bucket and Object Access Permissions \(p. 110\)](#)
- [How Do I Set ACL Bucket Permissions? \(p. 119\)](#)

How Do I Allow Cross-Domain Resource Sharing with CORS?


This section explains how to use the Amazon S3 console to add a cross-origin resource sharing (CORS) configuration to an S3 bucket. CORS allows client web applications that are loaded in one domain to interact with resources in another domain.

To configure your bucket to allow cross-origin requests, you add CORS configuration to the bucket. A CORS configuration is an XML document that defines rules that identify the origins that you will allow to access your bucket, the operations (HTTP methods) supported for each origin, and other operation-specific information. For more information about CORS, see [Cross-Origin Resource Sharing \(CORS\)](#) in the *Amazon Simple Storage Service Developer Guide*.


When you enable CORS on the bucket, the access control lists (ACLs) and other access permission policies continue to apply.

To add a CORS configuration to an S3 bucket


1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. In the **Bucket name** list, choose the name of the bucket that you want to create a bucket policy for.



3. Choose **Permissions**, and then choose **CORS configuration**.



4. In the **CORS configuration editor** text box, type or copy and paste a new CORS configuration, or edit an existing configuration. The CORS configuration is an XML file. The text that you type in the editor must be valid XML.



5. Choose **Save**.

Note

Amazon S3 displays the Amazon Resource Name (ARN) for the bucket next to the **CORS configuration editor** title. For more information about ARNs, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#) in the *Amazon Web Services General Reference*.

More Info

- [Setting Bucket and Object Access Permissions](#) (p. 110)
- [How Do I Set ACL Bucket Permissions?](#) (p. 119)
- [How Do I Add an S3 Bucket Policy?](#) (p. 123)

Document History

Latest documentation update: November 15, 2018

The following table describes the important changes in each release of the *Amazon Simple Storage Service Console User Guide* from June 19, 2018, onward. For notification about updates to this documentation, you can subscribe to an RSS feed.

update-history-change	update-history-description	update-history-date
Blocking public access to S3 buckets (p. 126)	Amazon S3 block public access prevents the application of any settings that allow public access to data within S3 buckets. For more information, see Blocking Public Access to S3 Buckets .	November 15, 2018
Filtering enhancements in cross-region replication (CRR) rules (p. 126)	In a CRR rule, you can specify an object filter to choose a subset of objects to apply the rule to. Previously, you could filter only on an object key prefix. In this release, you can filter on an object key prefix, one or more object tags, or both. For more information, see How Do I Add a Cross-Region Replication (CRR) Rule to an S3 Bucket? .	September 19, 2018
Updates now available over RSS (p. 126)	You can now subscribe to an RSS feed to receive notifications about updates to the Amazon Simple Storage Service Console User Guide guide.	June 19, 2018

Earlier Updates

The following table describes the important changes in each release of the *Amazon Simple Storage Service Console User Guide* before June 19, 2018.

Change	Description	Date Changed
New storage class	Amazon S3 now offers a new storage class, ONEZONE_IA (IA, for infrequent access) for storing objects. For more information, see Storage Classes in the <i>Amazon Simple Storage Service Developer Guide</i> .	April 4, 2018
Support for ORC-formatted Amazon S3 inventory files	Amazon S3 now supports the Apache optimized row columnar (ORC) format in addition to comma-separated values (CSV) file format for inventory output files. For more information, see How Do I Configure Amazon S3 Inventory? (p. 101).	November 17, 2017

Change	Description	Date Changed
Bucket permissions check	Bucket permissions check in the Amazon S3 console checks bucket policies and bucket access control lists (ACLs) to identify publicly accessible buckets. Bucket permissions check makes it easier to identify S3 buckets that provide public read and write access.	November 06, 2017
Default encryption for S3 buckets	Amazon S3 default encryption provides a way to set the default encryption behavior for an S3 bucket. You can set default encryption on a bucket so that all objects are encrypted when they are stored in the bucket. The objects are encrypted using server-side encryption with either Amazon S3-managed keys (SSE-S3) or AWS KMS-managed keys (SSE-KMS). For more information, see How Do I Enable Default Encryption for an S3 Bucket? (p. 13) .	November 06, 2017
Encryption status in Amazon S3 inventory	Amazon S3 now supports including encryption status in Amazon S3 inventory so you can see how your objects are encrypted at rest for compliance auditing or other purposes. You can also configure to encrypt Amazon S3 inventory with server-side encryption (SSE) or SSE-KMS so that all inventory files are encrypted accordingly. For more information, see How Do I Configure Amazon S3 Inventory? (p. 101) .	November 06, 2017
Cross-region replication enhancements	Cross-region replication now supports the following: <ul style="list-style-type: none"> By default, Amazon S3 does not replicate objects in your source bucket that are created using server-side encryption using AWS KMS-managed keys. You can now configure a replication rule to replicate these objects. For more information, see How Do I Add a Cross-Region Replication (CRR) Rule to an S3 Bucket? (p. 81). In a cross-account scenario, you can configure a replication rule to change replica ownership to the AWS account that owns the destination bucket. For more information, see Adding a CRR Rule When the Destination Bucket Is in a Different AWS Account (p. 89). 	November 06, 2017
Added functionality and documentation	The Amazon S3 console now supports enabling object-level logging for an S3 bucket with AWS CloudTrail data events logging. For more information, see How Do I Enable Object-Level Logging for an S3 Bucket with AWS CloudTrail Data Events? (p. 19) .	October 19, 2017
Old Amazon S3 console no longer available	The old version of the Amazon S3 console is no longer available and the old user guide was removed from the Amazon S3 documentation site.	August 31, 2017
General availability of New Amazon S3 console	Announced the general availability of the new Amazon S3 console.	May 15, 2017

AWS Glossary

For the latest AWS terminology, see the [AWS Glossary](#) in the *AWS General Reference*.